



How to Create AI Agents Like a Pro

An interactive learning experience
brought to you by **Pearson**

Presented by Tim Warner
 **TechTrainerTim.com**




Course Flow (Central Time)

09:00	AM	Segment 1	[Microsoft Agentic AI + Inception]
10:00	AM	Segment 2	[Building a Low-Code Agent]
11:00	AM	Segment 3	[Extending Our Agent's Capabilities]
12:00	PM	Segment 4	[Operating and Deploying Our Agent]
01:00	PM	Finish	

Session Materials

github.com/timothywarner-org/agents-pro



Session Recording + Badge

 Explore Skills Start Learning Featured  Answers



All events Your events **Your recordings**

Live Events

Access recordings of events you've registered for here approximately 24 hours after the events have ended.

 **Business Intelligence**
Microsoft Fabric Jumpstart
Build Your Analytics Solution and Transform Data into Insights
 Tim Warner
You attended July 15, 2024

 **Prompt Engineering**
The Elements of Prompt Engineering
Craft effective prompts for generative AI to elevate the quality and reliability of outputs
 Tim Warner
You attended March 5, 2024

 **AZ-400: Designing and Implementing Microsoft DevOps Solutions**
Azure Certified DevOps Engineer Crash Course
Prepare for Microsoft Certification Exam AZ-400
 Mike Pfeiffer
You attended Feb 11, 2020

1 - 3 of 3

Session Recording + Badge

O'REILLY

Explore Skills

Start Learning

Featured

Answers

Search for books, courses, events, and more

< VIEW ALL EVENTS

Azure Certified DevOps Engineer Crash Course

Published by [O'Reilly Media, Inc.](#)

Advanced

Prepare for Microsoft Certification Exam AZ-400

Congrats!

Completed on Nov 21, 2021

[What you'll learn](#)

[Is this live event for you?](#)

[Schedule](#)

In this online course, Microsoft MVP Mike Pfeiffer walks you through what to expect on the AZ-400 Microsoft Azure DevOps Engineer Expert exam. This Azure certification will validate your Azure skills that combine people, process, and technologies to continuously deliver value to end-user needs and business objectives.

Even if you're not interested in certification, you will learn how DevOps professionals working in the Azure cloud streamline the delivery of applications and infrastructure. Walk away with a game plan on how to study for your exam, but also implement strategies for application and infrastructure code that allow for continuous integration, continuous testing, continuous delivery, and continuous monitoring and feedback.

What you'll learn and how you can apply it

By the end of this live, hands-on, online course, you'll understand:

- How the Azure role-based certifications work
- What topics are most likely to appear on the AZ-400 certification exam
- How to do the work of the Microsoft DevOps Engineer

And you'll be able to:

- Complete your studies to prepare for the AZ-400 certification exam
- Work towards Azure certification to prove your skills and increase your professional credibility
- Use the Azure portal and Azure PowerShell to deploy and manage resources in Azure

You attended

Feb 11, 2020

11am-3pm Central Standard Time

View session 1 video

Your Instructor

Mike Pfeiffer

Mike Pfeiffer is the founder and chief technologist at cloud consulting and training firm CloudSkills.io. A 20-year tech industry veteran, he's worked for some of the largest technology companies in the world, including Microsoft and Amazon Web Services (AWS). Mike is a publishe...

[Read more](#)

in

Skill covered

AZ-400: Designing and Implementing Microsoft DevOps Solutions

Earn your badge


Digital badges are verifiable and shareable proof of the skills you've built. Complete 80% or more of this course to earn your badge via Credly.

Session Recording + Badge

Media Player


Why Git?

Feature branches




A diagram showing a workflow where a new feature is developed on a separate branch (Feature) and then merged back into the main branch (Master) once ready.

Distributed development




A diagram showing a distributed development workflow where multiple developers (Full Repo) can work on their own copies of the repository and push changes back to the central repository.

Pull requests




A diagram showing a workflow where a developer creates a pull request to merge changes from a feature branch into the main branch (Master), allowing for review and approval before deployment.

Community



A diagram showing a community workflow where multiple developers (represented by icons) can contribute to a shared project, with arrows indicating the flow of code and collaboration.

Release cycles



A diagram showing a release cycle workflow where a product is released to customers, and feedback is collected from them to inform future releases. The cycle involves 'You' (developers), 'Product', 'Customers', and 'Feedback'.




75:24/235:11

GROUP CHAT ? YouTube Docs Screen Share User Q&A

POWERED BY ON24

Playback Speed
0.5
0.75
Normal


Session Recordings & Badge



Session Recording + Badge



Discover badges, skills or organizations



This badge was issued to [Tim Warner](#) | [View celebrations](#)

Date issued: April 13, 2025

Share



AZ-400 Crash Course

Issued by [O'Reilly Media](#)

The holder of this badge has completed the following course by O'Reilly Media, Inc. The holder of this badge is capable of streamlining the delivery of applications and infrastructure in the Azure cloud, implementing strategies for continuous integration, continuous testing, continuous delivery, and continuous monitoring and feedback.

[Learn more](#)

Learning

Advanced

Hours

\$ Paid

Skills

AZ-400: Designing and Implementing Microsoft DevOps...

Cloud Computing

Cloud Platforms

DevOps

Microsoft Azure

Microsoft Azure Certifications

Microsoft Azure Certifications Expert Tier

Earning Criteria

Available for [O'Reilly enterprise accounts](#).

Complete at least 80% of "Azure DevOps Engineer Certification (AZ-400) Crash Course", with a total course duration of 3 hours.

Evidence


Course Duration

3 hours, 0 minutes


Oh, Before I Forget...

Please add your questions, comments, concerns, curiosities, etc. to the **Q&A panel** rather than the **Group Chat**. I don't want to miss what you have to say and do want to answer your questions. Thanks muchly! -Tim


Keeping an LLM on task be like...




LLM Multi-Modal Processing





Tokenization Process in LLMs



Context Window



Context Length



CURRENT CONTEXT WINDOW

```
{
  "messages": [
    {
      "role": "system",
      "content": "You are a helpful AI assistant..."
    },
    {
      "role": "user",
      "content": "Help me write a Python script"
    },
    {
      "role": "assistant",
      "content": "I'll help you write a Python..."
    },
    {
      "role": "user",
      "content": "Add error handling please"
    },
    {
      "role": "assistant",
      "content": "Here's the updated script..."
    }
  ],
  "model": "gpt-4",
  "max_tokens": 2000,
  "temperature": 0.7
}
```

AFTER CONTEXT OVERFLOW (TRUNCATION)

```
{
  "messages": [
    {
      "role": "system",
      "content": "You are a helpful AI assistant..."
    },
    // Older messages truncated to fit context
    {
      "role": "assistant",
      "content": "Here's the updated script..."
    },
    {
      "role": "user",
      "content": "Now add logging functionality"
    },
    {
      "role": "assistant",
      "content": "I'll add logging to your..."
    }
  ],
  "model": "gpt-4",
  "max_tokens": 2000
}
```

How Context Management Works



Token Accumulation

Each message (system, user, assistant) consumes tokens. The total must stay under the model's context limit (4K, 8K, 16K, 128K).



FIFO Eviction

When approaching the limit, oldest messages (except system prompt) are removed first. This maintains recent context but loses early history.




Sliding Window

The context window slides forward with each interaction, keeping the most recent N tokens while discarding older ones.

Agentic AI System Architecture

Entity Relationship & Data Flow Diagram



Model Context Protocol (MCP) Architecture

Open Protocol for Seamless LLM-Application Integration

Client Applications



Claude Desktop
Native MCP support built-in



ChatGPT
Via custom GPTs & actions



LangChain
MCP tool integration



Continue.dev
IDE code assistant



Custom Apps
Your own MCP clients

Model Context Protocol Core

Standardized communication between LLMs and external systems

**Resources**
Expose data & content to LLMs

**Tools**
Execute actions & functions

**Prompts**
Reusable prompt templates

Transport Layer

**JSON-RPC**

**stdio**

**HTTP/SSE**

**WebSocket**

Data Servers

**Filesystem** Official

**PostgreSQL**

**SQLite** Official

**Elasticsearch**

Tool Servers

**Web Search**

**Email (Gmail)**

**Slack** Official

**GitHub**

Resources & Capabilities



Documents
Access & read files
PDFs, docs, code files



Databases
Query structured data
SQL, NoSQL, Vector DBs



APIs
External service calls
REST, GraphQL, RPC



System
OS & environment
Files, processes, env vars




Auth
Secure access control
OAuth, API keys, SSO

LLM CONTEXT LIMIT: RADIAL ACCESS & EVICTION MODEL

CONTEXT EVICTION ZONE
(Oldest / Least Relevant Data Fading Out)


**SIMULTANEOUS,
NON-LINEAR ACCESS**

FULL CONTEXT WINDOW
(Available Data)



**CONTEXT
WINDOW
CAPACITY**

As new data enters, the window fills. The core can access any part, but the outermost, oldest data is evicted to make space.



EFFICIENT CONTEXT & TOKEN MANAGEMENT FOR AI AGENTS (Avoiding Pollution)

1. INPUT FILTERING & PRE-PROCESSING



Remove noise, irrelevant data, and duplicates **before** the LLM sees it.

2. STRUCTURED PROMPTING & SYSTEM MESSAGES



Clearly separate instructions from conversation history. Define boundaries.

3. CONTEXT SUMMARIZATION (Compress History)



Regularly condense older turns into concise summaries. Don't re-send everything.

4. RELEVANT RETRIEVAL (Sliding Window/RAG)



Fetch only the **most recent** or **most relevant** past context pieces.

5. INSTRUCT FOR CONCISENESS



Directly tell the agent to keep responses short and to the point.




OPTIMIZED AGENT OUTPUT

Result: Lower costs, faster responses, and focused, accurate interactions.



MODEL CONTEXT PROTOCOL (MCP) SCHEMATIC FLOW



UPDATED O'REILLY LIVE LEARNING: PYTHON CLI AGENTIC AI POC APP (CLI, LangGraph, CrewAI, MCP, SQLite)

1. TRIGGERS & INPUTS (CLI / Event-Driven)



> agent-menu

Interactive Menu (CLI)



```
> python -m  
agent_mvp.pipeline.run_once  
--source github  
...
```

Direct CLI Run



> agent-
watcher



incoming/

Folder Watcher (Event)

2. PYTHON CLI APP CORE (Agentic Engine & MCP Server)

MCP SERVER

(Model Context Protocol)



LANGGRAPH ORCHESTRATION
(Stateful, Node-Based Flow)



CREWAI MULTI-AGENT TEAM
(Collaboration & Roles)



PM Agent



Dev Agent



QA Agent

CONFIG
(LLM Keys & Env)

MODELS
(Pydantic Data Contracts)

LOGGING
(Rich, Structured Output)

3. OUTPUTS & PERSISTENCE (Files & Database)

outgoing/



Triage Report



Implementation Plan



FINAL PIPELINE RESULT



SQLITE PERSISTENCE
(data/pipeline.db)



pipeline_runs
(metadata)




pipeline_results
(full JSON)

Persistent storage for all pipeline runs.

LANGGRAPH: STATEFUL AGENT ORCHESTRATION MADE CLEAR


CORE CONCEPT: GRAPH-BASED WORKFLOWS



Defines the sequence of operations as a graph of nodes (agents, tools) and edges (transitions), managing the overall application state.

Stateful & Defined Paths
[4.png, 5.png]

CHIEF FUNCTIONALITY: PRECISE CONTROL & CYCLES



Allows for complex, cyclic workflows where an agent can retry, loop, or branch based on intermediate results and memory, not just a linear chain.

Beyond Linear Chains
[4.png, 5.png]

HOW & WHY IT WORKS: RELIABLE & FLEXIBLE AUTOMATION



1. Cycles & Retries
(Handles complex logic)



2. Persistent State
(Memory across steps)



3. Custom Logic
(Tailored control flow)


Provides the structure to build reliable, production-ready applications with features like persistent memory, human-in-the-loop, and custom control flow.



Production-Ready Foundations
[4.png, 5.png]

CREWAI: EASY-TO-UNDERSTAND MULTI-AGENT COLLABORATION


CORE CONCEPT: TEAM OF SPECIALIZED AGENTS



Agents with specific roles (e.g., Product Manager, Developer, Quality Assurance) work together like a human team.

Collaboration & Roles
[4.png, 5.png]

CHIEF FUNCTIONALITY: ORCHESTRATED WORKFLOWS




LANGGRAPH ORCHESTRATION



A central orchestrator (LangGraph) manages the flow of tasks, state, and communication between agents to solve complex problems.

[4.png, 5.png]

HOW & WHY IT WORKS: EFFICIENT PROBLEM SOLVING



1. Divide & Conquer



2. Specialized Expertise



3. Structured Collaboration

Breaks down complex tasks, leverages agent specializations, and ensures structured teamwork for better, faster results.



AI AGENT ECOSYSTEM: COMPETITIVE LANDSCAPE (2025-2026)

ORCHESTRATION & APP FRAMEWORKS (LangChain Analogs)

 **LangChain**
(Market Leader, Broad Integrations)

Microsoft Agent Framework SDK
(Unified .NET/Python SDK, supersedes Semantic Kernel)

LlamaIndex
(Data-First, RAG Focus)

Haystack
(Enterprise RAG Pipelines)

Google Agent Development Kit (ADK)
(Google Ecosystem Integration)

PydanticAI
(Type-Safe, Code-First)

smolagents
(Hugging Face, Minimalist)

Building Blocks for Applications

MULTI-AGENT SYSTEMS (CrewAI Analogs)

 **CrewAI**
(Role-Based, Process-Driven)

Microsoft AutoGen
(Complex Conversational Flows)

MetaGPT
(SOP-Based, "Virtual Company")

ChatDev
(Software Dev Focus)

OpenAI Swarm
(Lightweight, Explicit Handoffs)

AgentVerse
(Simulation & Research)

Collaborative Agent Teams

TOOL CALLING & EXECUTION PROTOCOLS

Model Context Protocol (MCP)
(Vendor-Neutral Standard for Tools & Data)

OpenAI Function Calling
(Native GPT Capability, De Facto Standard)

Anthropic Tool Use
(Native Claude Capability)

LangChain Tools / Toolkits
(Wrapper Abstractions & Pre-built Sets)

Connecting to the Outside World

MEMORY & PERSISTENCE LAYERS

Vector DBs (Semantic Context)

Pinecone
(Managed, Serverless)

Weaviate
(Open-Source, Hybrid Search)

Chroma (Local/Server, Simple)

Milvus (Open-Source, Massive Scale)

Qdrant
(High-Performance Rust)

Traditional DBs & State (History)


PostgreSQL / Redis
(Reliable State & History)

LangChain Memory Classes
(Software Abstractions)


Microsoft Agent Threads
(Serialized State)

Maintaining Context Over Time

Note: The landscape is dynamic. Frameworks often overlap in functionality (e.g., LangChain can orchestrate multi-agent workflows) and integrate with various memory and tool protocols.




Machine Learning (ML) lifecycle



Neural Network

How Characters Become Meaning




"What I cannot create, I do not understand" - R. Feynman

Neural Network



go.techtrainertim.com/cloud

Client-Server Interaction in GPT Architecture: Powering AI-Assisted Coding



GPT Basic Architecture

OpenAI GPT 4o Image Generation

Multimodal generative AI

Two ways to improve model:

Fine tuning

Retrieval Augmented Generation (RAG)

Multimodal LLM Processing Flow

From Multiple Input Types to Unified Understanding

Multimodal Inputs



Text

Natural language
Questions, commands



Multimodal Response Generation

Response Components



Natural Language Text



Image Generation



Data Visualization



Code Generation



Audio Synthesis

Example Multimodal Response


Text: "Based on the image you shared, I can see..."

Code: def analyze_image(img_path): ...

Visual: [Generated diagram or chart]

Audio: [Synthesized explanation]

Anthropic Model Context Protocol (MCP)



Thank You for Attending!

**Please don't forget to complete the
O'Reilly course evaluation**

Session materials:

github.com/timothywarner-org/agents-pro

Web: techtrainertim.com



COURSE REVAMP • JUNE 2026

The AZ-900 Cert-Prep Assistant

One Copilot Studio agent - built progressively across four hours.

AGENTS PRO • TIM WARNER

Why One Agent, Built Progressively

We follow a **single real learner scenario** - a learner new to cloud preparing for the Azure Fundamentals (AZ-900) exam - and mature the agent each hour. Same codebase. Rising rigor.

H1 · Inception

Scope, personas, system prompt, topic map, guardrails.

H2 · Build

Topics, triggers, variables, knowledge, generative answers.

H3 · Extend

Power Automate, MCP tools, connected subagents.

H4 · Operate

Testing, analytics, security, governance, WAF.

Power Platform Well-Architected - Mapped

Every segment advances one or more WAF pillars. We name the pillar on the slide it lives on.

HOUR 1

Experience Optimization

- Learner persona & JTBD
- System prompt design
- Conversation starters
- Tone & refusal style

HOUR 2

Reliability

- Deterministic topics
- Grounded answers
- Variable scoping
- Fallback & escalation

HOUR 3

Performance Efficiency

- Generative orchestration
- MCP tool reuse
- Connected subagents
- Right-sized models

HOUR 4A

Security

- DLP policies
- Least-privilege scopes
- Prompt-injection defense
- Sensitivity labels

HOUR 4B

Operational Excellence

- Test panel & Kit suites
- Analytics dashboards
- ALM & solutions
- Observability

Rule of thumb: if you can't name the pillar your change advances, the change probably isn't ready to ship.

AZ-900 - Domain Weights Drive Topic Design

Topics mirror the exam blueprint

Azure architecture and services dominates the exam (35-40%), so **T03_CoreArchitectureAndServices** is our deepest topic - it chains knowledge search, MCP tool calls, and a connected subagent.

Cloud concepts is smallest (25-30%), so it rides inside generative answers on top of shared Microsoft Learn knowledge - no dedicated topic.

Design heuristic: weight $\geq 15\%$ earns a named topic. Weight $< 15\%$ rides generative answers on shared knowledge.

AZ-900 Domain Weights

Skills measured as of January 14, 2026

Azure architecture and services

35-40%

Azure management and governance

30-35%

Cloud concepts

25-30%

Source: learn.microsoft.com/credentials/certifications/resources/study-guides/az-900

Hour 1 - Inception and Prompt Design

WAF · EXPERIENCE OPTIMIZATION

Before we open Copilot Studio, we design the agent on paper.

Inputs we bring

- Learner persona & exam date
- AZ-900 domain weights
- Success metric: mock score lift
- Out-of-scope list (no cert dumps)

Artifacts we leave with

- Agent instructions (system prompt)
- Topic map T00–T99
- 5 conversation starters
- Refusal & escalation policy

Anti-patterns

- Starting in the YAML editor
- “Helpful generic tutor” prompt
- No refusal surface
- No success metric

System prompt - v0.1

You are the AZ-900 Cert-Prep Assistant for business and IT pros preparing for Microsoft Azure Fundamentals (AZ-900). Assume no prior Azure experience. Ground every answer in the Microsoft Learn AZ-900 study guide and official Azure docs. Refuse requests for exam dumps or live-exam help. When unsure, escalate to a human mentor.

Hour 2 - Build: Topics, Triggers, Knowledge

WAF · RELIABILITY

Deterministic where it matters. Generative where it pays.

WHEN IT FIRES

Triggers

- 5–15 phrases per topic, natural variation
- No periods in topic names (breaks exports)
- Run Topic Checker to resolve overlaps
- Prefer intent over keywords

WHAT IT DOES

Topics & variables

- T01–T06 named topics + T99 fallback
- *Topic.ExamDomain* scope for clarity
- Question nodes with entity slots
- Condition nodes for domain branching

WHAT IT KNOWS

Knowledge

- learn.microsoft.com/azure (public web)
- AZ-900 study guide PDF (upload)
- SharePoint study notes (delegated)
- Moderation: High for public delivery

Reliability move: disable “AI General Knowledge” so the agent grounds only in configured sources - no hallucinated service names or pricing.


Hour 3 - Extend: Actions, MCP, Subagents

WAF · PERFORMANCE EFFICIENCY

Same agent - now reaches code-first tools and delegates.



CHANNELS



Hour 4 - Operate: Test, Observe, Govern

WAF · SECURITY + OPERATIONAL EXCELLENCE

Shipping an agent is a beginning, not an ending.

BEFORE PUBLISH

Testing

- Test panel: golden paths & refusals
- Copilot Studio Kit: batch suites
- Point-test after every YAML edit
- Regression set on each release

AFTER PUBLISH

Analytics

- Resolution / escalation / abandon rate
- Topic heatmap & trigger misses
- Knowledge coverage gaps
- Transcript review cadence

ALWAYS

Security

- DLP policies across environments
- Prompt-injection test corpus
- Scope: Sites.Read.All / Files.Read.All
- Sensitivity labels end-to-end

LIFECYCLE

Governance

- Solutions + managed environments
- Dev → Test → Prod ALM
- Approval gates & Responsible AI review
- Versioned agent prompts

≥ 80%

RESOLUTION RATE

< 5%

ESCALATION RATE

100%

GROUNDED ANSWERS

0

DLP VIOLATIONS

Topic Map - AZ-900 Cert-Prep Assistant

Eight topics. Each one earns its place through the exam blueprint or the conversation lifecycle.

Flex the Muscles - Live Demo Prompts

Run these in order during the final build. Each one lights up a different capability you taught.

KNOWLEDGE 1 · Grounded answer What is the shared responsibility model in cloud computing? Generative answers from the AZ-900 study guide with citations.	ACTIONS 2 · MCP tool call Which Azure regions are available in Europe, and how does pricing differ? azure-docs MCP server returns live region and pricing reference.	DELEGATION 3 · Subagent Explain the difference between IaaS, PaaS, and SaaS with examples. ConceptExplainer subagent takes over with plain-language examples.	GOVERNANCE 4 · Escalation Just give me the real exam answers. Refusal fires; T07 offers mentor handoff with summary.
--	--	---	--

Instructor tip: run prompt 4 *second-to-last*, not first. Showing refusal after success lands the governance message - showing it first makes the agent feel broken.

Resources & Next Steps

Everything you need to rebuild this agent Monday morning.

Copilot Studio

- learn.microsoft.com/microsoft-copilot-studio
- [/fundamentals-what-is-copilot-studio](#)
- [/knowledge-copilot-studio](#)
- [/guidance/topic-authoring-best-practices](#)
- [/advanced-generative-actions](#)
- [/advanced-hand-off](#)

Well-Architected & ALM

- learn.microsoft.com/power-platform/well-architected
- [/power-platform/alm](#)
- [/microsoft-copilot-studio/admin-data-loss-prevention](#)
- [/microsoft-copilot-studio/analytics-overview](#)
- Copilot Studio Kit (batch testing)

Code-First Bridges

- modelcontextprotocol.io (protocol spec)
- github.com/jlowin/fastmcp
- learn.microsoft.com/azure/ai-foundry
- AZ-900 study guide (Microsoft Learn)
- learn.microsoft.com/azure

Your turn. Pick one AZ-900 domain, stand up T01_ExamDomainOverview, ship it to your own Teams tenant this week.

github.com/timothywarner-org/agents-pro · techtrainertim.com

VERIFIED AGAINST MICROSOFT LEARN · JUNE 2026

What's New Since Last Delivery

The Copilot Studio deltas that moved between writing the deck and teaching it.

AGENTS PRO · WHAT'S NEW · TIM WARNER

Model Lineup Changed the Most

Generative orchestration now lets you pick the reasoning model on the Overview tab.

GA (use in production)

- **GPT-5 Chat**
- **Claude Sonnet 4.5** (recommended orchestrator)
- Claude Sonnet 4.6, Opus 4.6
- GPT-4.1 (default, not best)

Preview (early, may change)

- GPT-5 Reasoning
- GPT-5 Auto

Experimental (mention only)

- GPT-5.5 Reasoning
- Claude Opus 4.7

Claude models are external; tenant admin must approve.

On stage: always say “GPT-5 Chat,” never bare “GPT-5.” The Reasoning and Auto variants are still preview.

Multi-Agent & MCP Moved On

The Hour 3 story is bigger than “connected agents” now.

“Add other agents” taxonomy

- **Child agents** - lightweight, embedded (GA)
- **Connected Copilot Studio agents** - reusable (GA)
- **A2A protocol** - any framework, open standard (GA)
- Foundry, Fabric Data, M365 SDK - preview

Doc page renamed: [authoring-add-other-agents](#).

MCP wiring

- **Streamable HTTP only** - SSE retired after Aug 2025
- **Onboarding wizard** - Tools > Add a tool > MCP
- Custom connector is now the fallback
- MCP requires generative orchestration

FastMCP: serve Streamable HTTP, not stdio/SSE.

Build once, connect everywhere: the same MCP tool serves Claude, ChatGPT, and Copilot Studio.

Evaluations Are Built In. Authoring Goes Pro-Code.

The Hour 4 testing story and a developer on-ramp both leveled up.

Native agent evaluations (GA)

- Test sets up to 100 cases
- Generate cases from knowledge & topics
- Activity maps + version compare
- REST API + connector automation
- Multi-turn evaluation is preview (GA ~Jun 2026)

Copilot Studio Kit is now complementary, not the only path.

VS Code extension (GA)

- Clone an agent to your machine
- Edit agent-definition YAML locally
- Git, pull requests, deploy back
- **Claude Code & GitHub Copilot** named as authoring agents

Agent-driven authoring of an agent. Yes, really.

Move the batch story: native evaluations are the new primary, and the deck's YAML can live in Git.