How to Create AI Agents Like a Pro

An interactive learning experience brought to you by **Pearson**

Presented by Tim Warner
TechTrainerTim.com

Microsoft MVP 2025 Most Valuable Professional

Pearson

Course Flow (Central Time)

09:00 AM	Segment 1 [Microsoft Agentic AI + Inception]
10:00 AM	Segment 2 [Building a Low-Code Agent]
11:00 AM	Segment 3 [Extending Our Agent's Capabilities]
12:00 PM	Segment 4 [Operating and Deploying Our Agent]
01:00 PM	Finish

Session Materials

github.com/timothywarner-org/agents-pro

Session Recording + Badge

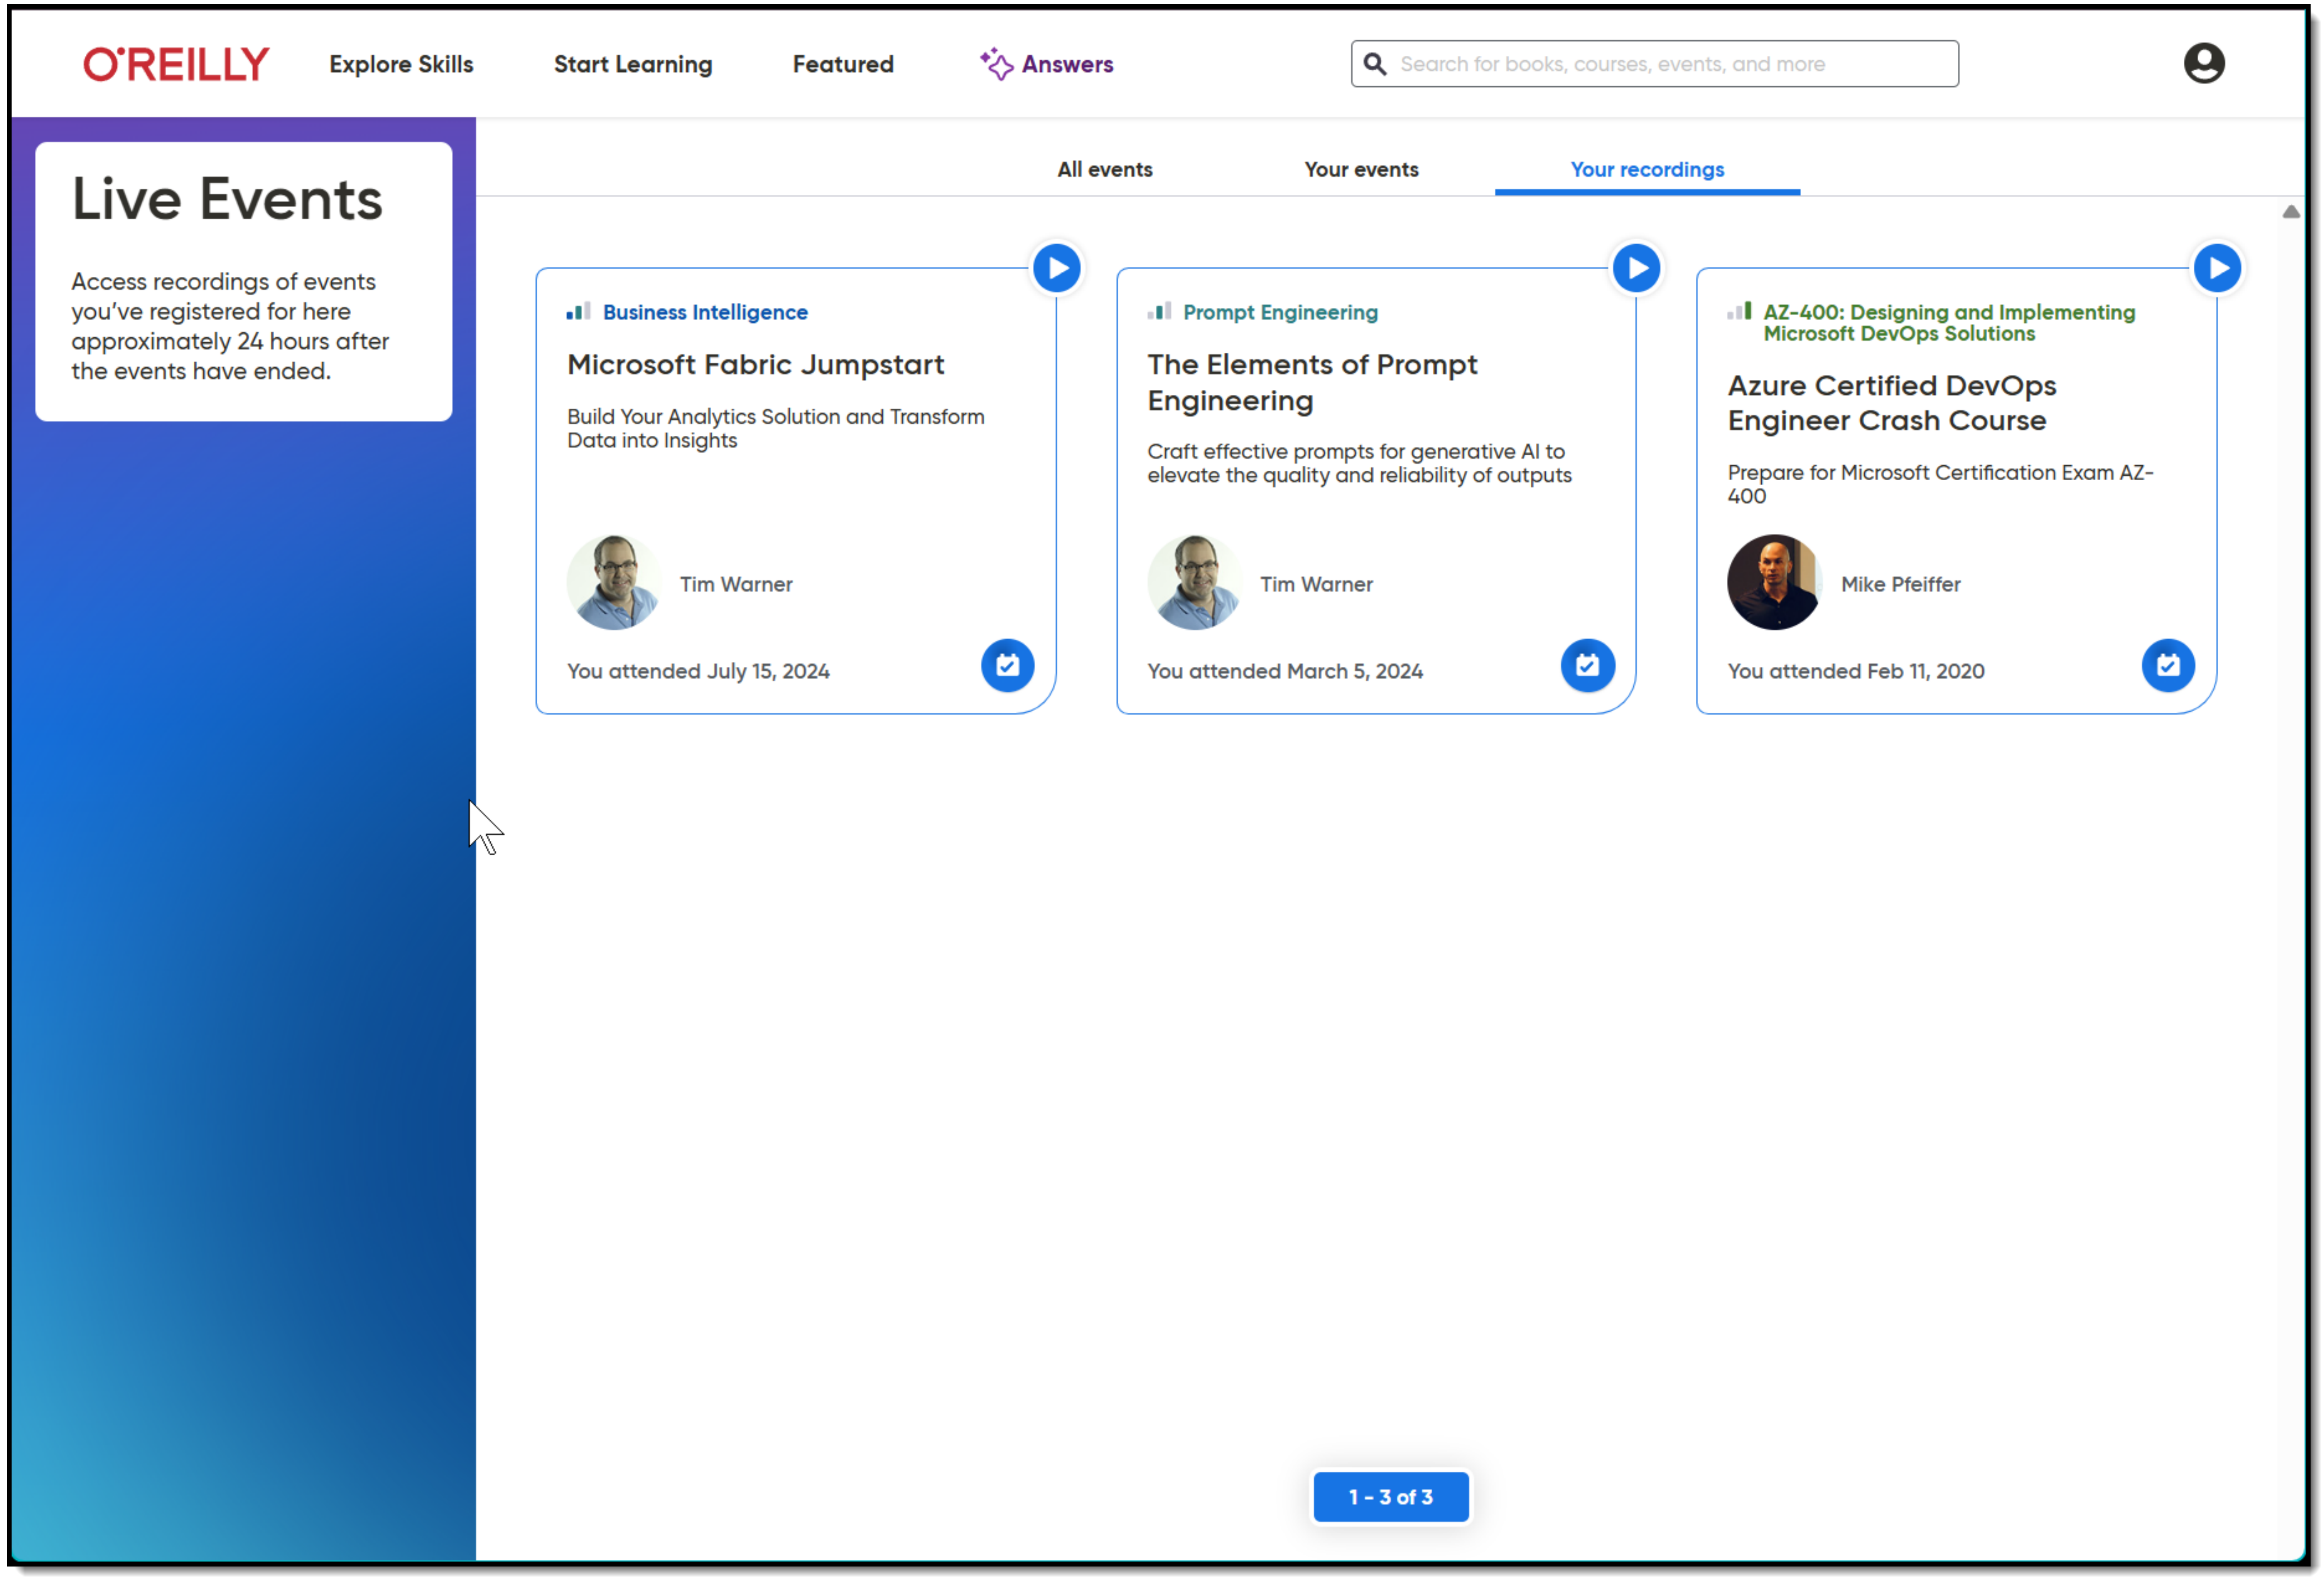

Session Recording + Badge

O'REILLY Explore Skills Start Learning Featured Answers Search for books, courses, events, and more

VIEW ALL EVENTS

Azure Certified DevOps Engineer Crash Course

Published by O'Reilly Media, Inc.

Advanced

Prepare for Microsoft Certification Exam AZ-400

Congrats!

Completed on Nov 21, 2021

What you'll learn Is this live event for you? Schedule

In this online course, Microsoft MVP Mike Pfeiffer walks you through what to expect on the AZ-400 Microsoft Azure DevOps Engineer Expert exam. This Azure certification will validate your Azure skills that combine people, process, and technologies to continuously deliver value to end-user needs and business objectives.

Even if you're not interested in certification, you will learn how DevOps professionals working in the Azure cloud streamline the delivery of applications and infrastructure. Walk away with a game plan on how to study for your exam, but also implement strategies for application and infrastructure code that allow for continuous integration, continuous testing, continuous delivery, and continuous monitoring and feedback.

What you'll learn and how you can apply it

By the end of this live, hands-on, online course, you'll understand:

- How the Azure role-based certifications work
- What topics are most likely to appear on the AZ-400 certification exam
- How to do the work of the Microsoft DevOps Engineer

And you'll be able to:

- Complete your studies to prepare for the AZ-400 certification exam
- Work towards Azure certification to prove your skills and increase your professional credibility
- Use the Azure portal and Azure PowerShell to deploy and manage resources in Azure

You attended

Feb 11, 2020

11am-3pm Central Standard Time

View session 1 video

Your Instructor

Mike Pfeiffer

Mike Pfeiffer is the founder and chief technologist at cloud consulting and training firm CloudSkills.io. A 20-year tech industry veteran, he's worked for some of the largest technology companies in the world, including Microsoft and Amazon Web Services (AWS). Mike is a publishe...

Read more

Skill covered

AZ-400: Designing and Implementing Microsoft DevOps Solutions

Earn your badge

Digital badges are verifiable and shareable proof of the skills you've built. Complete 80% or more of this course to earn your badge via Credly.

Session Recording + Badge

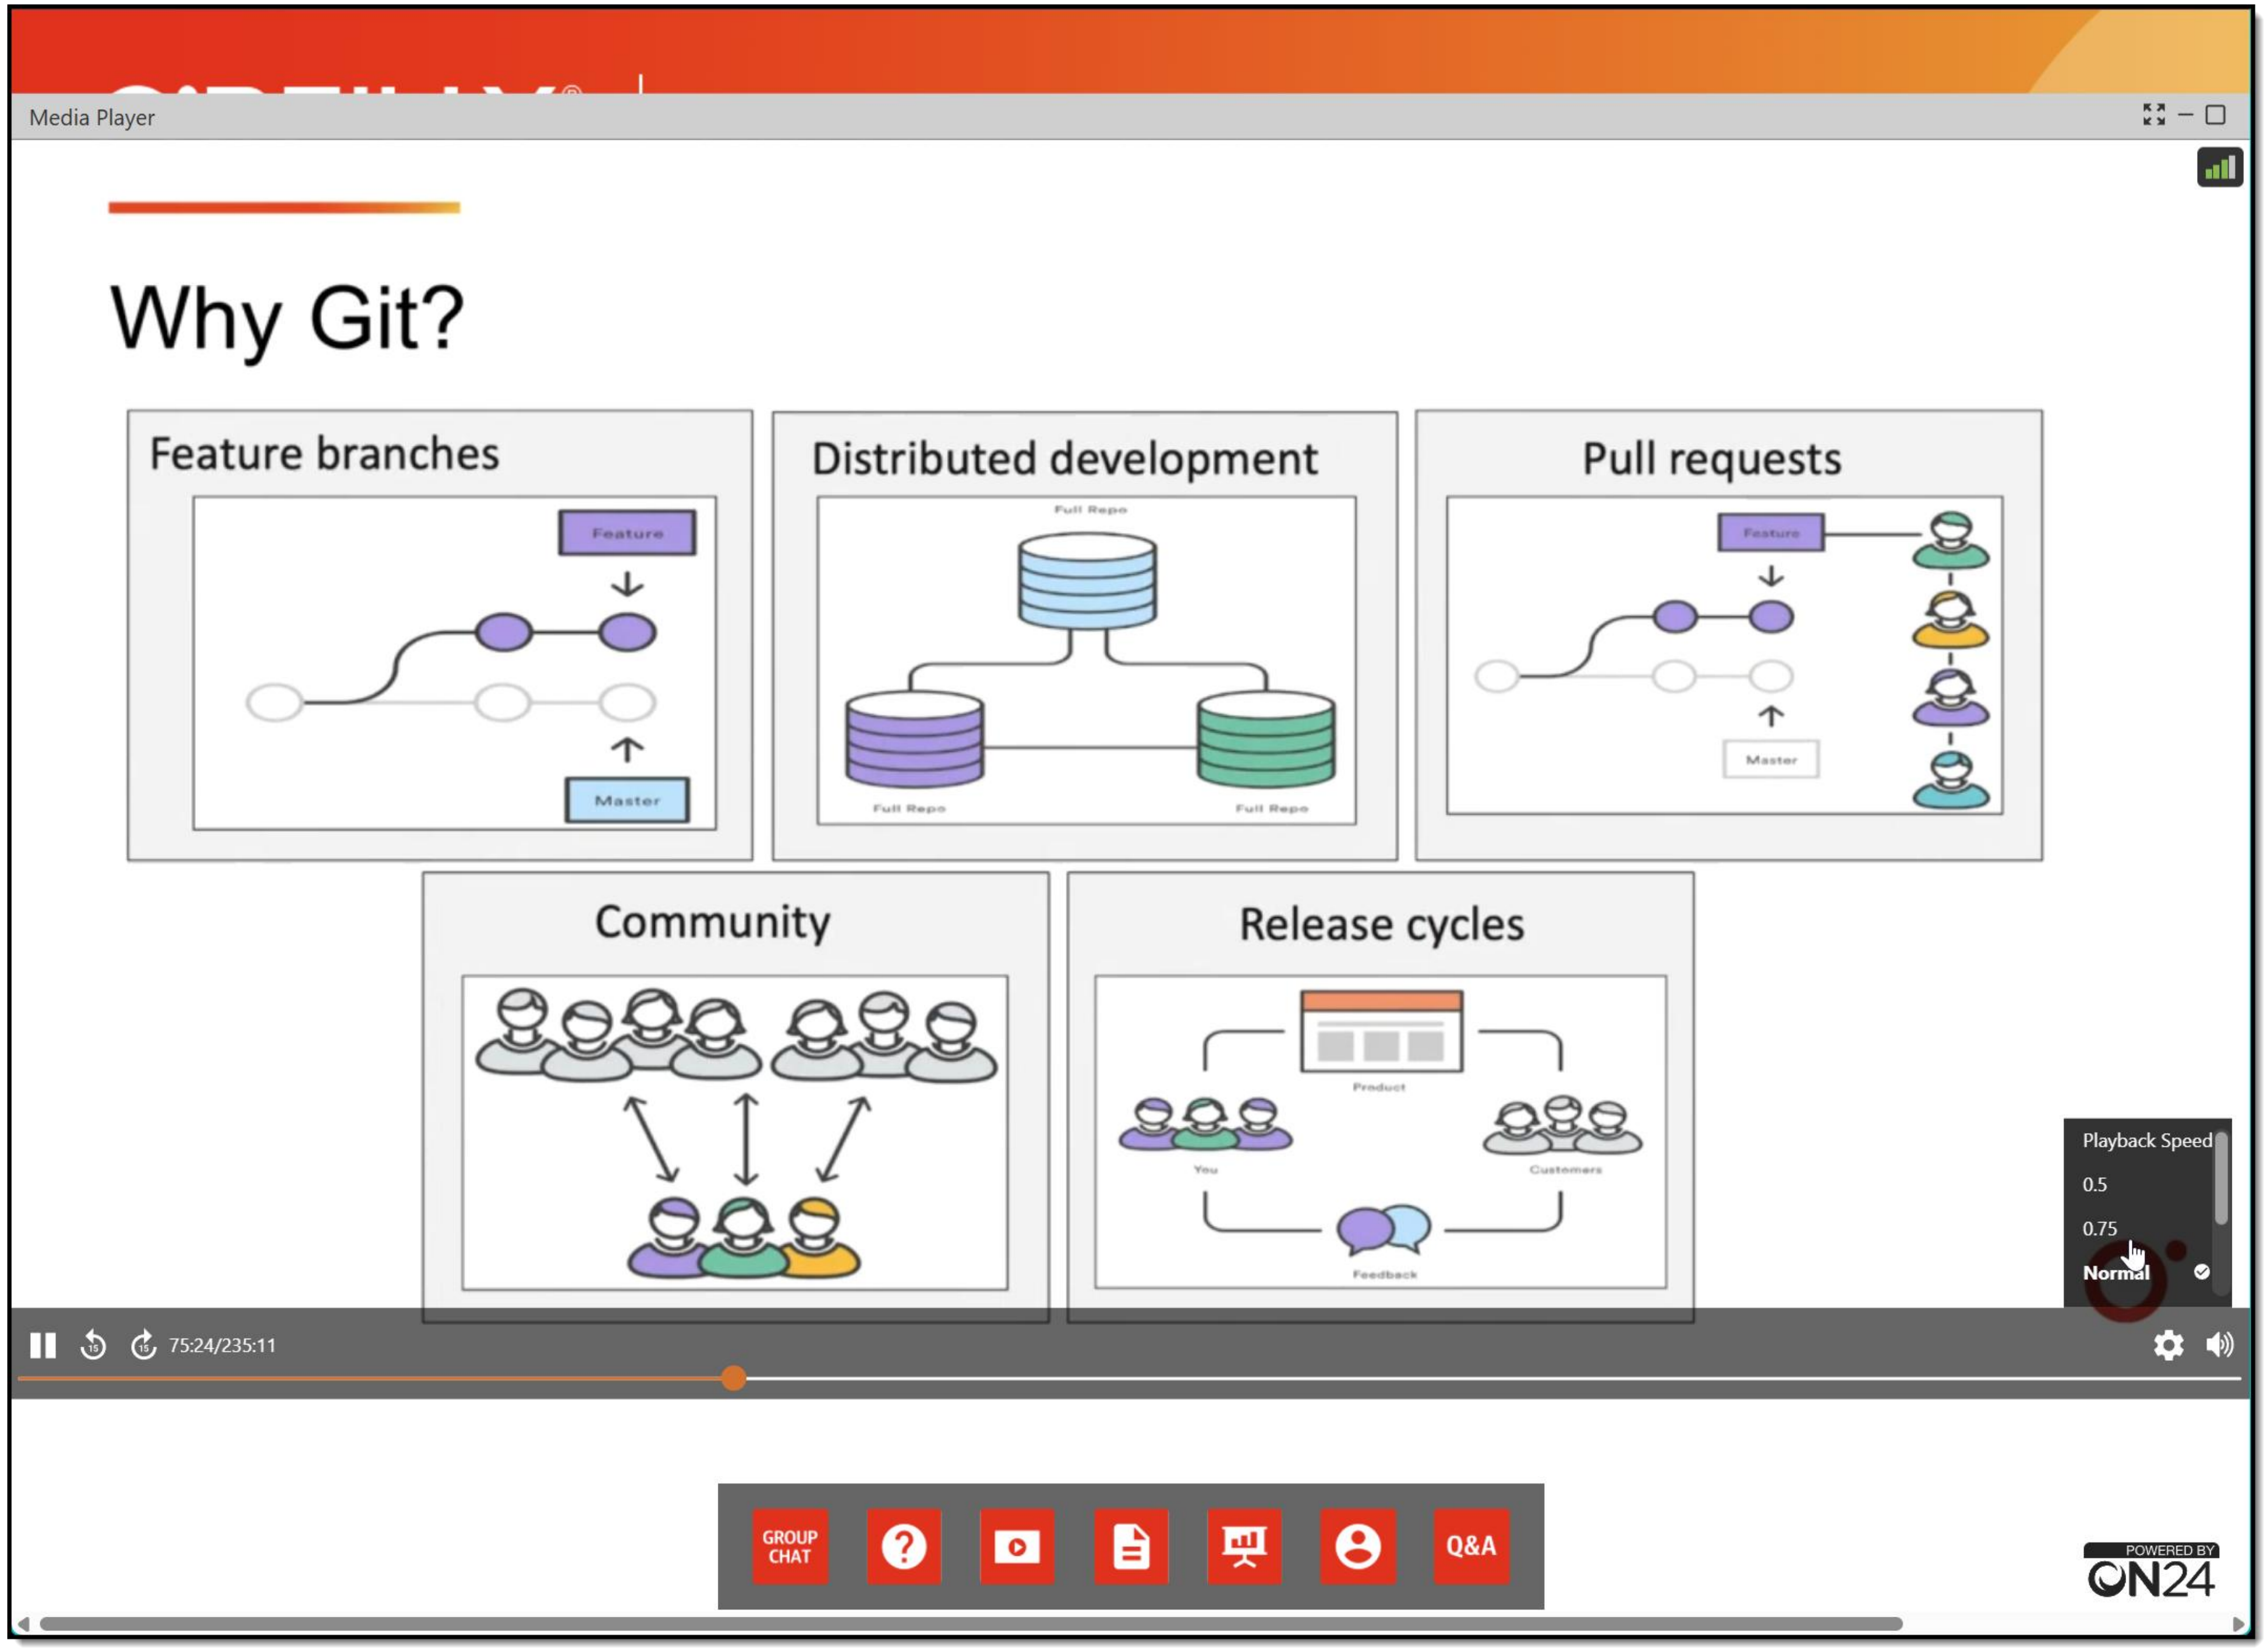

Session Recordings & Badge

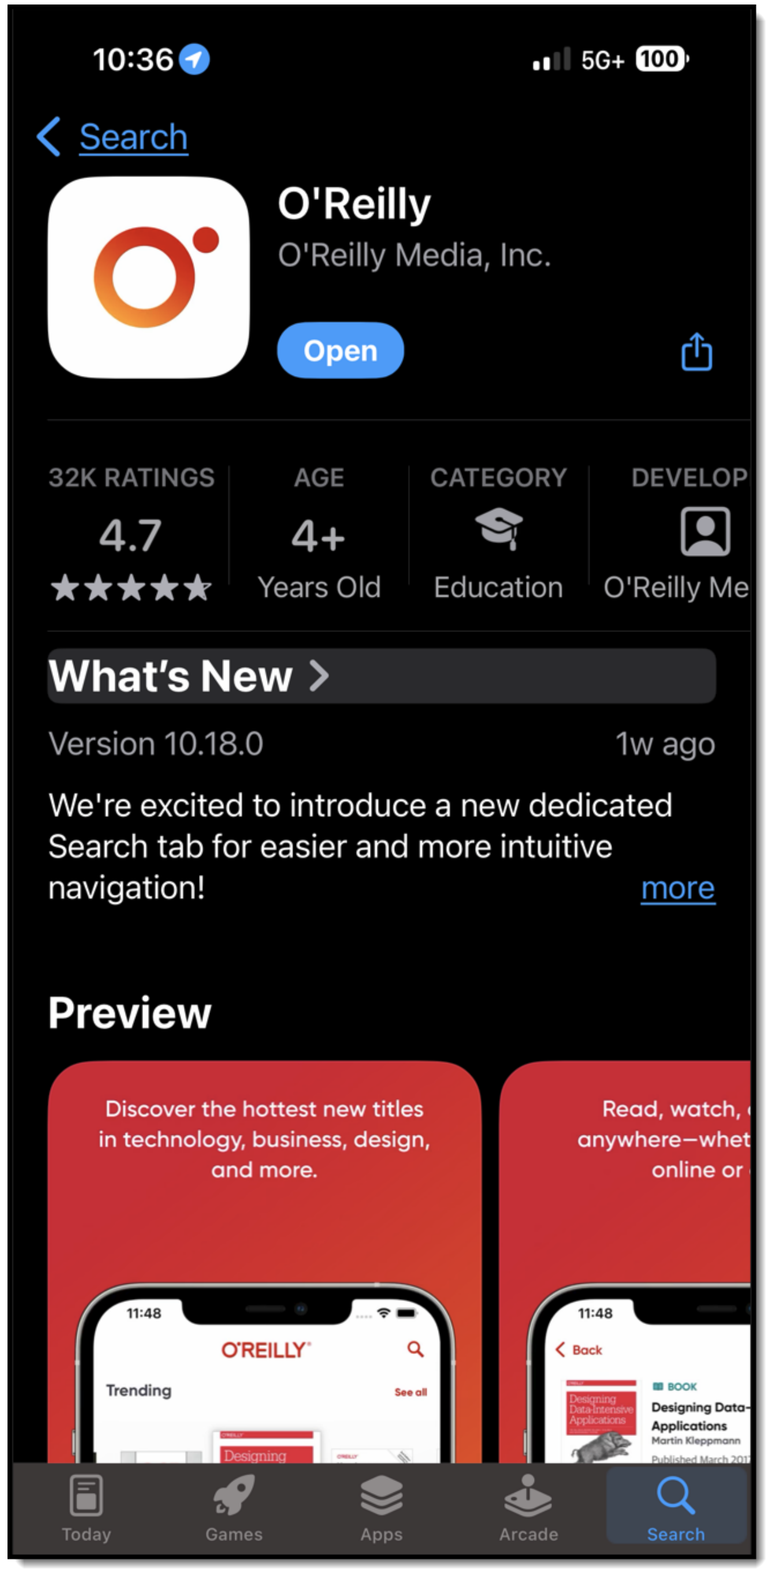

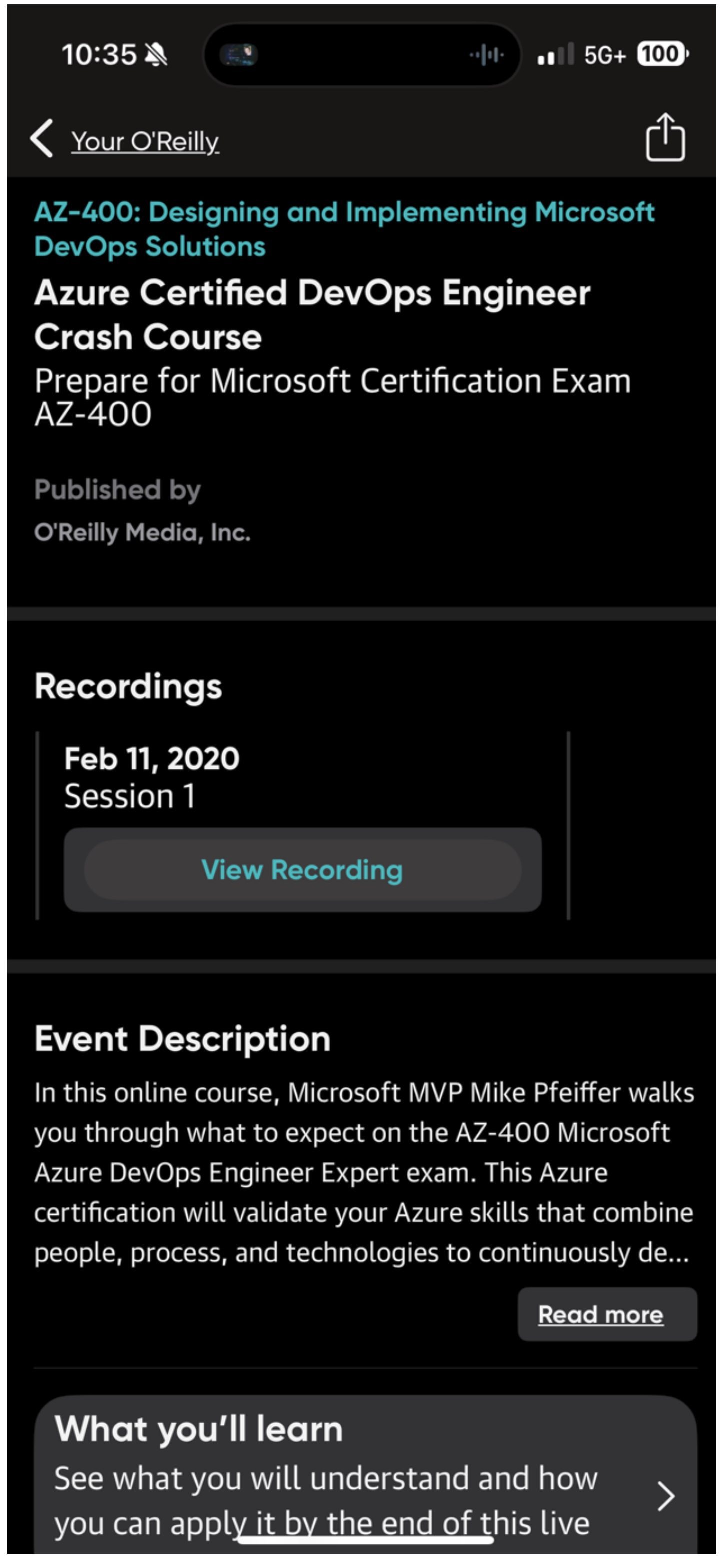

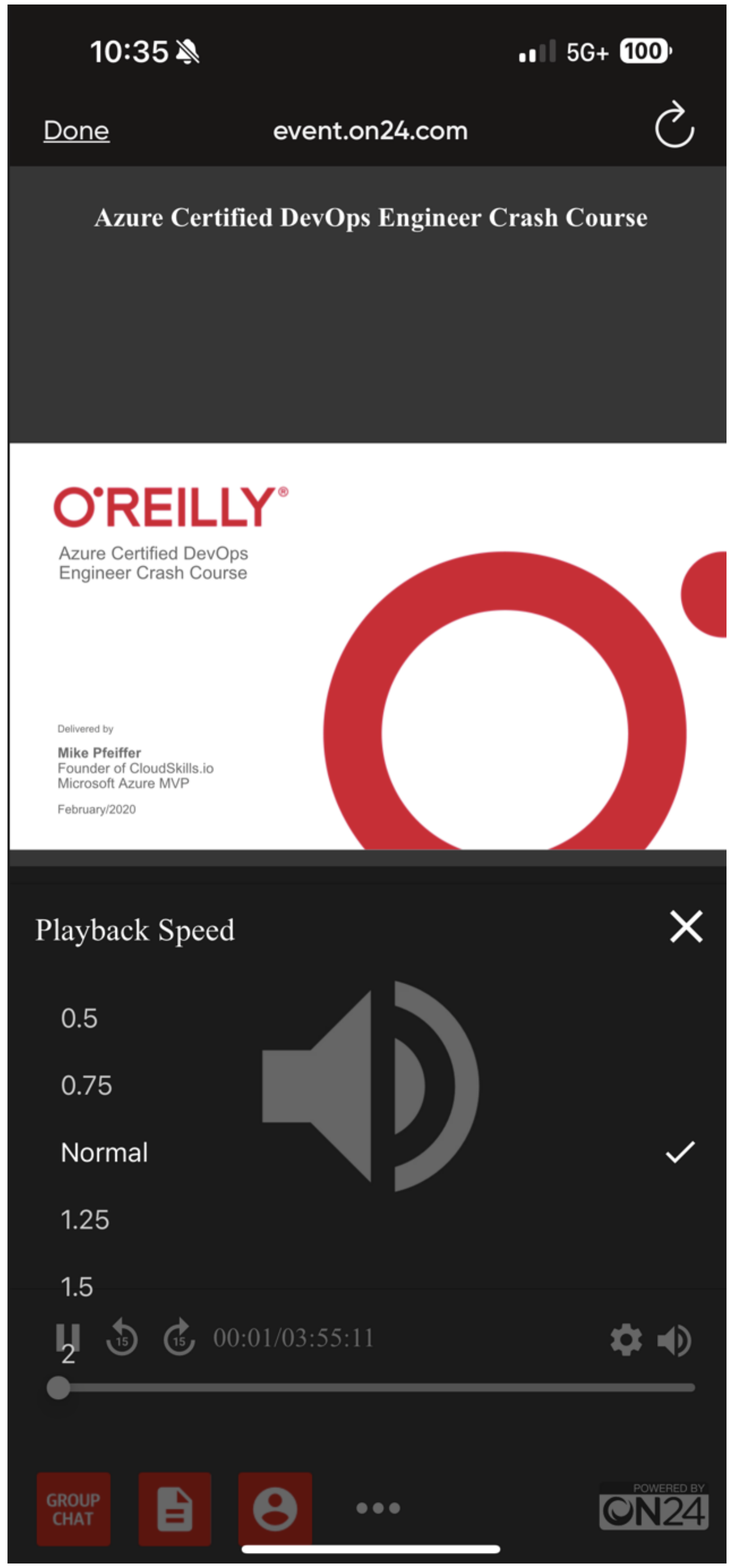

Session Recording + Badge

Credly by Pearson

Discover badges, skills or organizations

This badge was issued to Tim Warner | View celebrations
Date issued: April 13, 2025

Share

•••

AZ-400 Crash Course

Issued by O'Reilly Media

The holder of this badge has completed the following course by O'Reilly Media, Inc. The holder of this badge is capable of streamlining the delivery of applications and infrastructure in the Azure cloud, implementing strategies for continuous integration, continuous testing, continuous delivery, and continuous monitoring and feedback.

Learn more

Learning | Advanced | Hours | $ Paid

Skills

AZ-400: Designing and Implementing Microsoft DevOps... | Cloud Computing | Cloud Platforms | DevOps | Microsoft Azure | Microsoft Azure Certifications | Microsoft Azure Certifications Expert Tier

Earning Criteria

Available for O'Reilly enterprise accounts.

Complete at least 80% of "Azure DevOps Engineer Certification (AZ-400) Crash Course", with a total course duration of 3 hours.

Evidence

Course Duration
3 hours, 0 minutes

Oh, Before I Forget...

Please add your questions, comments, concerns, curiosities, etc. to the **Q&A panel** rather than the **Group Chat.** I don’t want to miss what you have to say and do want to answer your questions. Thanks muchly! -Tim

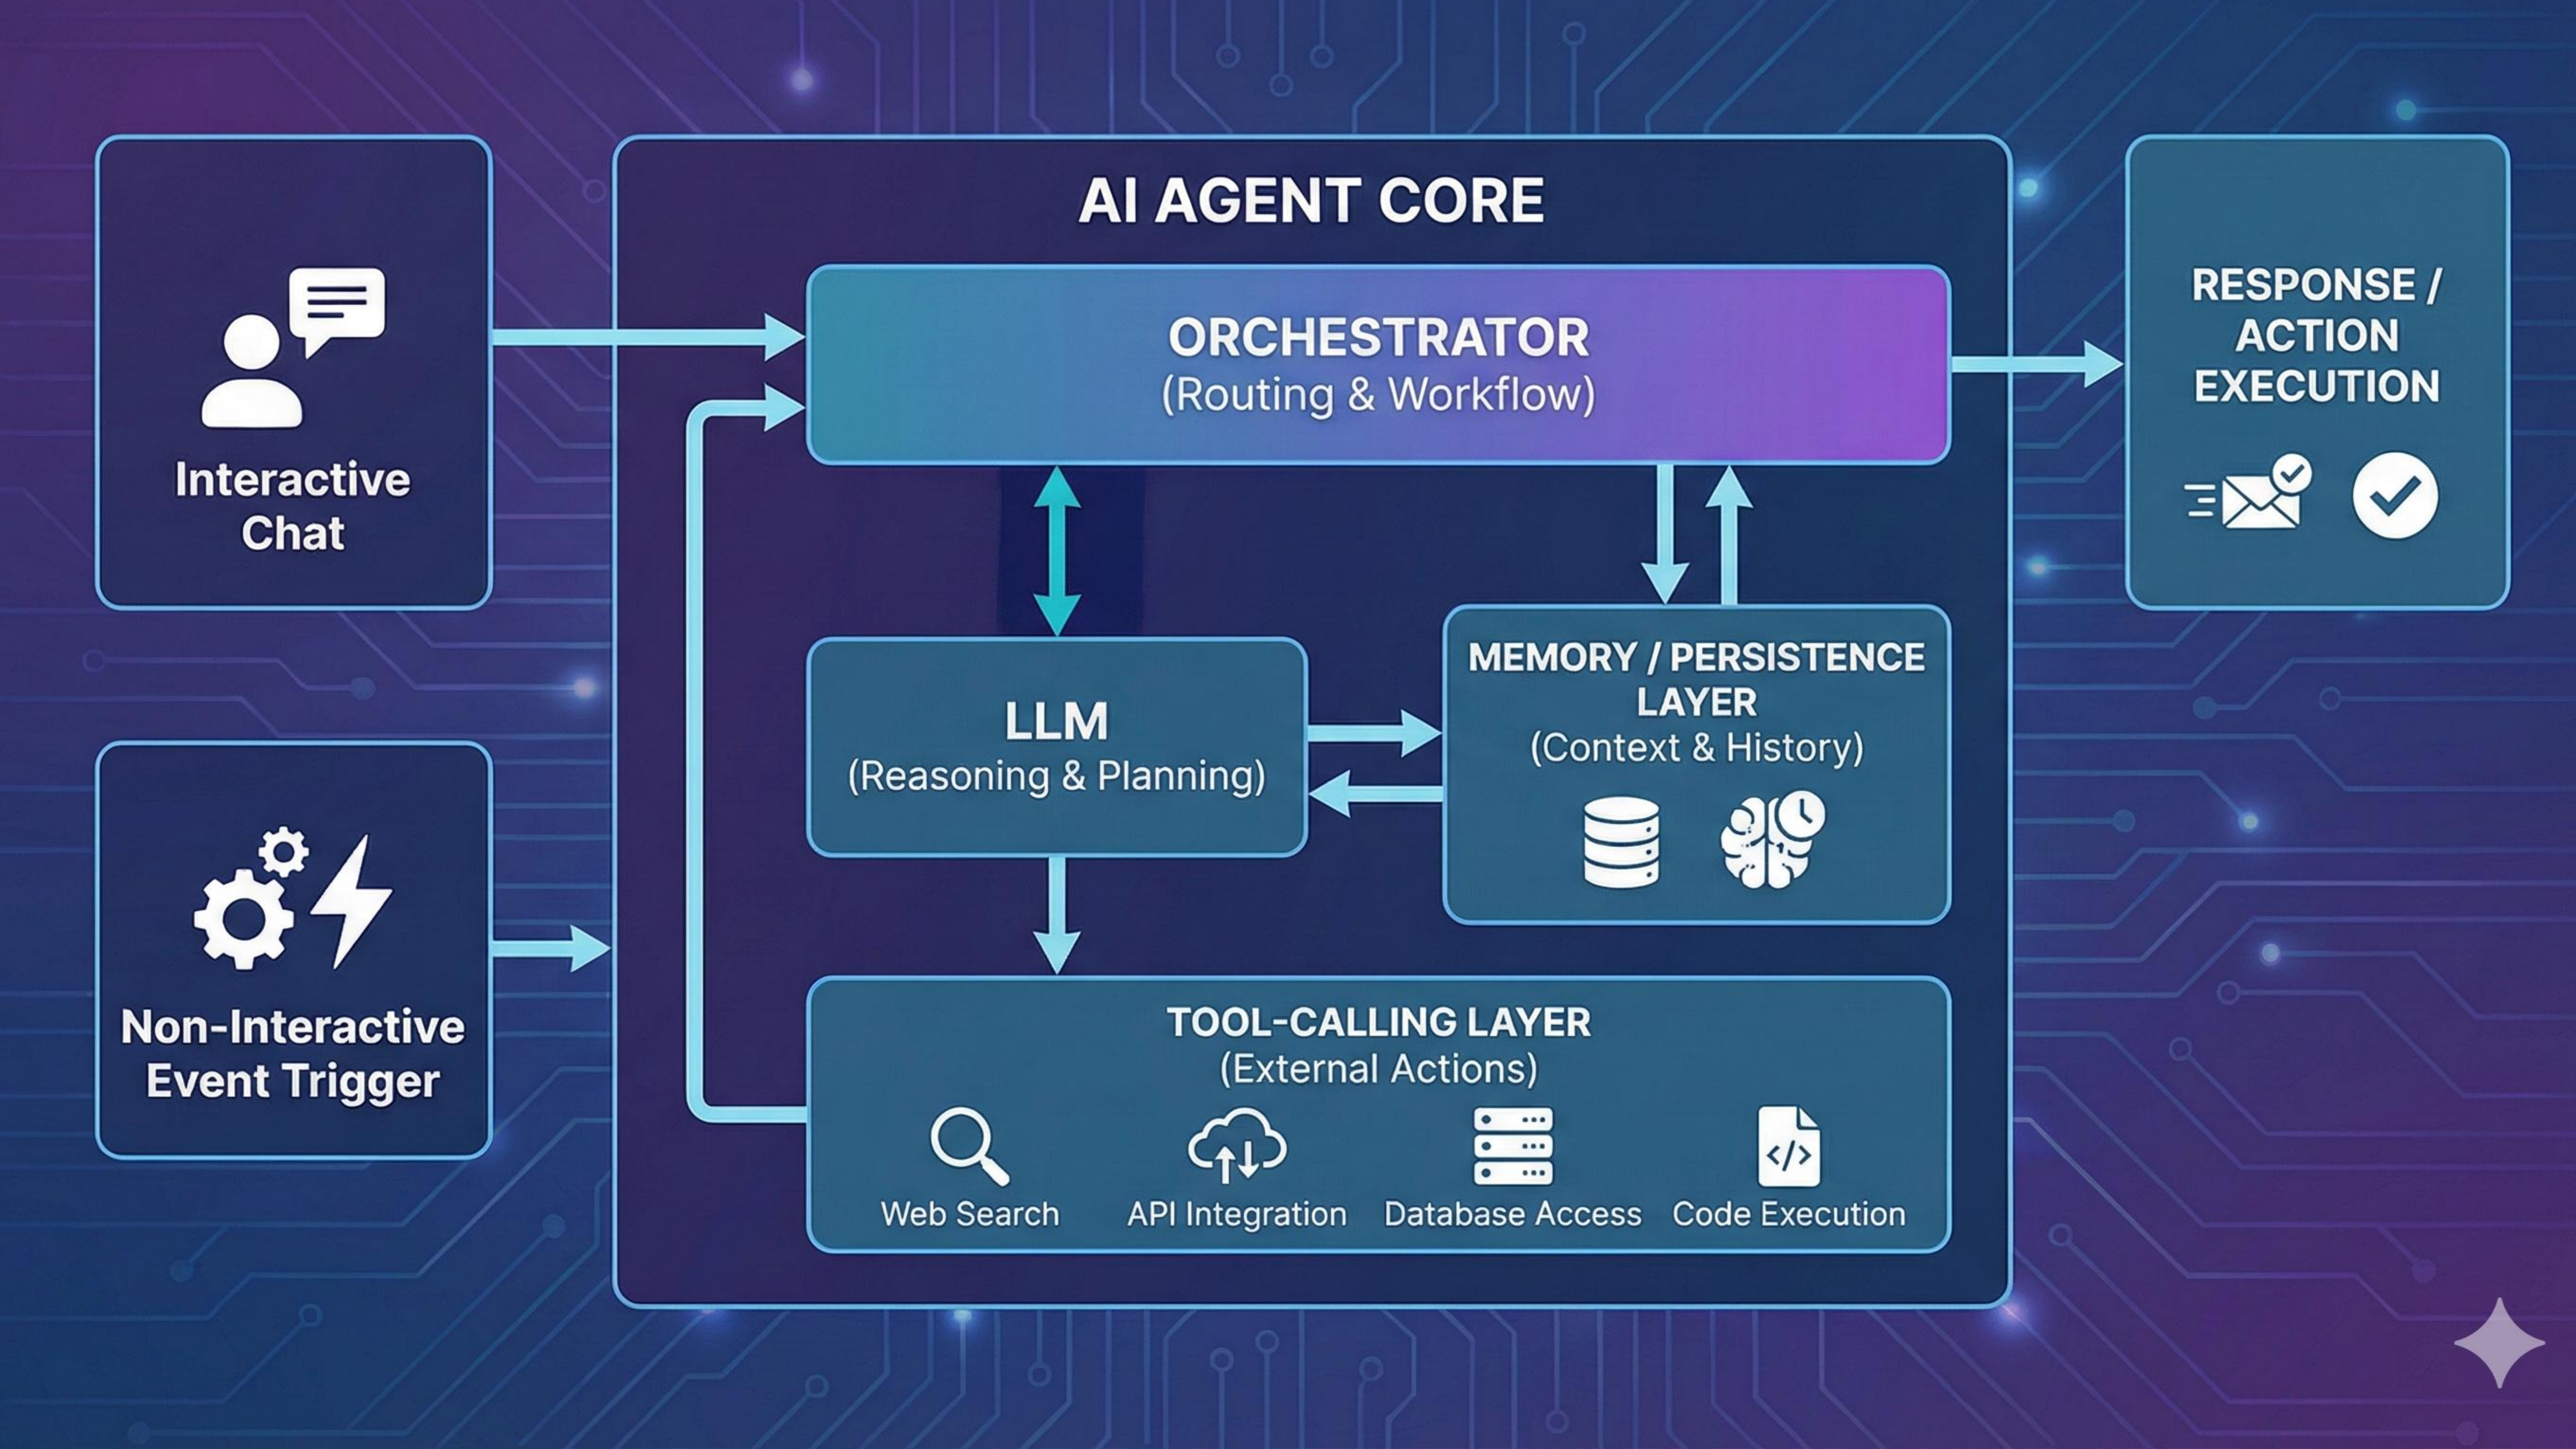
Interactive Chat
Non-Interactive Event Trigger
AI AGENT CORE
ORCHESTRATOR
(Routing & Workflow)
LLM
(Reasoning & Planning)
MEMORY / PERSISTENCE LAYER
(Context & History)
TOOL-CALLING LAYER
(External Actions)
Web Search
API Integration
Database Access
Code Execution
RESPONSE / ACTION EXECUTION

Keeping an LLM on task be like...

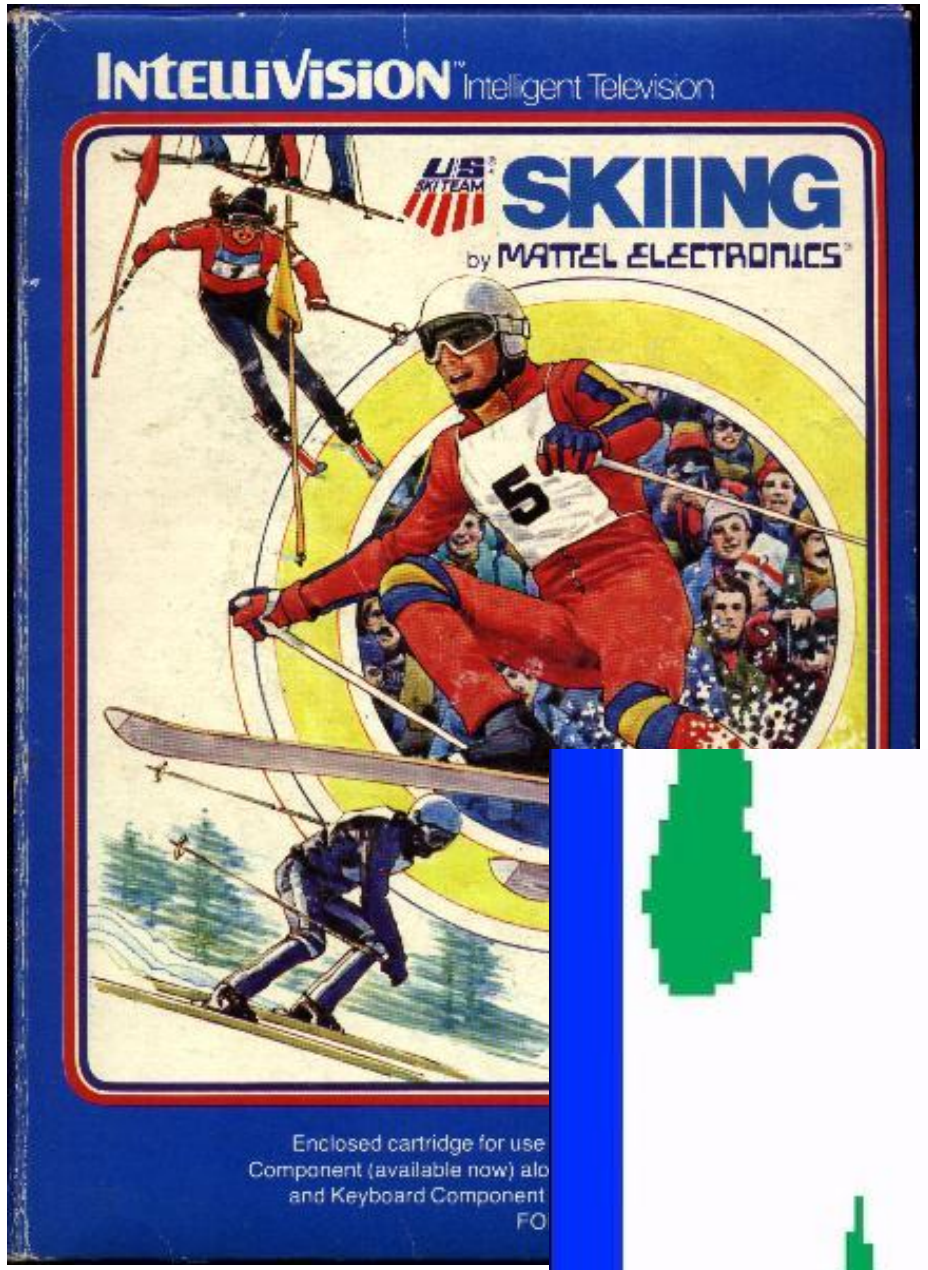

LLM Multi-Modal Processing

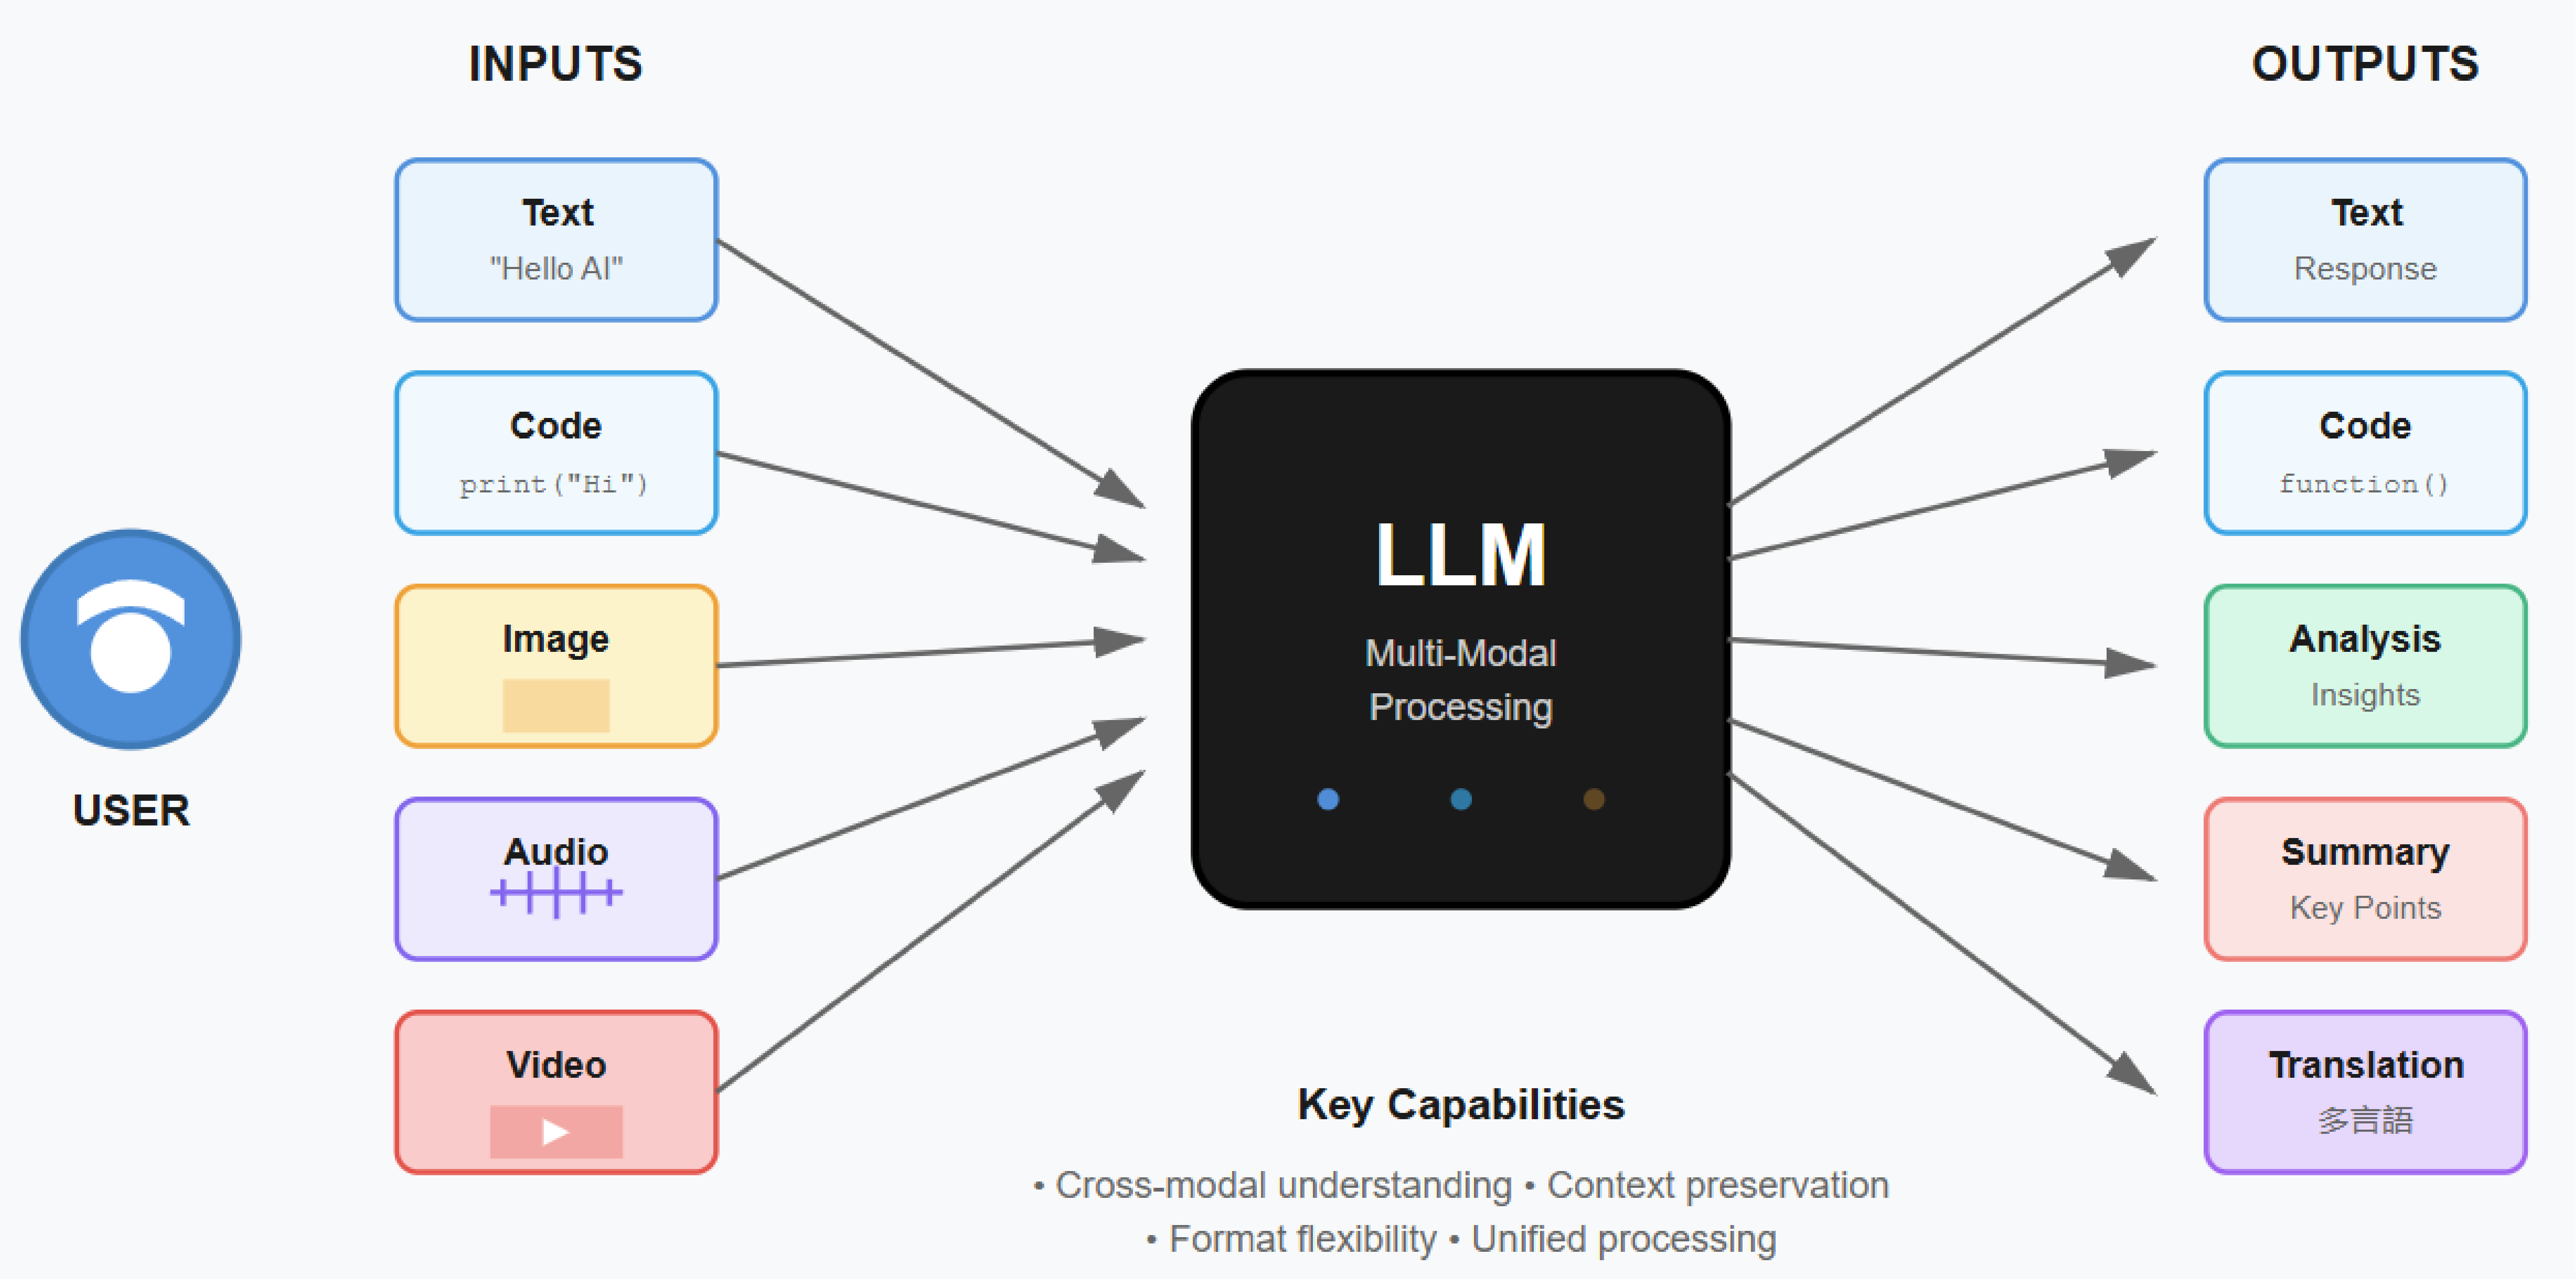

Tokenization Process in LLMs

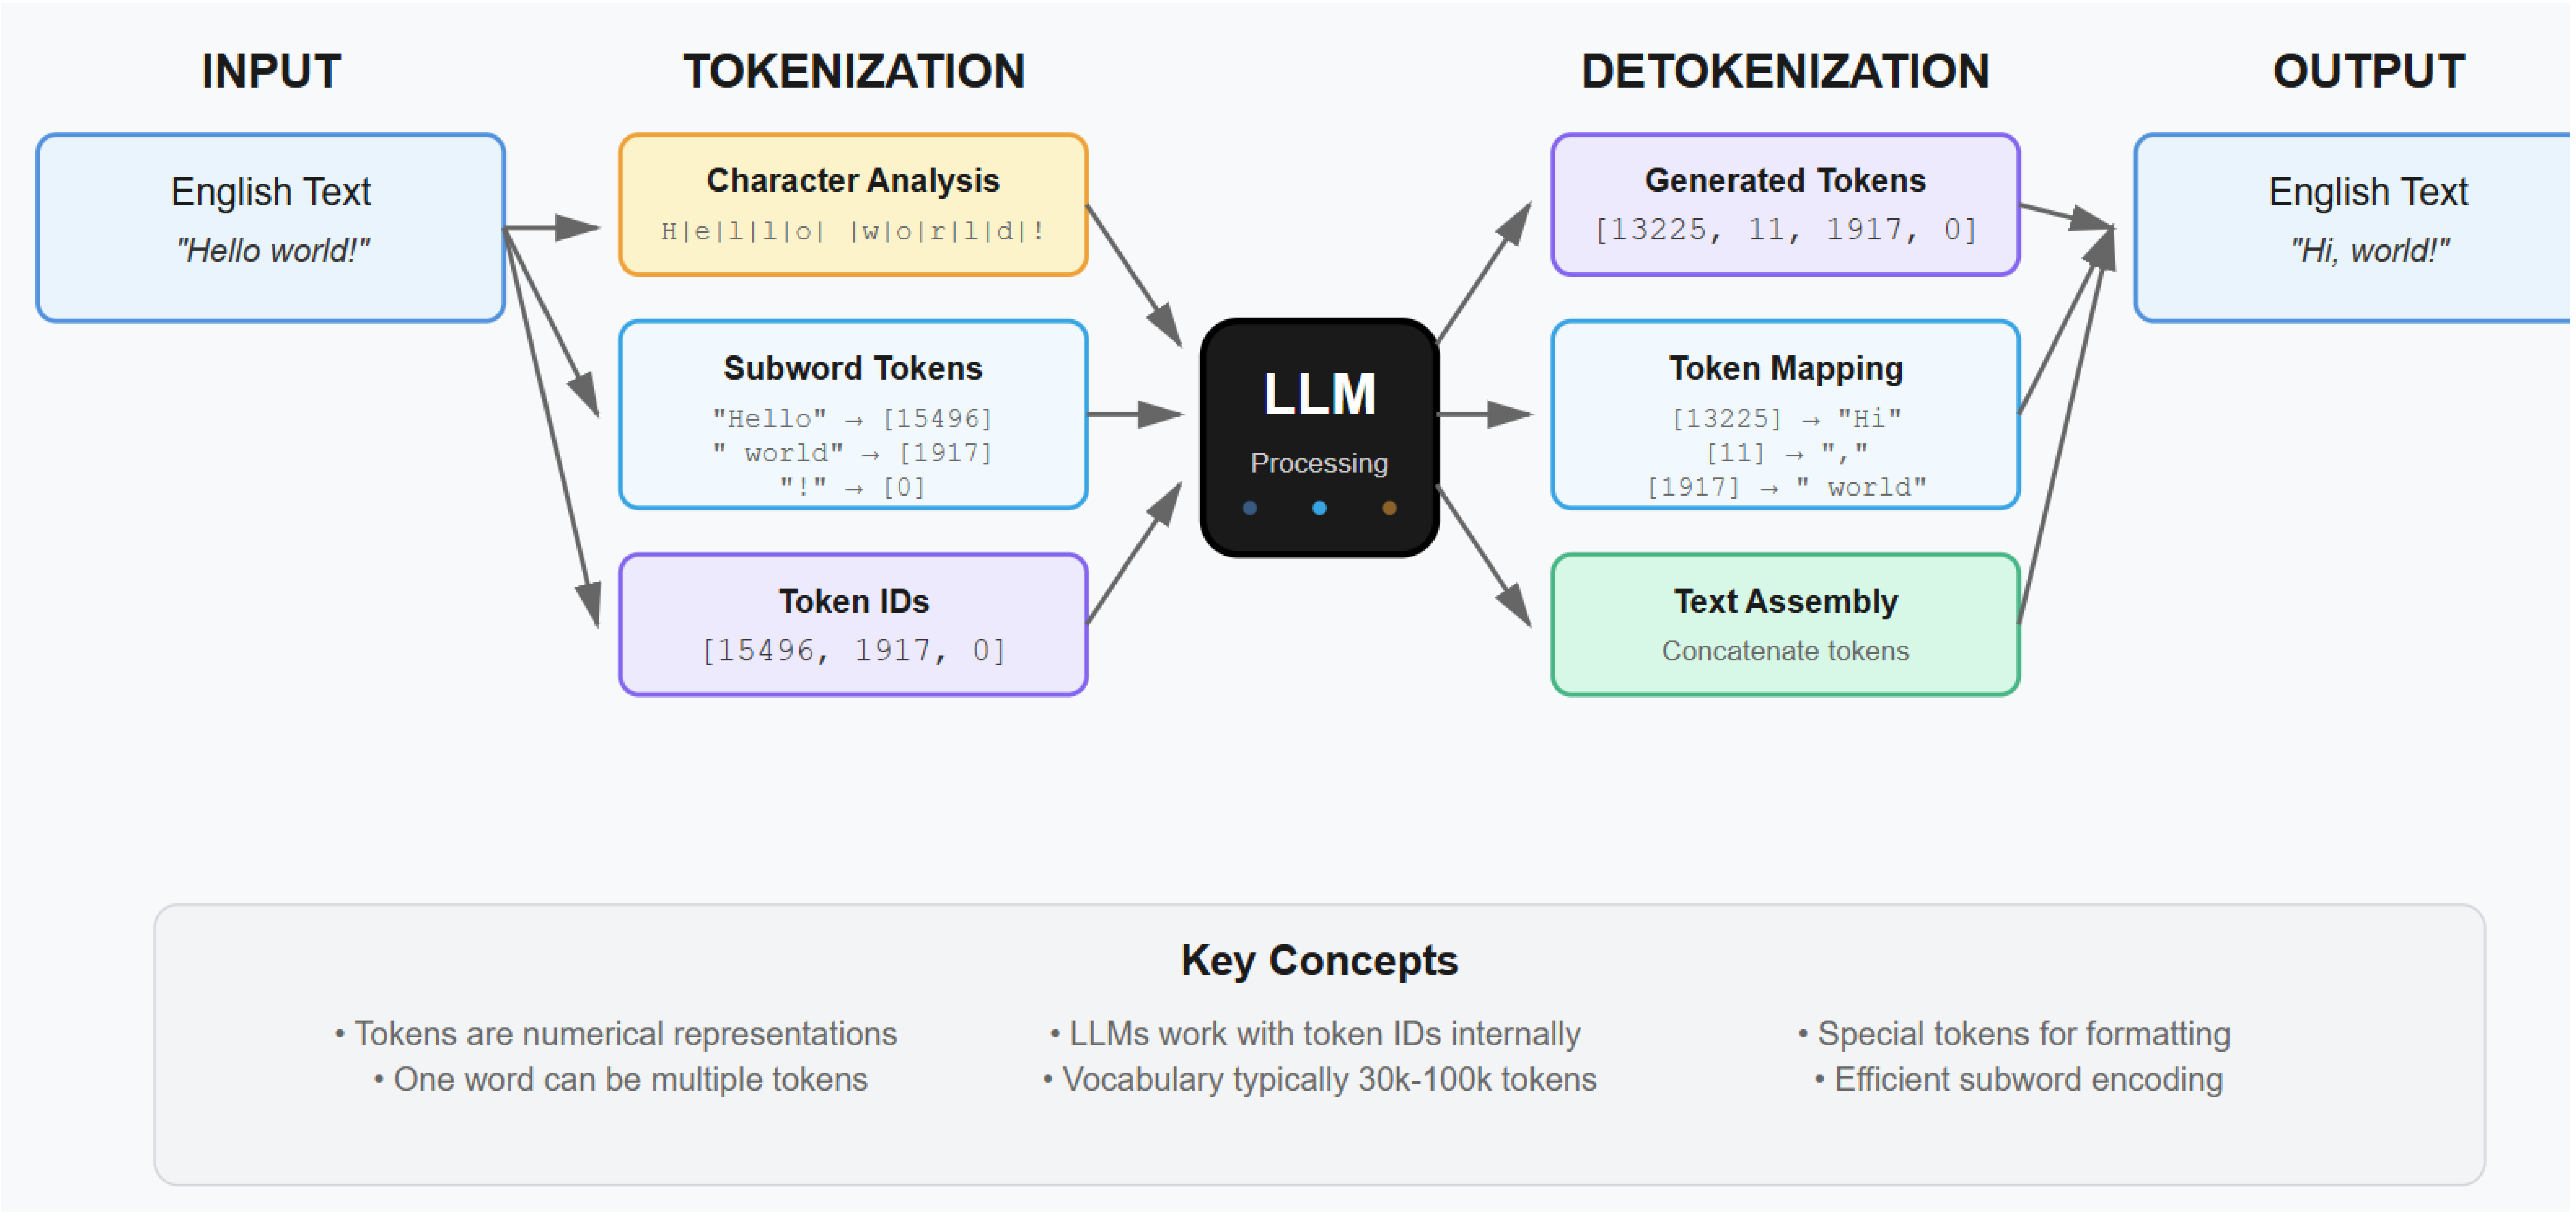

Context Window

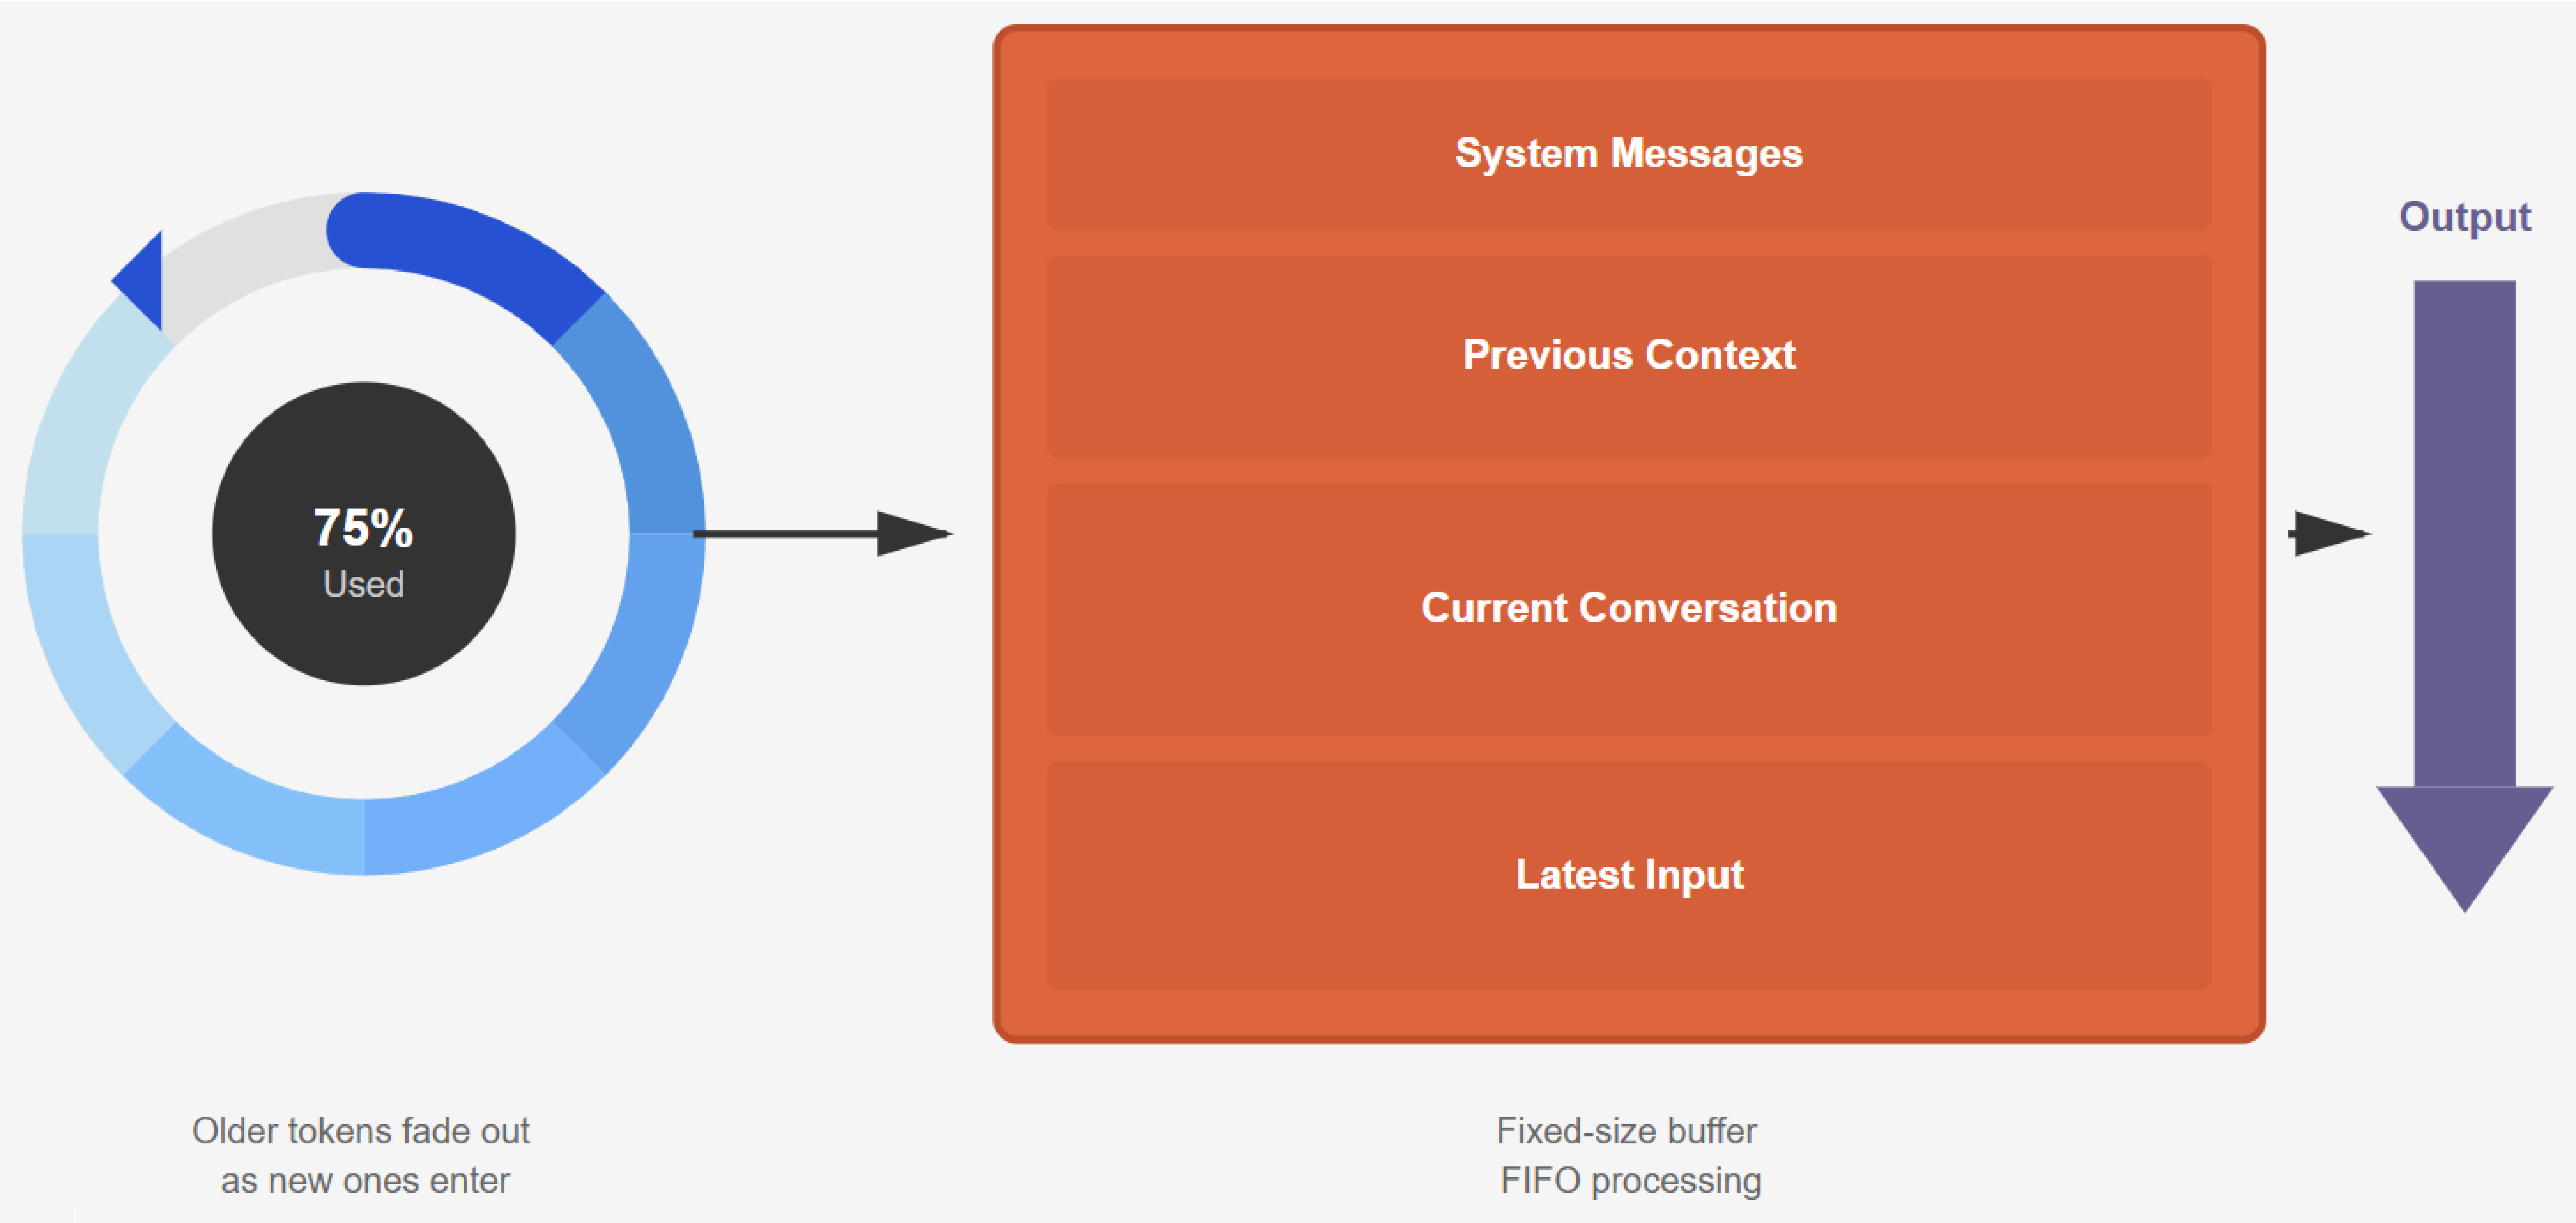

Context Length

CURRENT CONTEXT WINDOW

```
{
  "messages": [
    {
      "role": "system",
      "content": "You are a helpful AI assistant..."
    },
    {
      "role": "user",
      "content": "Help me write a Python script"
    },
    {
      "role": "assistant",
      "content": "I'll help you write a Python..."
    },
    {
      "role": "user",
      "content": "Add error handling please"
    },
    {
      "role": "assistant",
      "content": "Here's the updated script..."
    }
  ],
  "model": "gpt-4",
  "max_tokens": 2000,
  "temperature": 0.7
}
```

AFTER CONTEXT OVERFLOW (TRUNCATION)

```
{
  "messages": [
    {
      "role": "system",
      "content": "You are a helpful AI assistant..."
    },
    // Older messages truncated to fit context
    {
      "role": "assistant",
      "content": "Here's the updated script..."
    },
    {
      "role": "user",
      "content": "Now add logging functionality"
    },
    {
      "role": "assistant",
      "content": "I'll add logging to your..."
    }
  ],
  "model": "gpt-4",
  "max_tokens": 2000
}
```

How Context Management Works

Token Accumulation

Each message (system, user, assistant) consumes tokens. The total must stay under the model's context limit (4K, 8K, 16K, 128K).

FIFO Eviction

When approaching the limit, oldest messages (except system prompt) are removed first. This maintains recent context but loses early history.

Sliding Window

The context window slides forward with each interaction, keeping the most recent N tokens while discarding older ones.

Agentic AI System Architecture

Entity Relationship & Data Flow Diagram

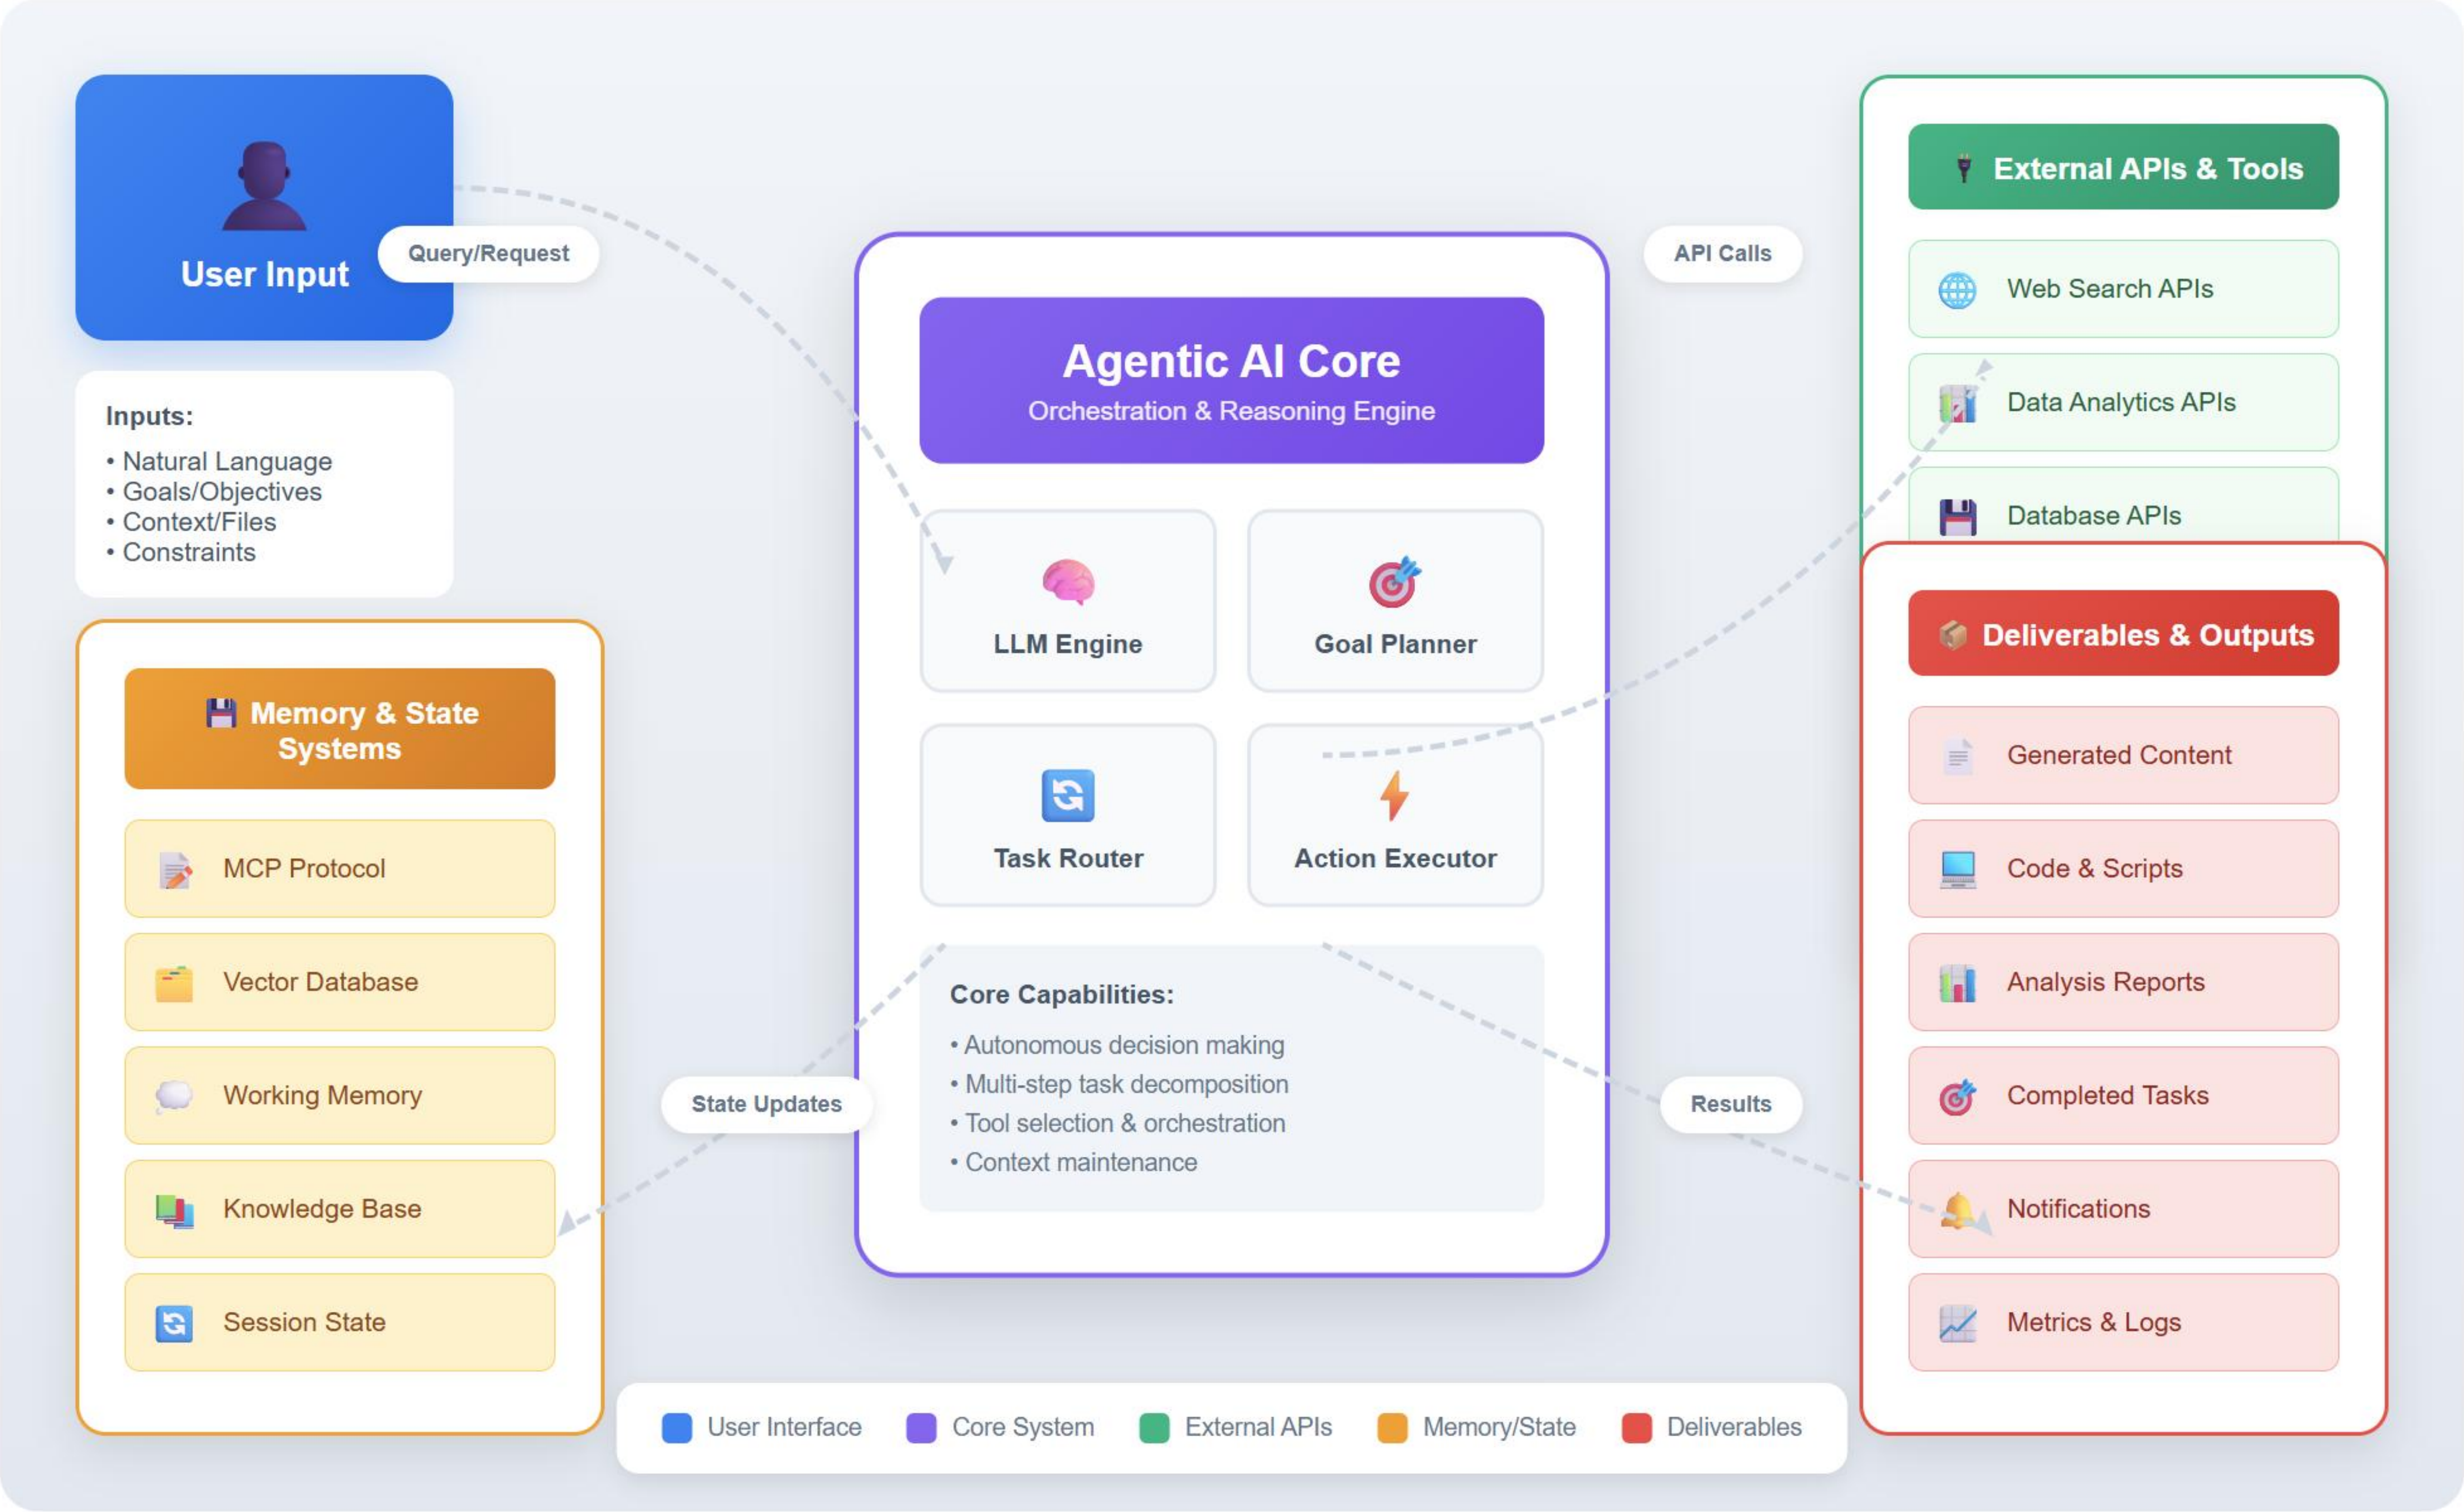

Model Context Protocol (MCP) Architecture

Open Protocol for Seamless LLM-Application Integration

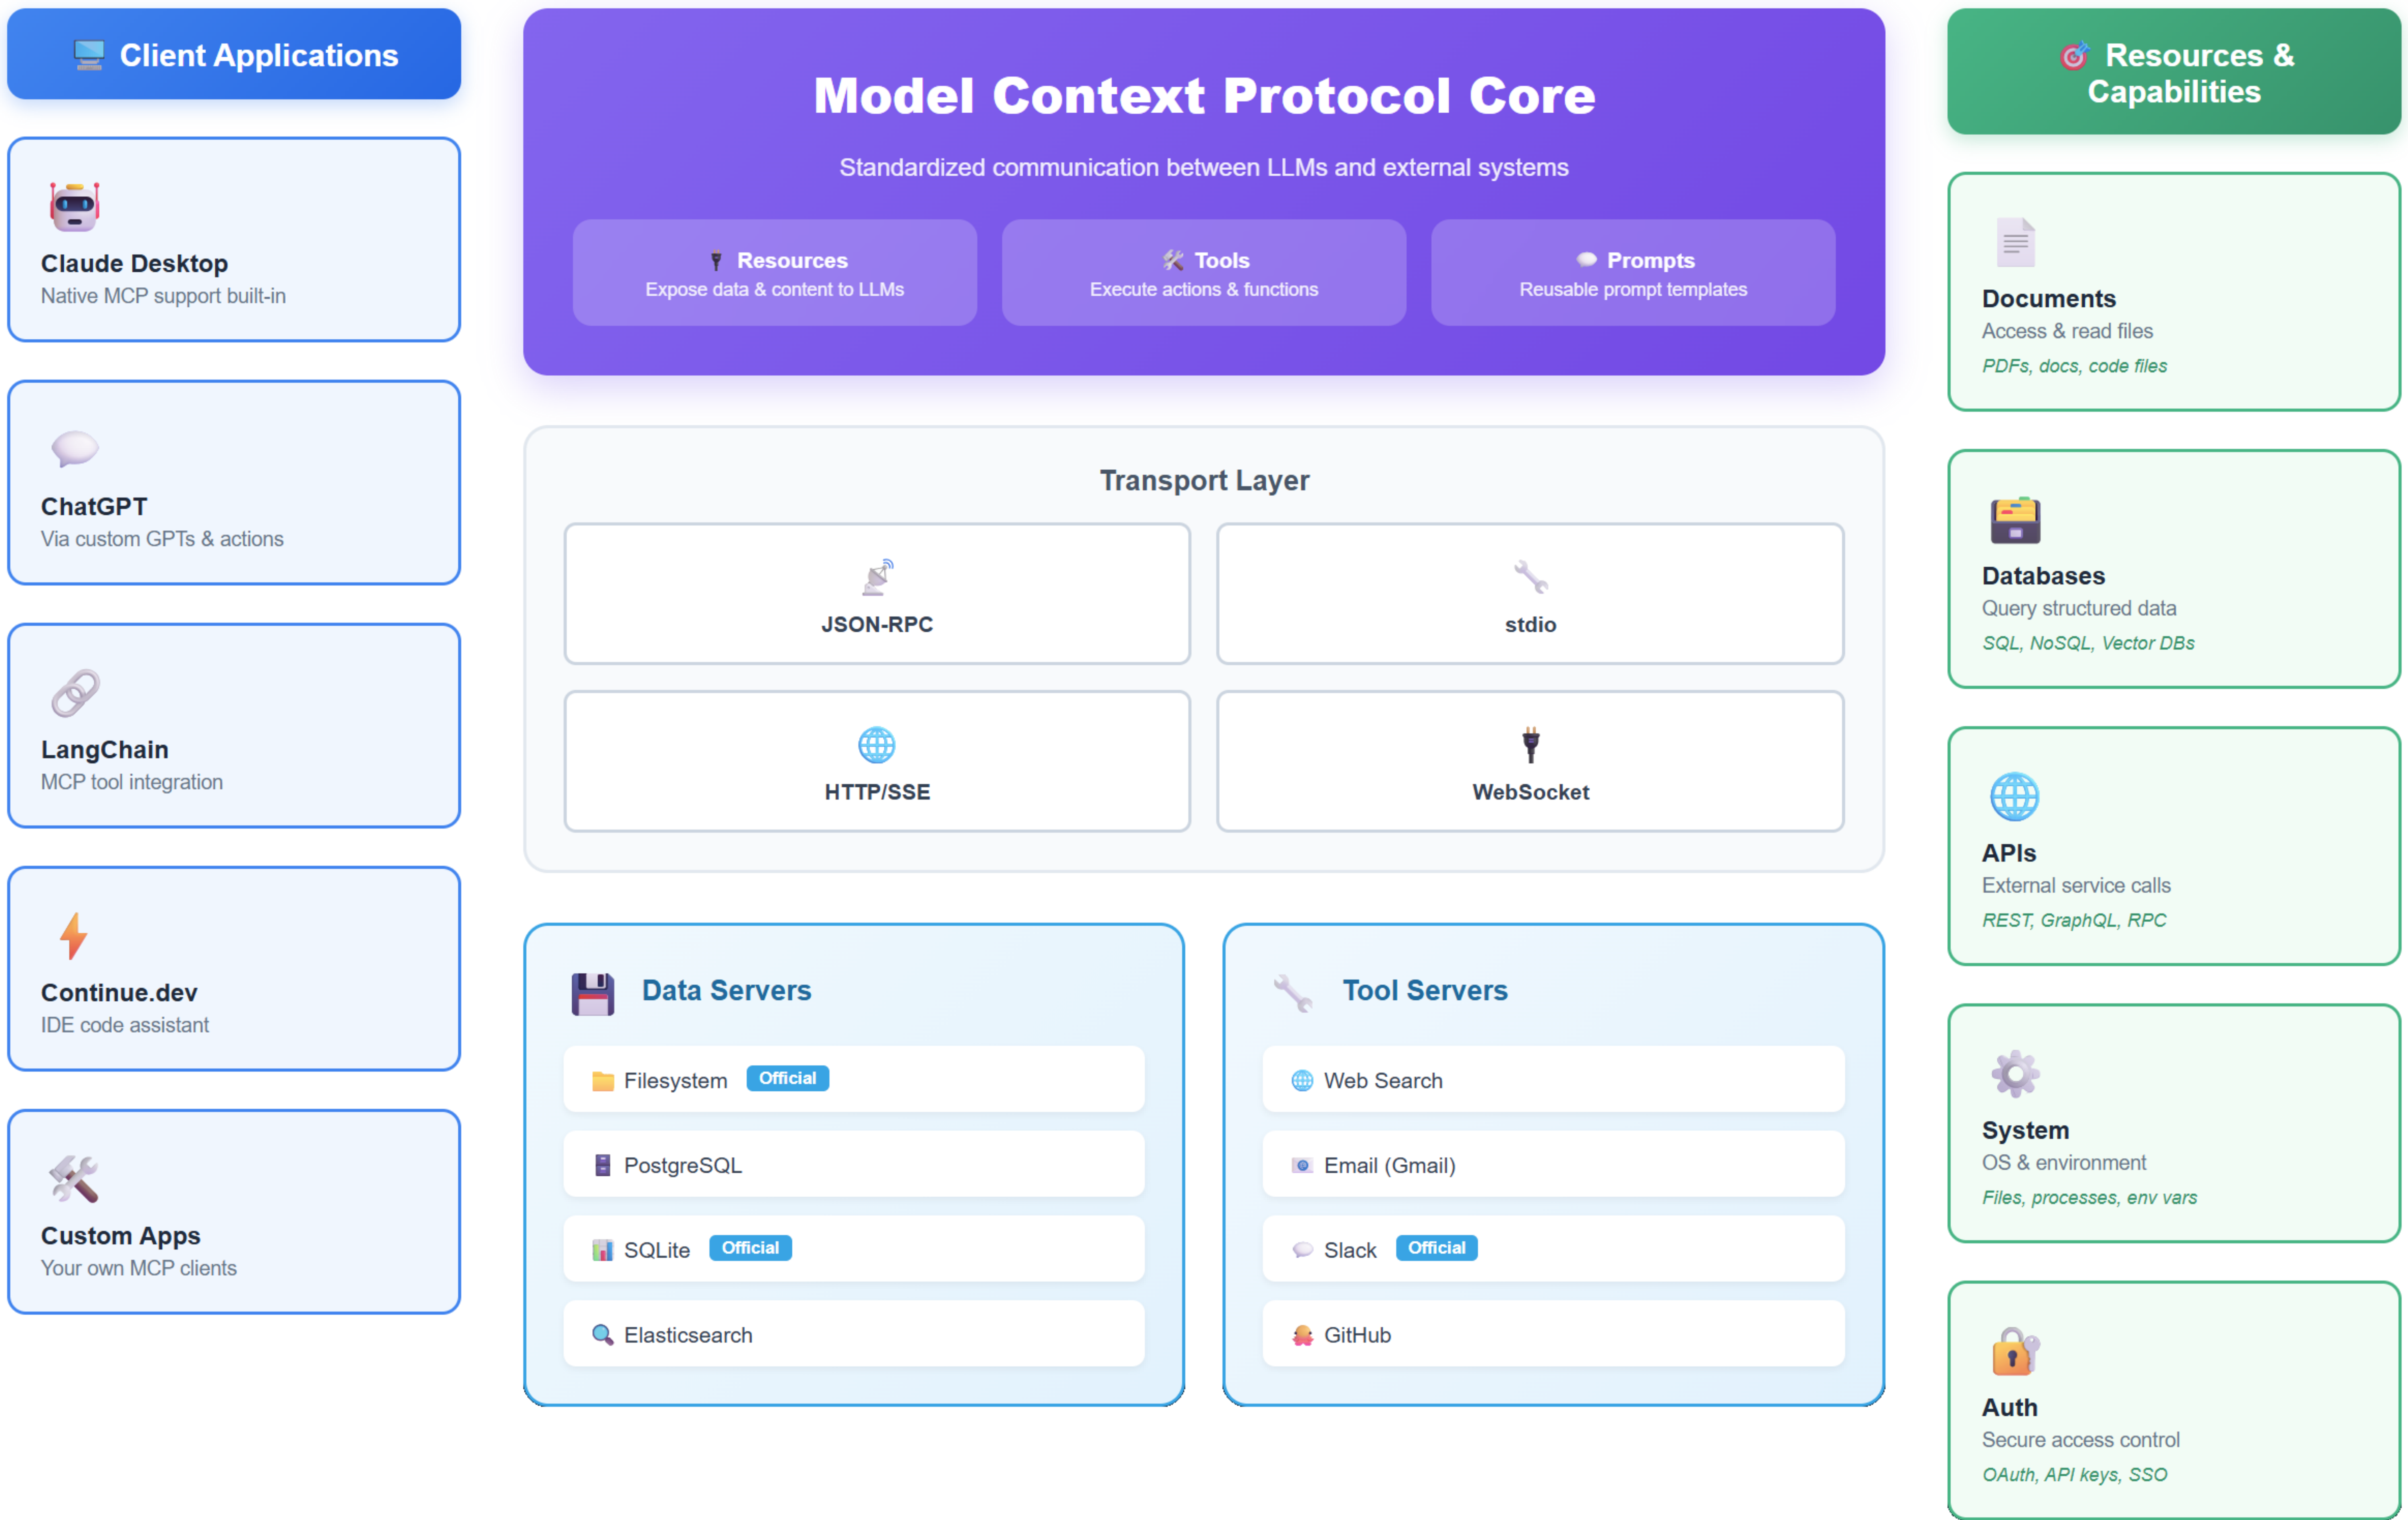

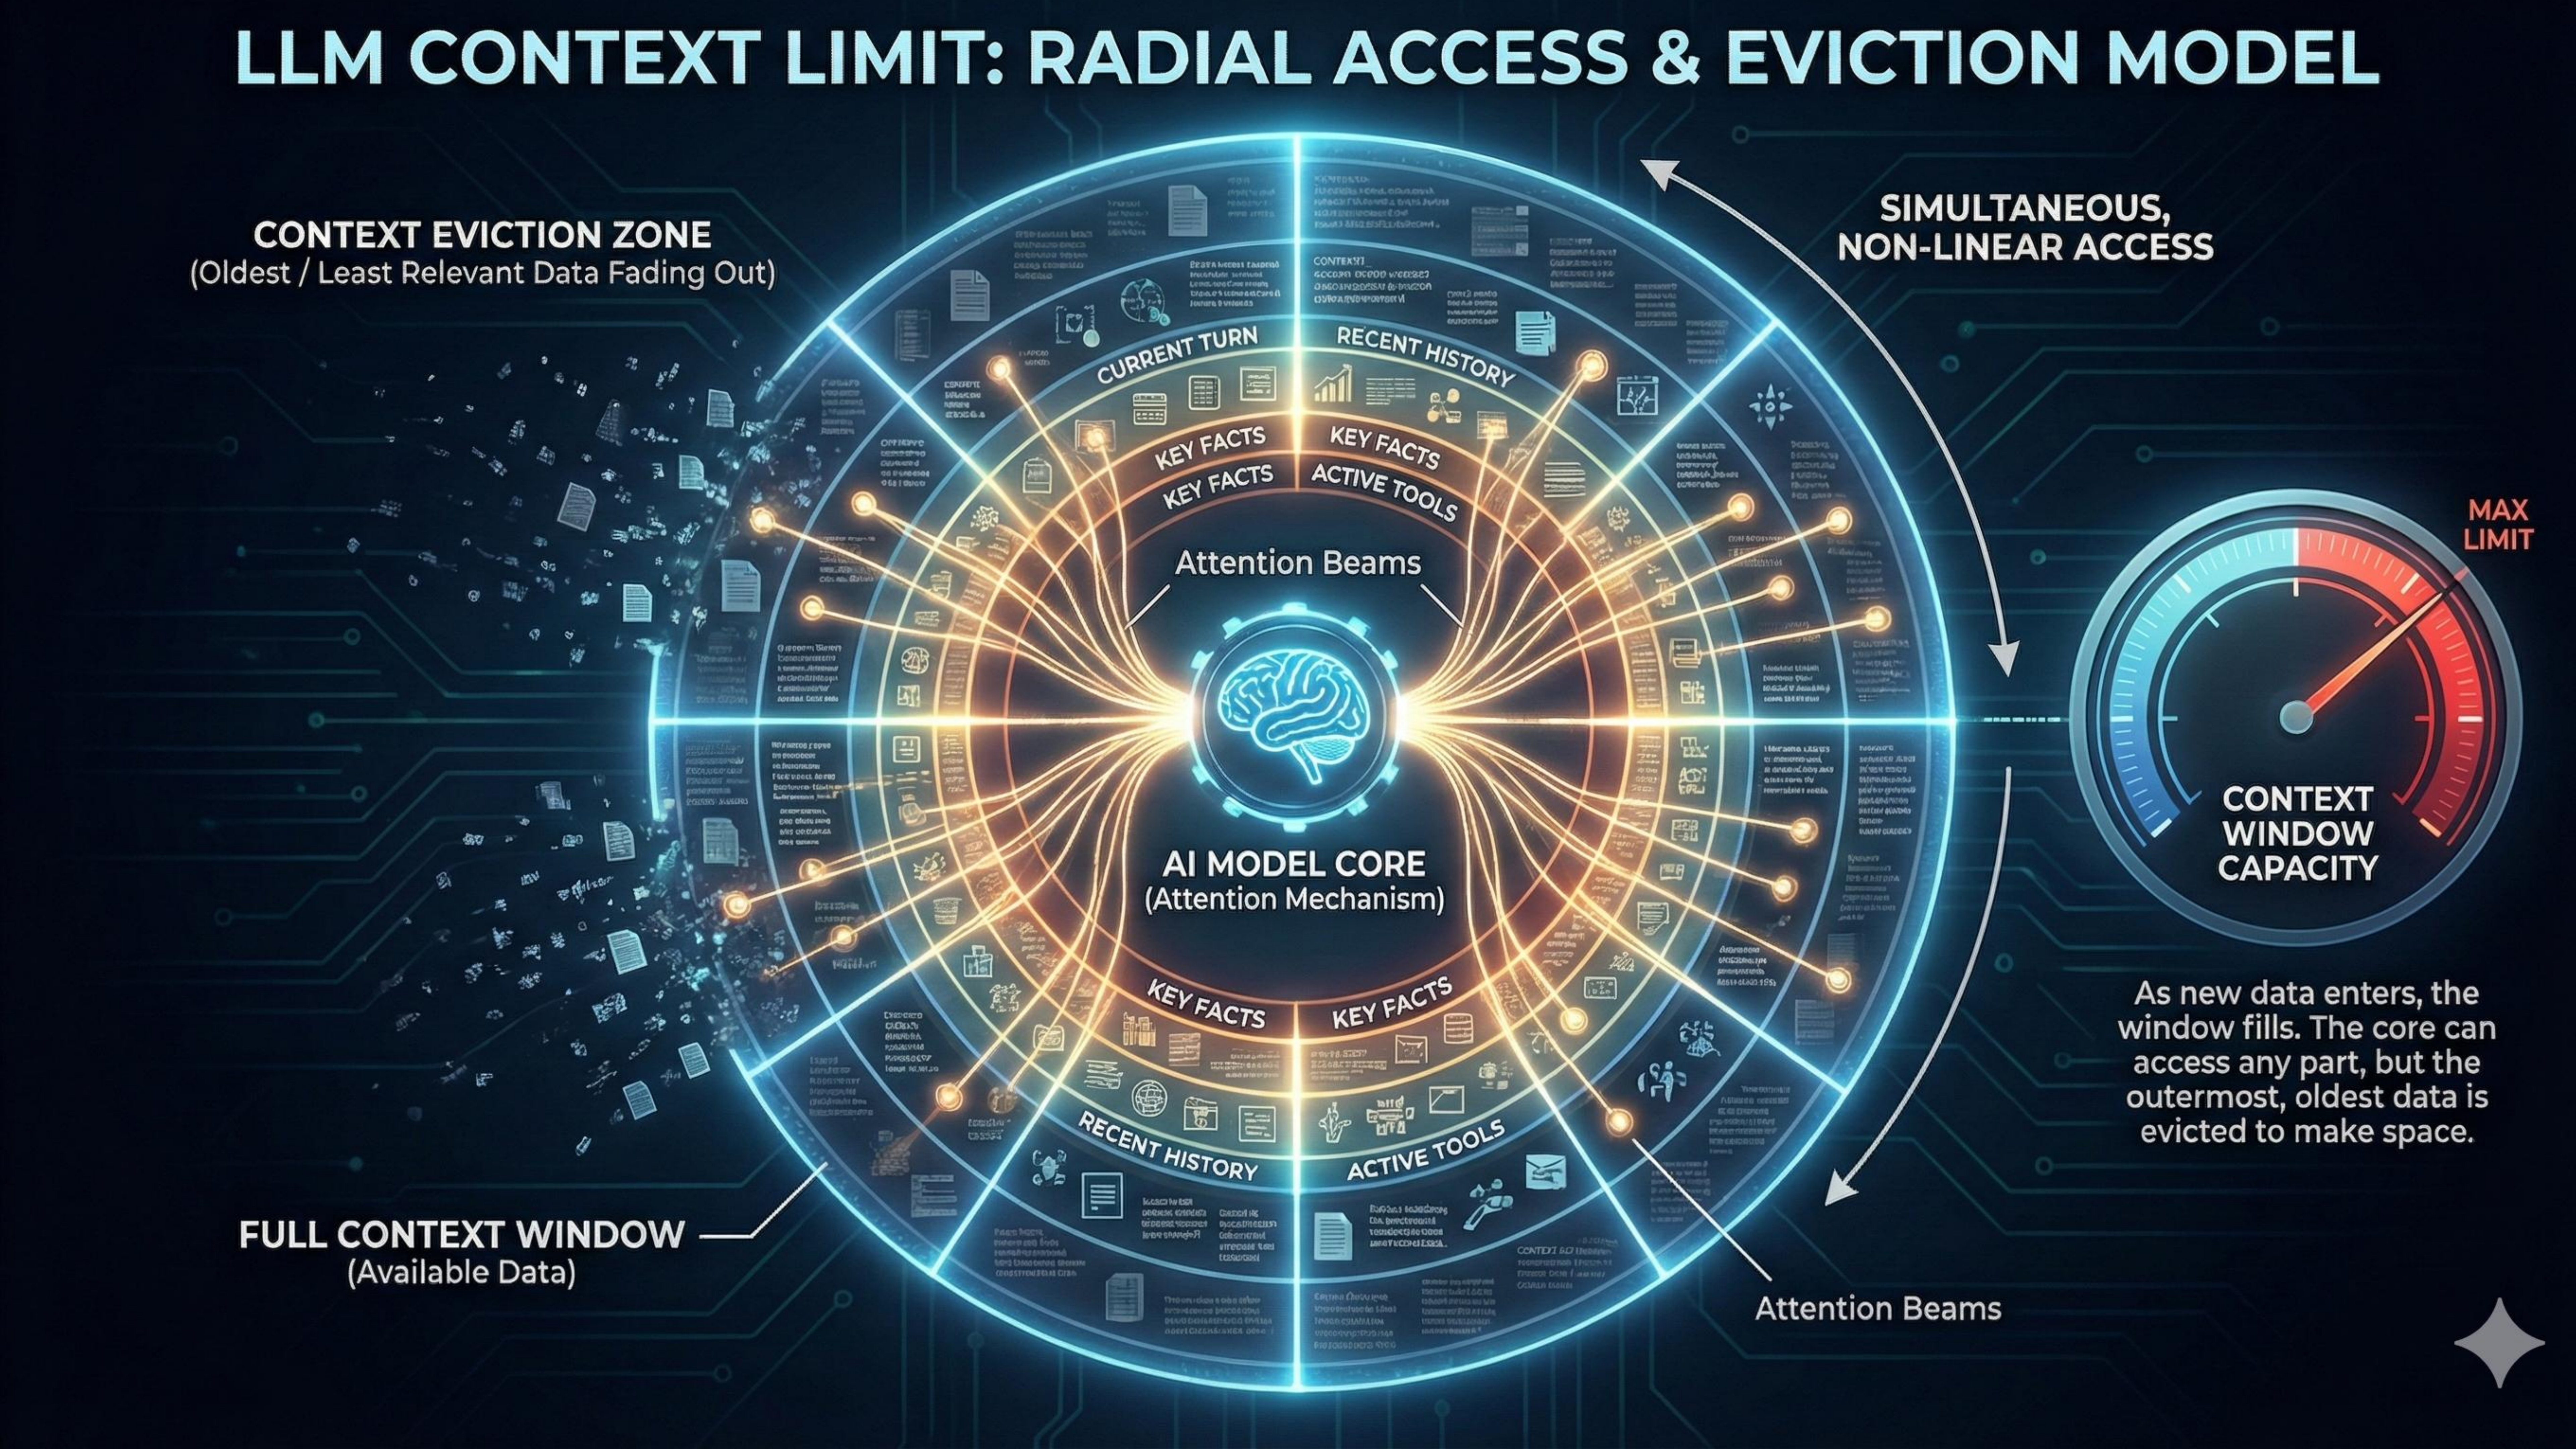
LLM CONTEXT LIMIT: RADIAL ACCESS & EVICTION MODEL
CONTEXT EVICTION ZONE
(Oldest / Least Relevant Data Fading Out)
SIMULTANEOUS, NON-LINEAR ACCESS
CURRENT TURN
RECENT HISTORY
KEY FACTS
KEY FACTS
KEY FACTS
ACTIVE TOOLS
Attention Beams
AI MODEL CORE
(Attention Mechanism)
KEY FACTS
KEY FACTS
RECENT HISTORY
ACTIVE TOOLS
MAX LIMIT
CONTEXT WINDOW CAPACITY
As new data enters, the window fills. The core can access any part, but the outermost, oldest data is evicted to make space.
FULL CONTEXT WINDOW
(Available Data)
Attention Beams

EFFICIENT CONTEXT & TOKEN MANAGEMENT FOR AI AGENTS (Avoiding Pollution)
1. INPUT FILTERING & PRE-PROCESSING
NOISY SOUP
TAGS
Remove noise, irrelevant data, and duplicates *before* the LLM sees it.
2. STRUCTURED PROMPTING & SYSTEM MESSAGES
{...}
[SYSTEM]
[CONTEXT]
[USER]
Clearly separate instructions from conversation history. Define boundaries.
3. CONTEXT SUMMARIZATION (Compress History)
Σ
Regularly condense older turns into concise summaries. Don't re-send everything.
4. RELEVANT RETRIEVAL (Sliding Window/RAG)
Fetch only the ***most recent*** or ***most relevant*** past context pieces.
5. INSTRUCT FOR CONCISENESS
Be Brief.
Directly tell the agent to keep responses short and to the point.
OPTIMIZED AGENT OUTPUT
Result: Lower costs, faster responses, and focused, accurate interactions.

MODEL CONTEXT PROTOCOL (MCP) SCHEMATIC FLOW

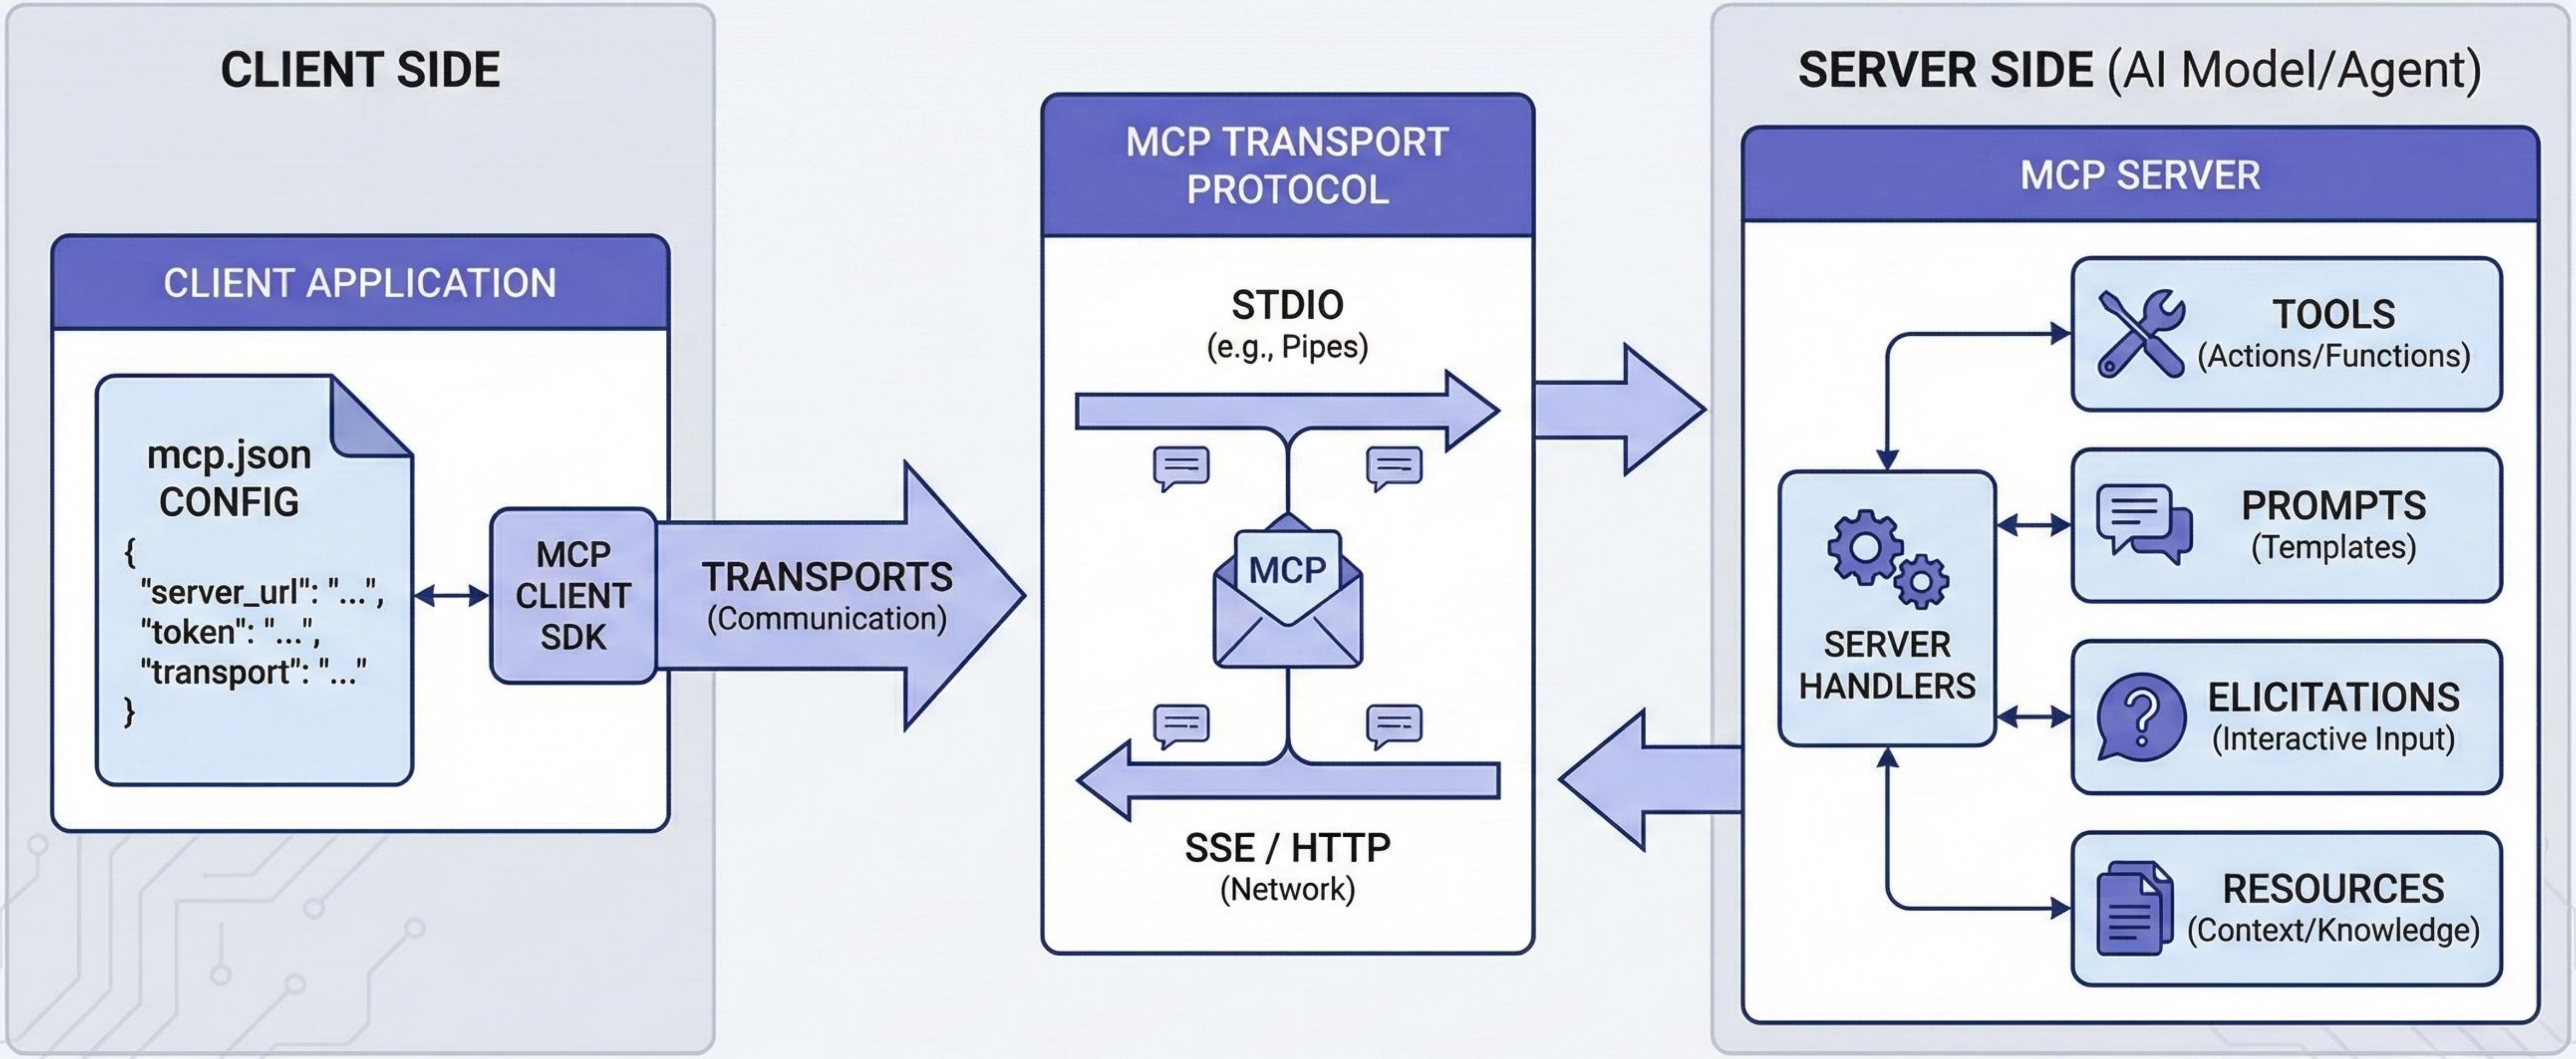

UPDATED O'REILLY LIVE LEARNING: PYTHON CLI AGENTIC AI POC APP (CLI, LangGraph, CrewAI, MCP, SQLite)

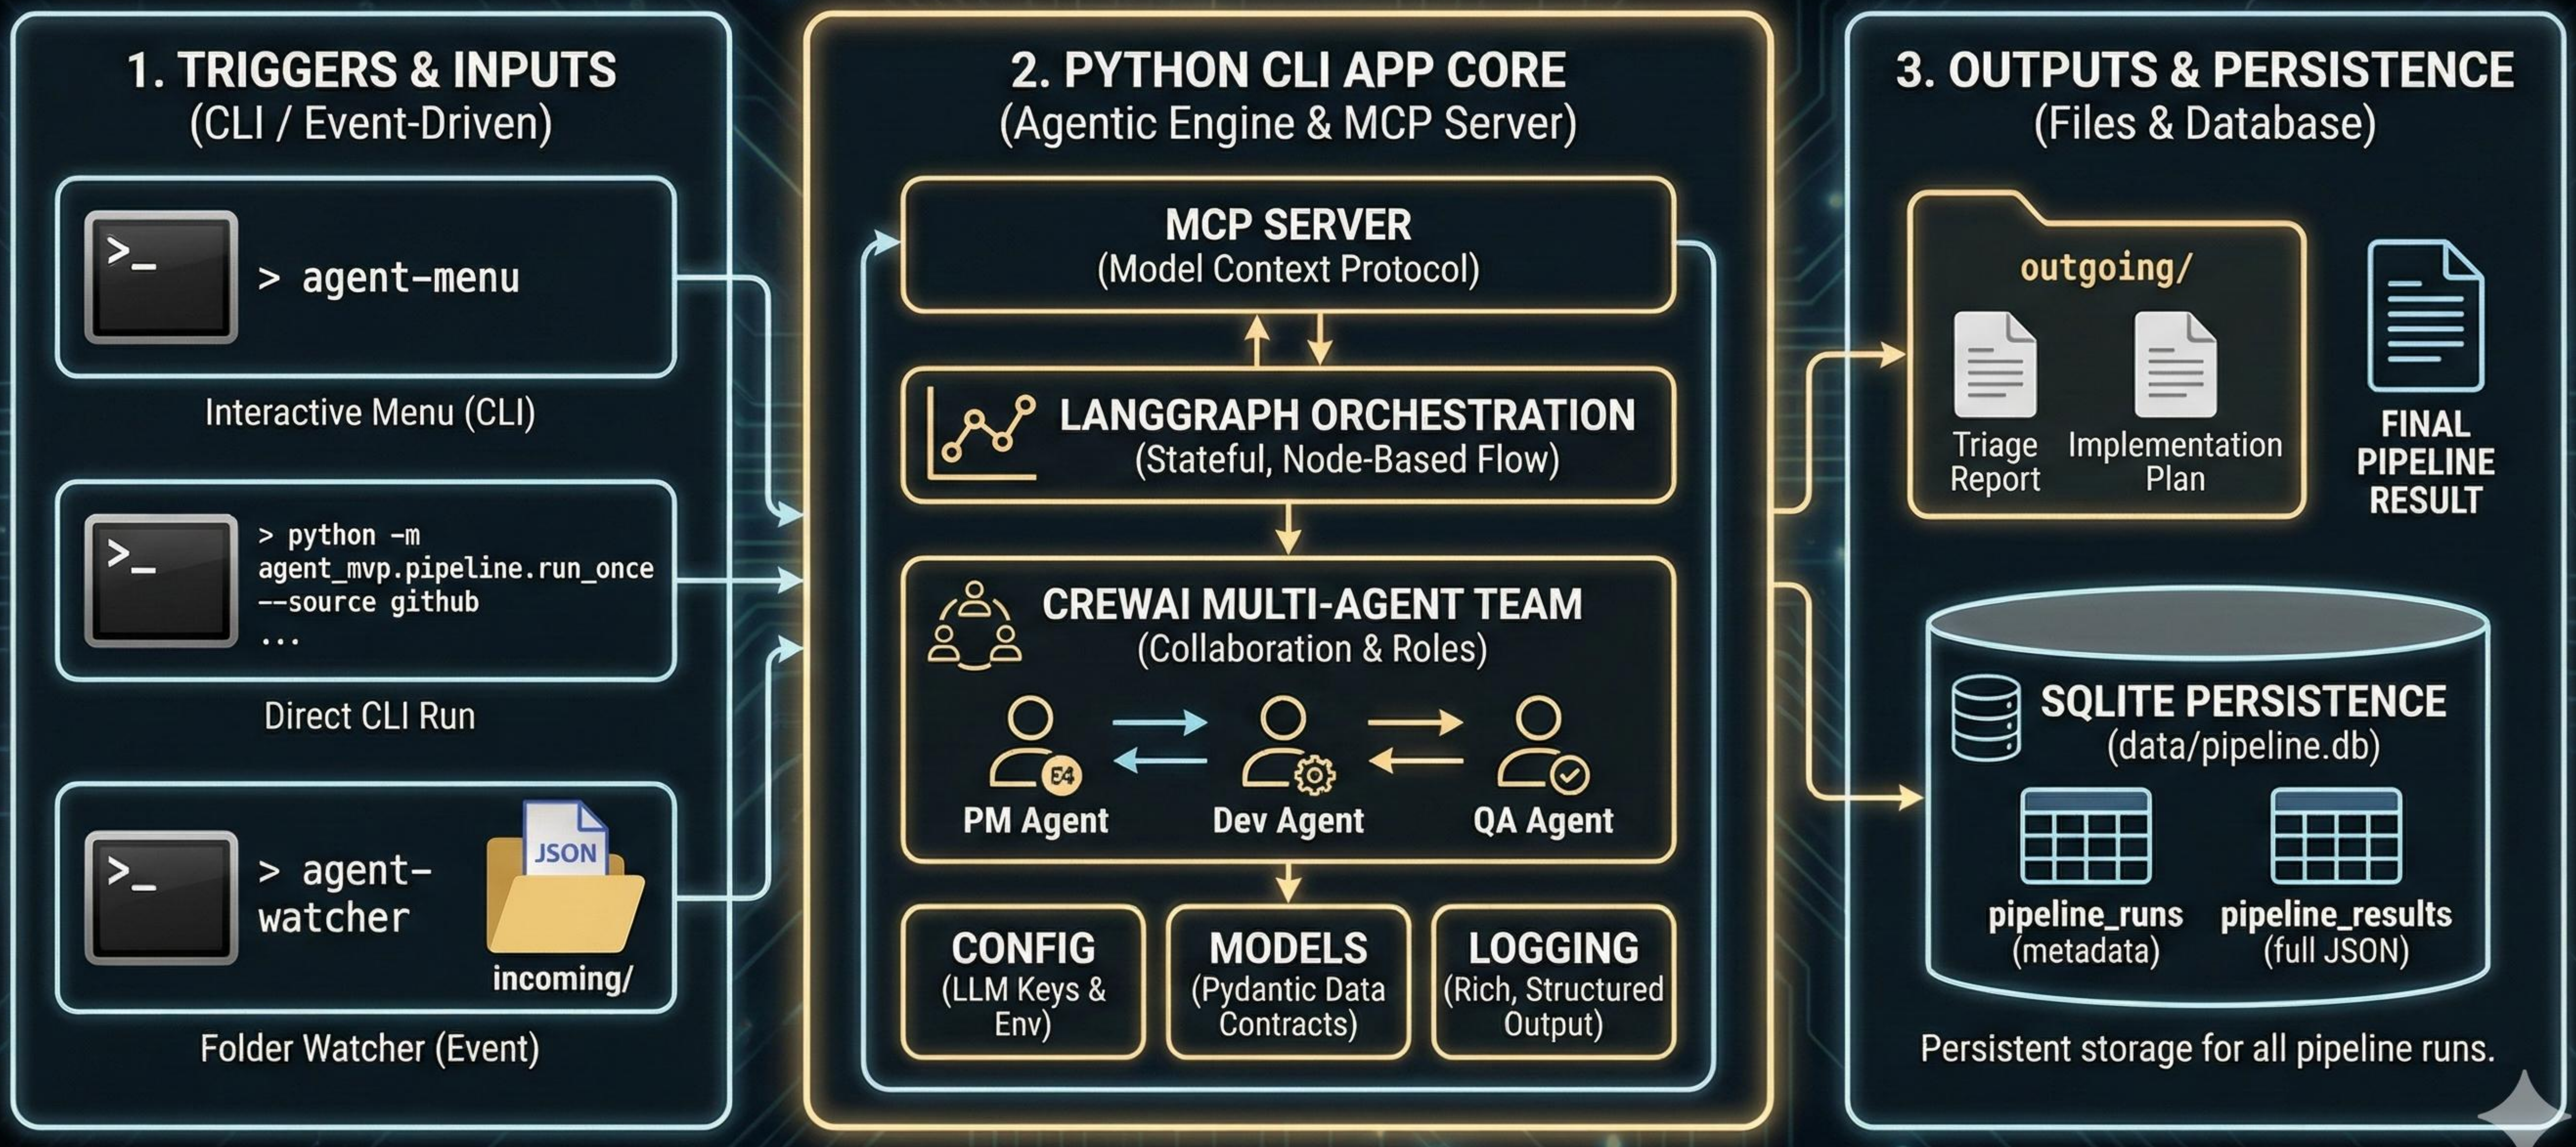

LANGGRAPH: STATEFUL AGENT ORCHESTRATION MADE CLEAR

CORE CONCEPT: GRAPH-BASED WORKFLOWS

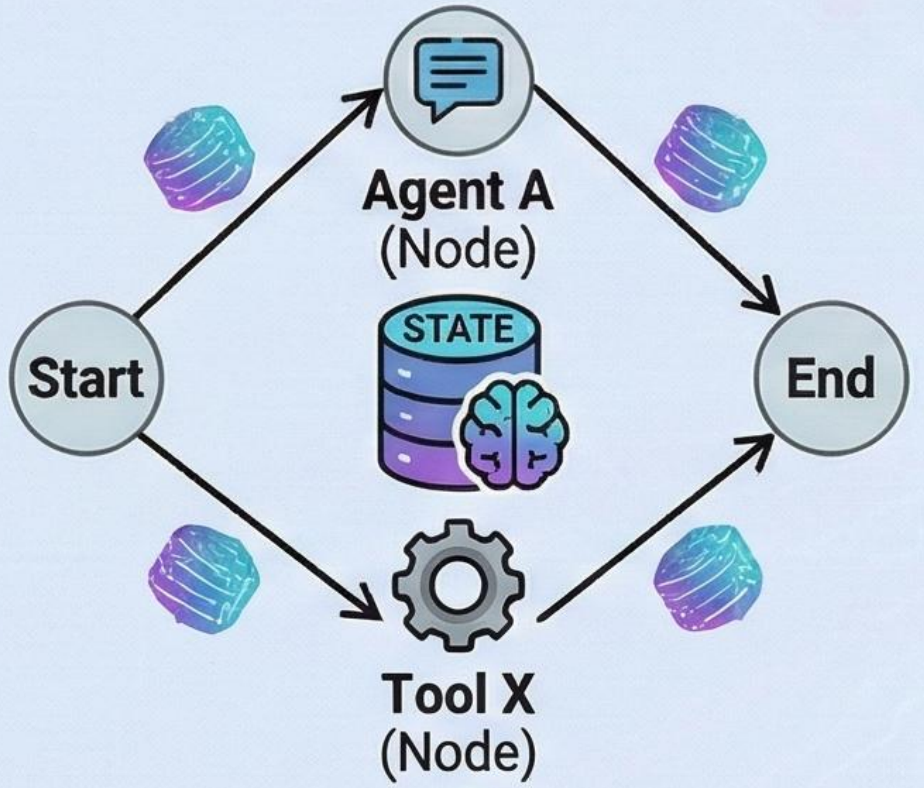

Defines the sequence of operations as a graph of nodes (agents, tools) and edges (transitions), managing the overall application state.

Stateful & Defined Paths
[4.png, 5.png]

CHIEF FUNCTIONALITY: PRECISE CONTROL & CYCLES

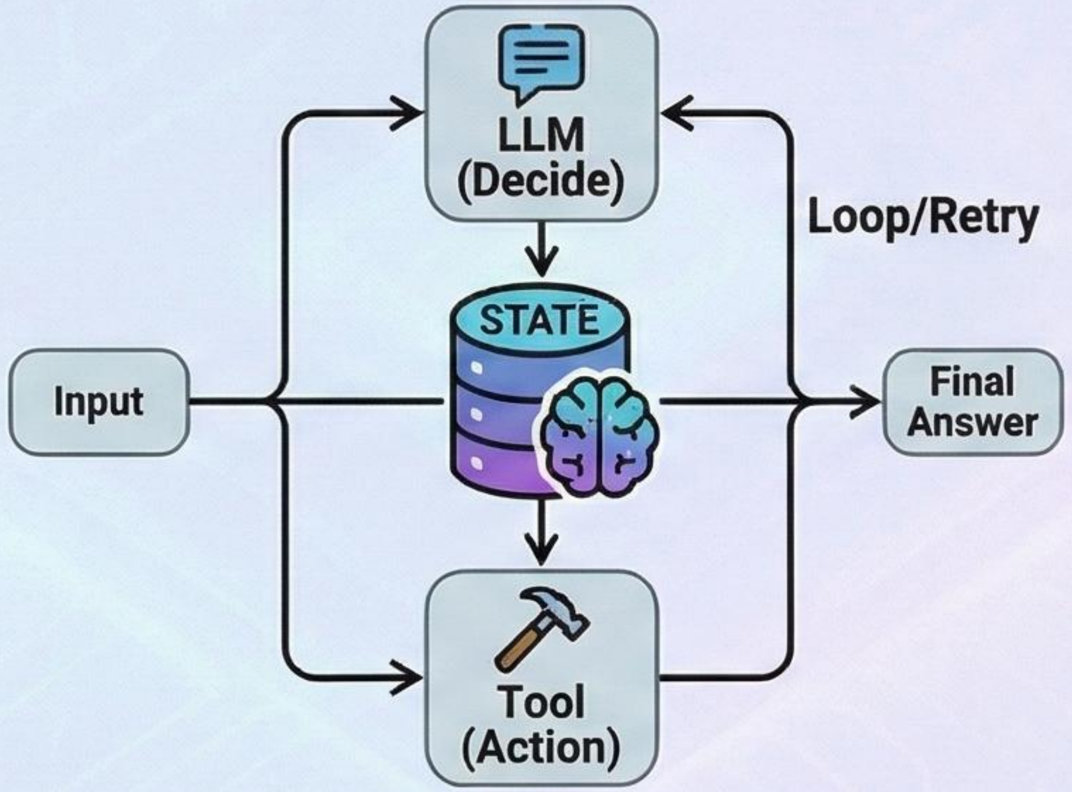

Allows for complex, cyclic workflows where an agent can retry, loop, or branch based on intermediate results and memory, not just a linear chain.

Beyond Linear Chains
[4.png, 5.png]

HOW & WHY IT WORKS: RELIABLE & FLEXIBLE AUTOMATION

redo

1. Cycles & Retries
(Handles complex logic)

2. Persistent State
(Memory across steps)

{/}

3. Custom Logic
(Tailored control flow)

Provides the structure to build reliable, production-ready applications with features like persistent memory, human-in-the-loop, and custom control flow.

Production-Ready Foundations
[4.png, 5.png]

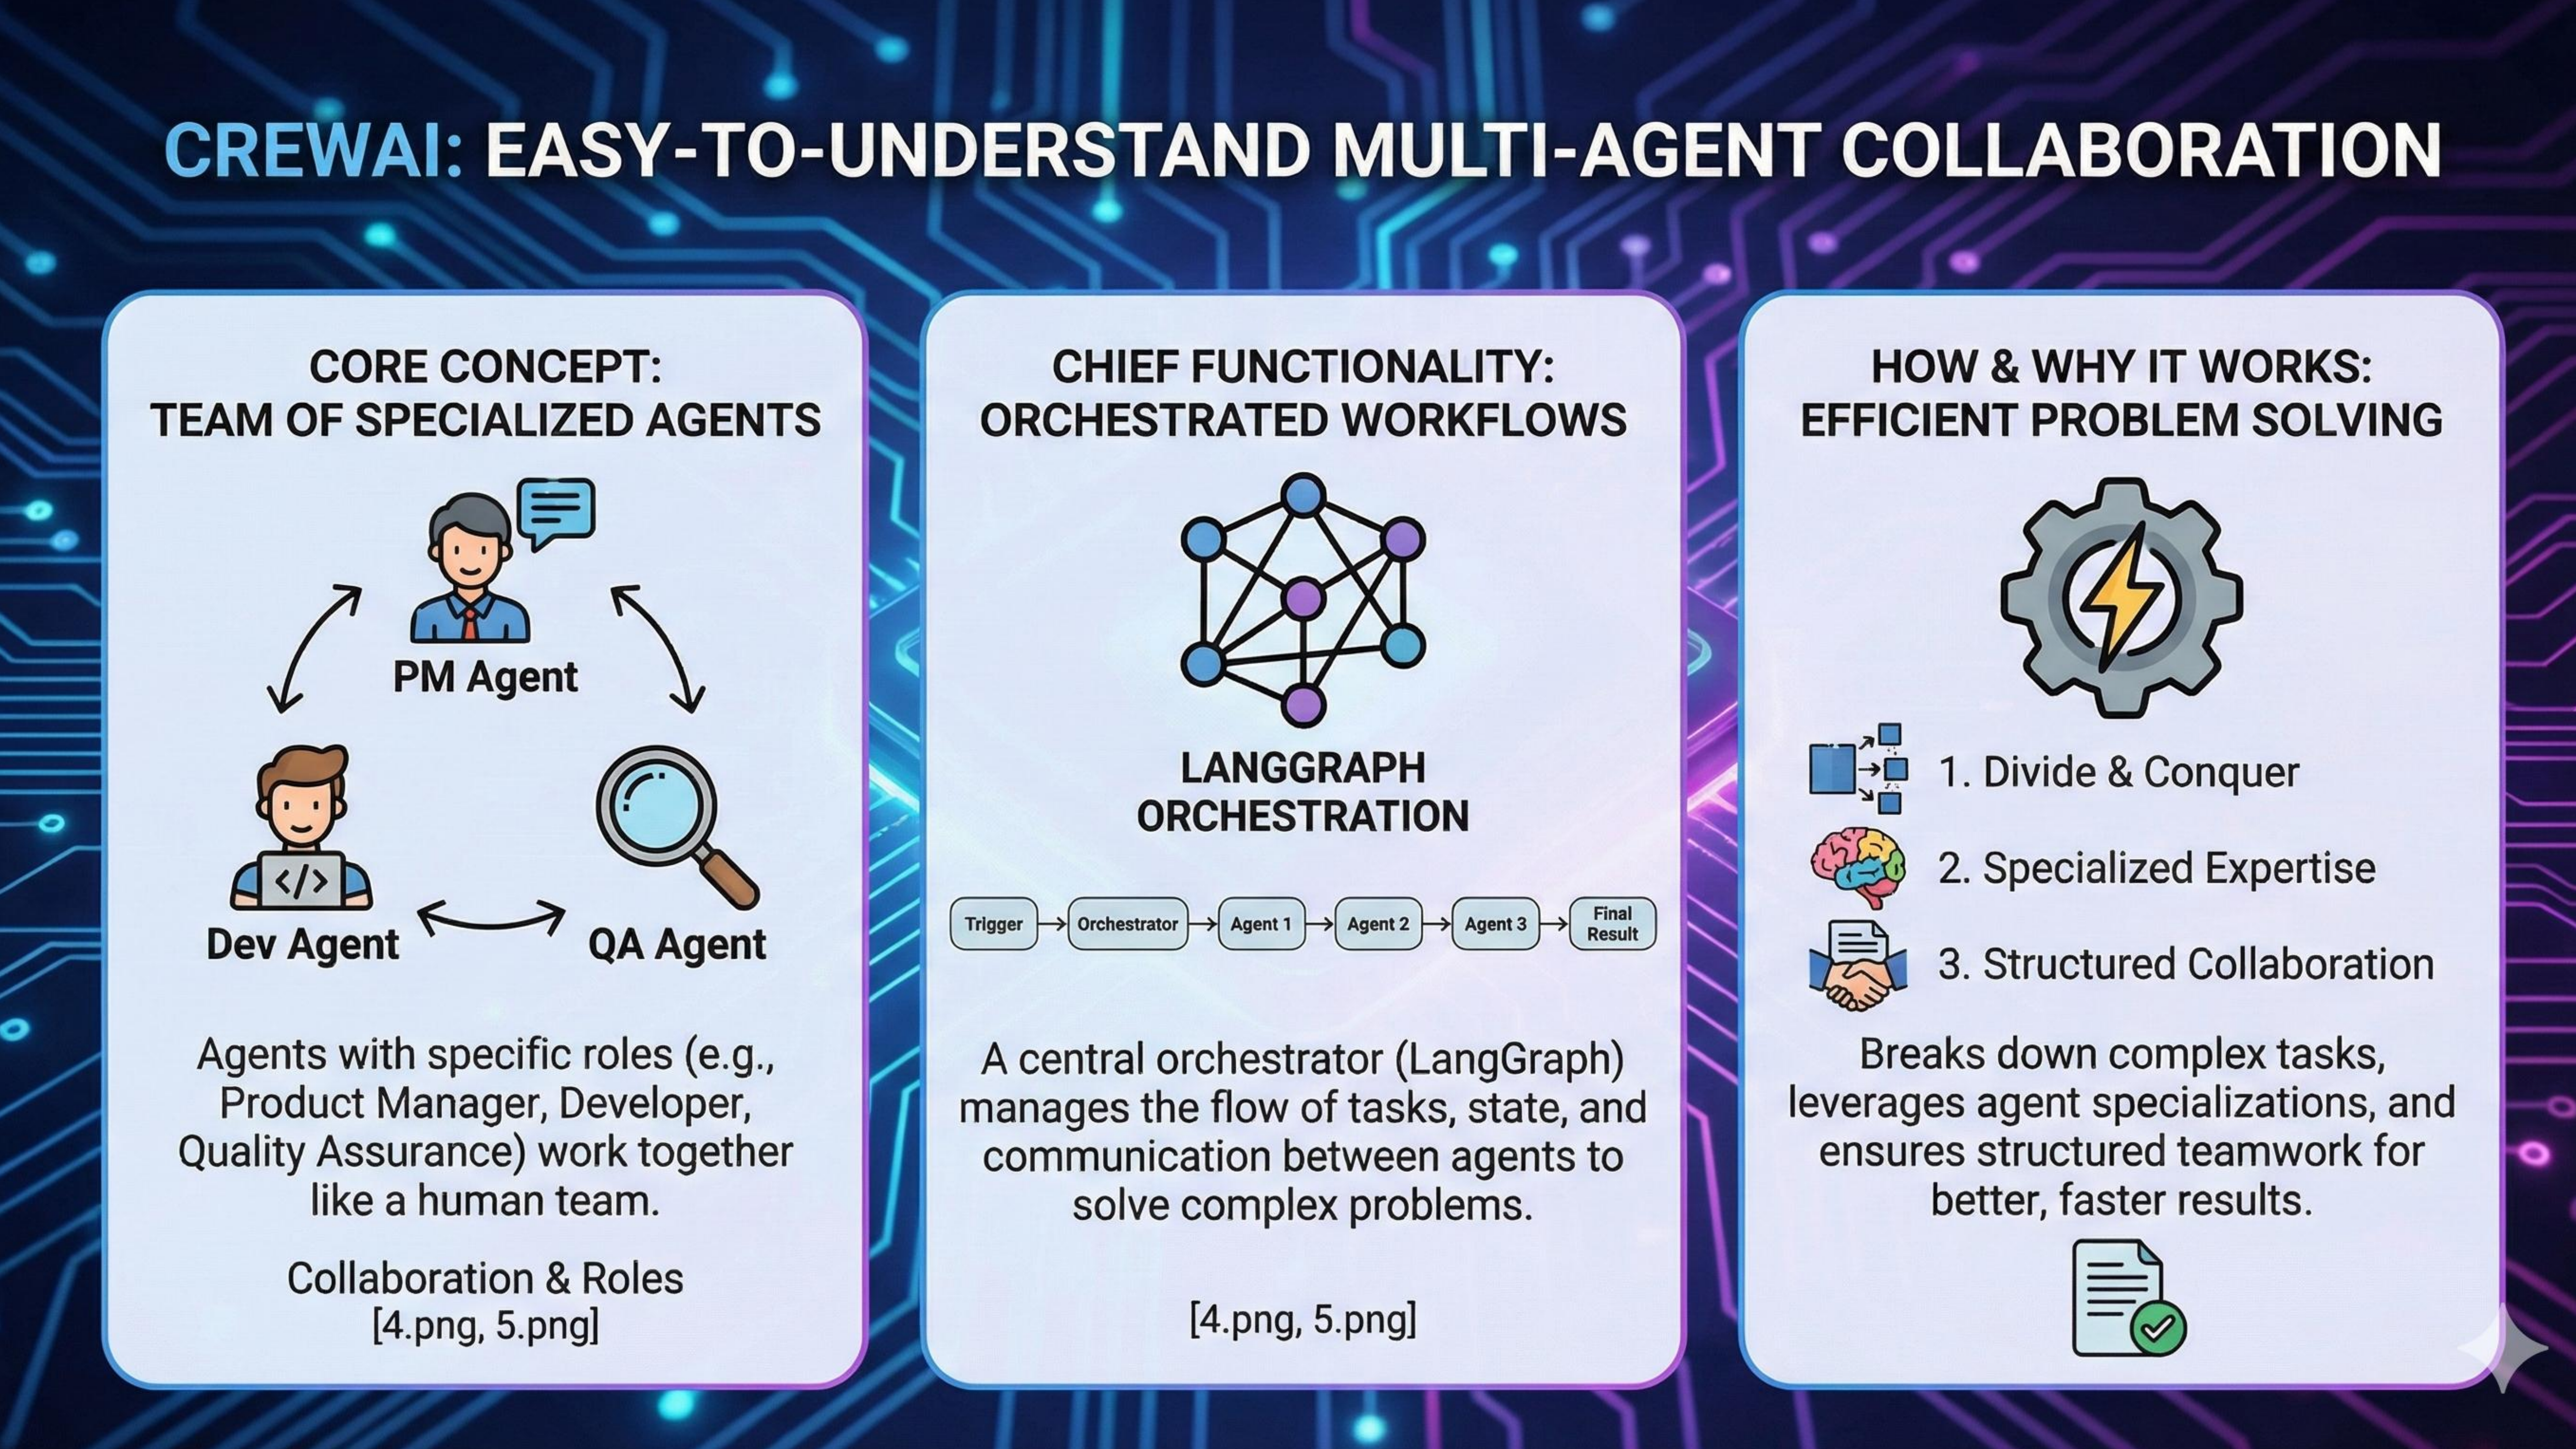
CREWAI: EASY-TO-UNDERSTAND MULTI-AGENT COLLABORATION
CORE CONCEPT:
TEAM OF SPECIALIZED AGENTS
PM Agent
Dev Agent
QA Agent
Agents with specific roles (e.g., Product Manager, Developer, Quality Assurance) work together like a human team.
Collaboration & Roles
[4.png, 5.png]
CHIEF FUNCTIONALITY:
ORCHESTRATED WORKFLOWS
LANGGRAPH ORCHESTRATION
Trigger
Orchestrator
Agent 1
Agent 2
Agent 3
Final Result
A central orchestrator (LangGraph) manages the flow of tasks, state, and communication between agents to solve complex problems.
[4.png, 5.png]
HOW & WHY IT WORKS:
EFFICIENT PROBLEM SOLVING
1. Divide & Conquer
2. Specialized Expertise
3. Structured Collaboration
Breaks down complex tasks, leverages agent specializations, and ensures structured teamwork for better, faster results.

AI AGENT ECOSYSTEM: COMPETITIVE LANDSCAPE (2025-2026)

ORCHESTRATION & APP FRAMEWORKS (LangChain Analogs)

- **🦜🔗 LangChain** (Market Leader, Broad Integrations)
- **Microsoft Agent Framework SDK** (Unified .NET/Python SDK, supersedes Semantic Kernel)
- **LlamaIndex** (Data-First, RAG Focus)
- **Haystack** (Enterprise RAG Pipelines)
- **Google Agent Development Kit (ADK)** (Google Ecosystem Integration)
- **PydanticAI** (Type-Safe, Code-First)
- **smolagents** (Hugging Face, Minimalist)

Building Blocks for Applications

MULTI-AGENT SYSTEMS (CrewAI Analogs)

- **CrewAI** (Role-Based, Process-Driven)
- **Microsoft AutoGen** (Complex Conversational Flows)
- **MetaGPT** (SOP-Based, "Virtual Company")
- **ChatDev** (Software Dev Focus)
- **OpenAI Swarm** (Lightweight, Explicit Handoffs)
- **AgentVerse** (Simulation & Research)

Collaborative Agent Teams

TOOL CALLING & EXECUTION PROTOCOLS

- **Model Context Protocol (MCP)** (Vendor-Neutral Standard for Tools & Data)
- **OpenAI Function Calling** (Native GPT Capability, De Facto Standard)
- **Anthropic Tool Use** (Native Claude Capability)
- **LangChain Tools / Toolkits** (Wrapper Abstractions & Pre-built Sets)

Connecting to the Outside World

MEMORY & PERSISTENCE LAYERS

Vector DBs (Semantic Context)

- **Pinecone** (Managed, Serverless)
- **Weaviate** (Open-Source, Hybrid Search)
- **Chroma** (Local/Server, Simple)
- **Milvus** (Open-Source, Massive Scale)
- **Qdrant** (High-Performance Rust)

Traditional DBs & State (History)

- **PostgreSQL / Redis** (Reliable State & History)
- **LangChain Memory Classes** (Software Abstractions)
- **Microsoft Agent Threads** (Serialized State)

Maintaining Context Over Time

Note: The landscape is dynamic. Frameworks often overlap in functionality (e.g., LangChain can orchestrate multi-agent workflows) and integrate with various memory and tool protocols.

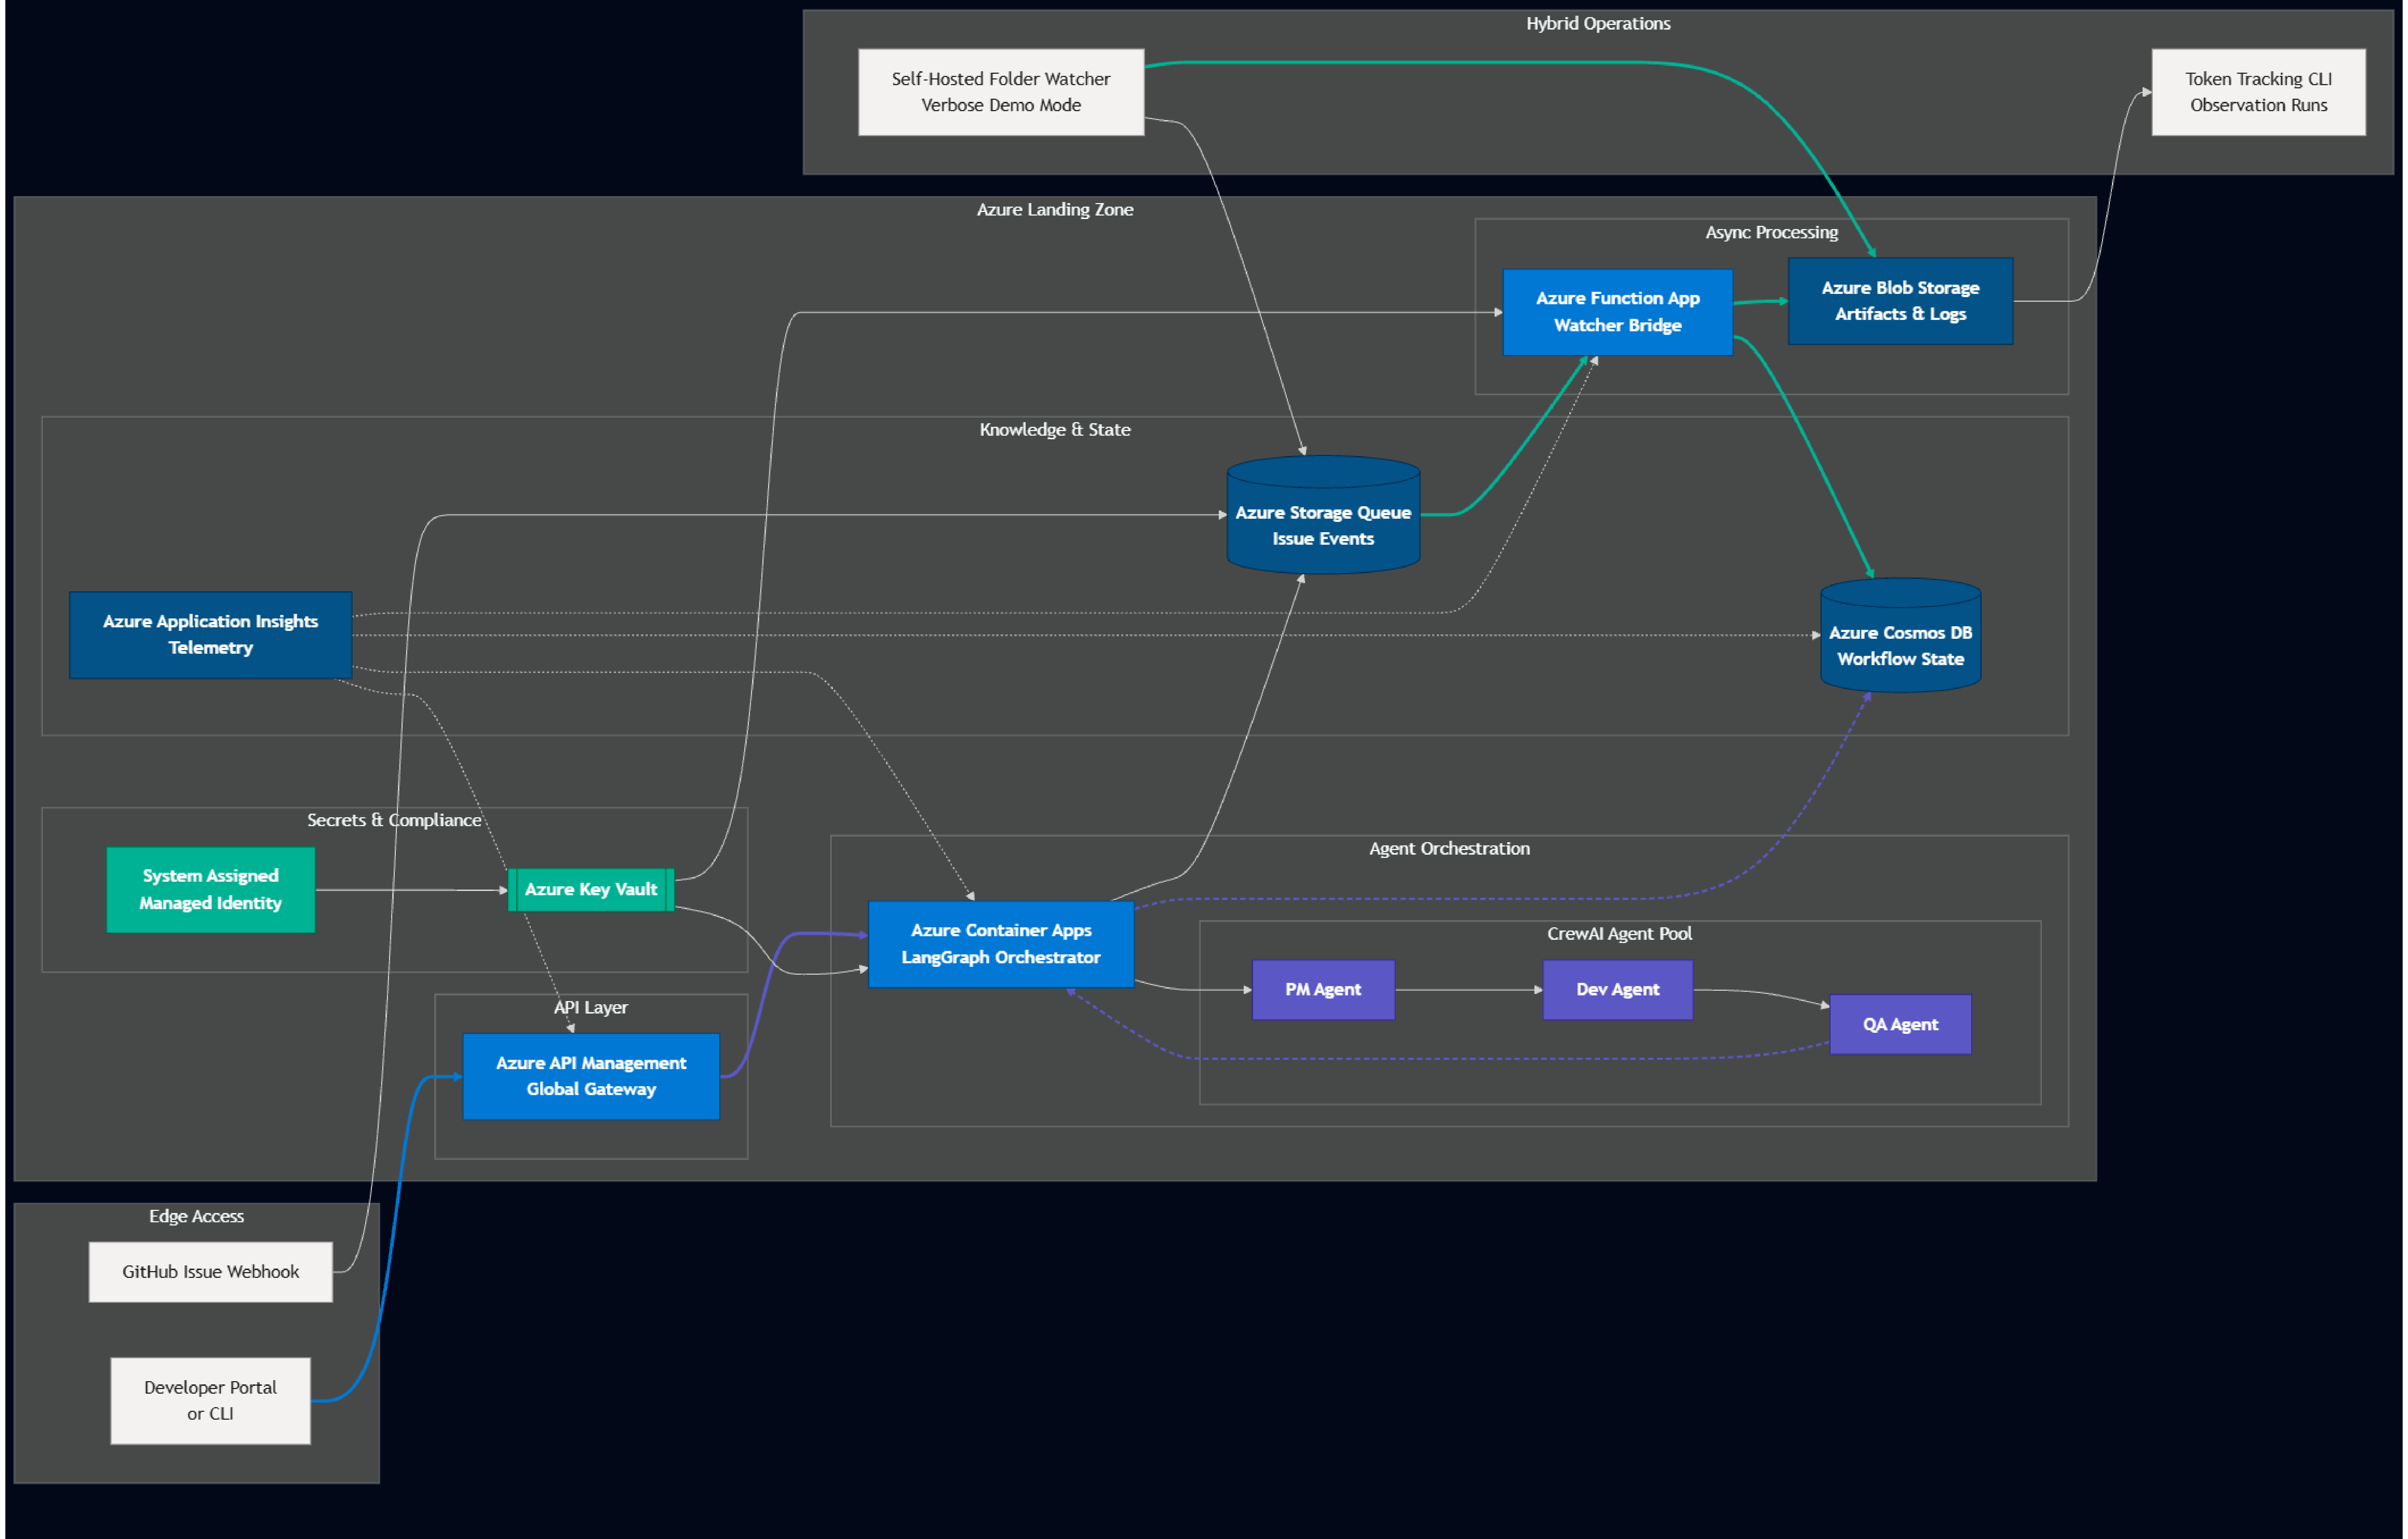
Hybrid Operations
Self-Hosted Folder Watcher
Verbose Demo Mode
Token Tracking CLI
Observation Runs
Azure Landing Zone
Async Processing
Azure Function App
Watcher Bridge
Azure Blob Storage
Artifacts & Logs
Knowledge & State
Azure Storage Queue
Issue Events
Azure Application Insights
Telemetry
Azure Cosmos DB
Workflow State
Secrets & Compliance
System Assigned
Managed Identity
Azure Key Vault
Agent Orchestration
Azure Container Apps
LangGraph Orchestrator
CrewAI Agent Pool
PM Agent
Dev Agent
QA Agent
API Layer
Azure API Management
Global Gateway
Edge Access
GitHub Issue Webhook
Developer Portal
or CLI

Machine Learning (ML) lifecycle

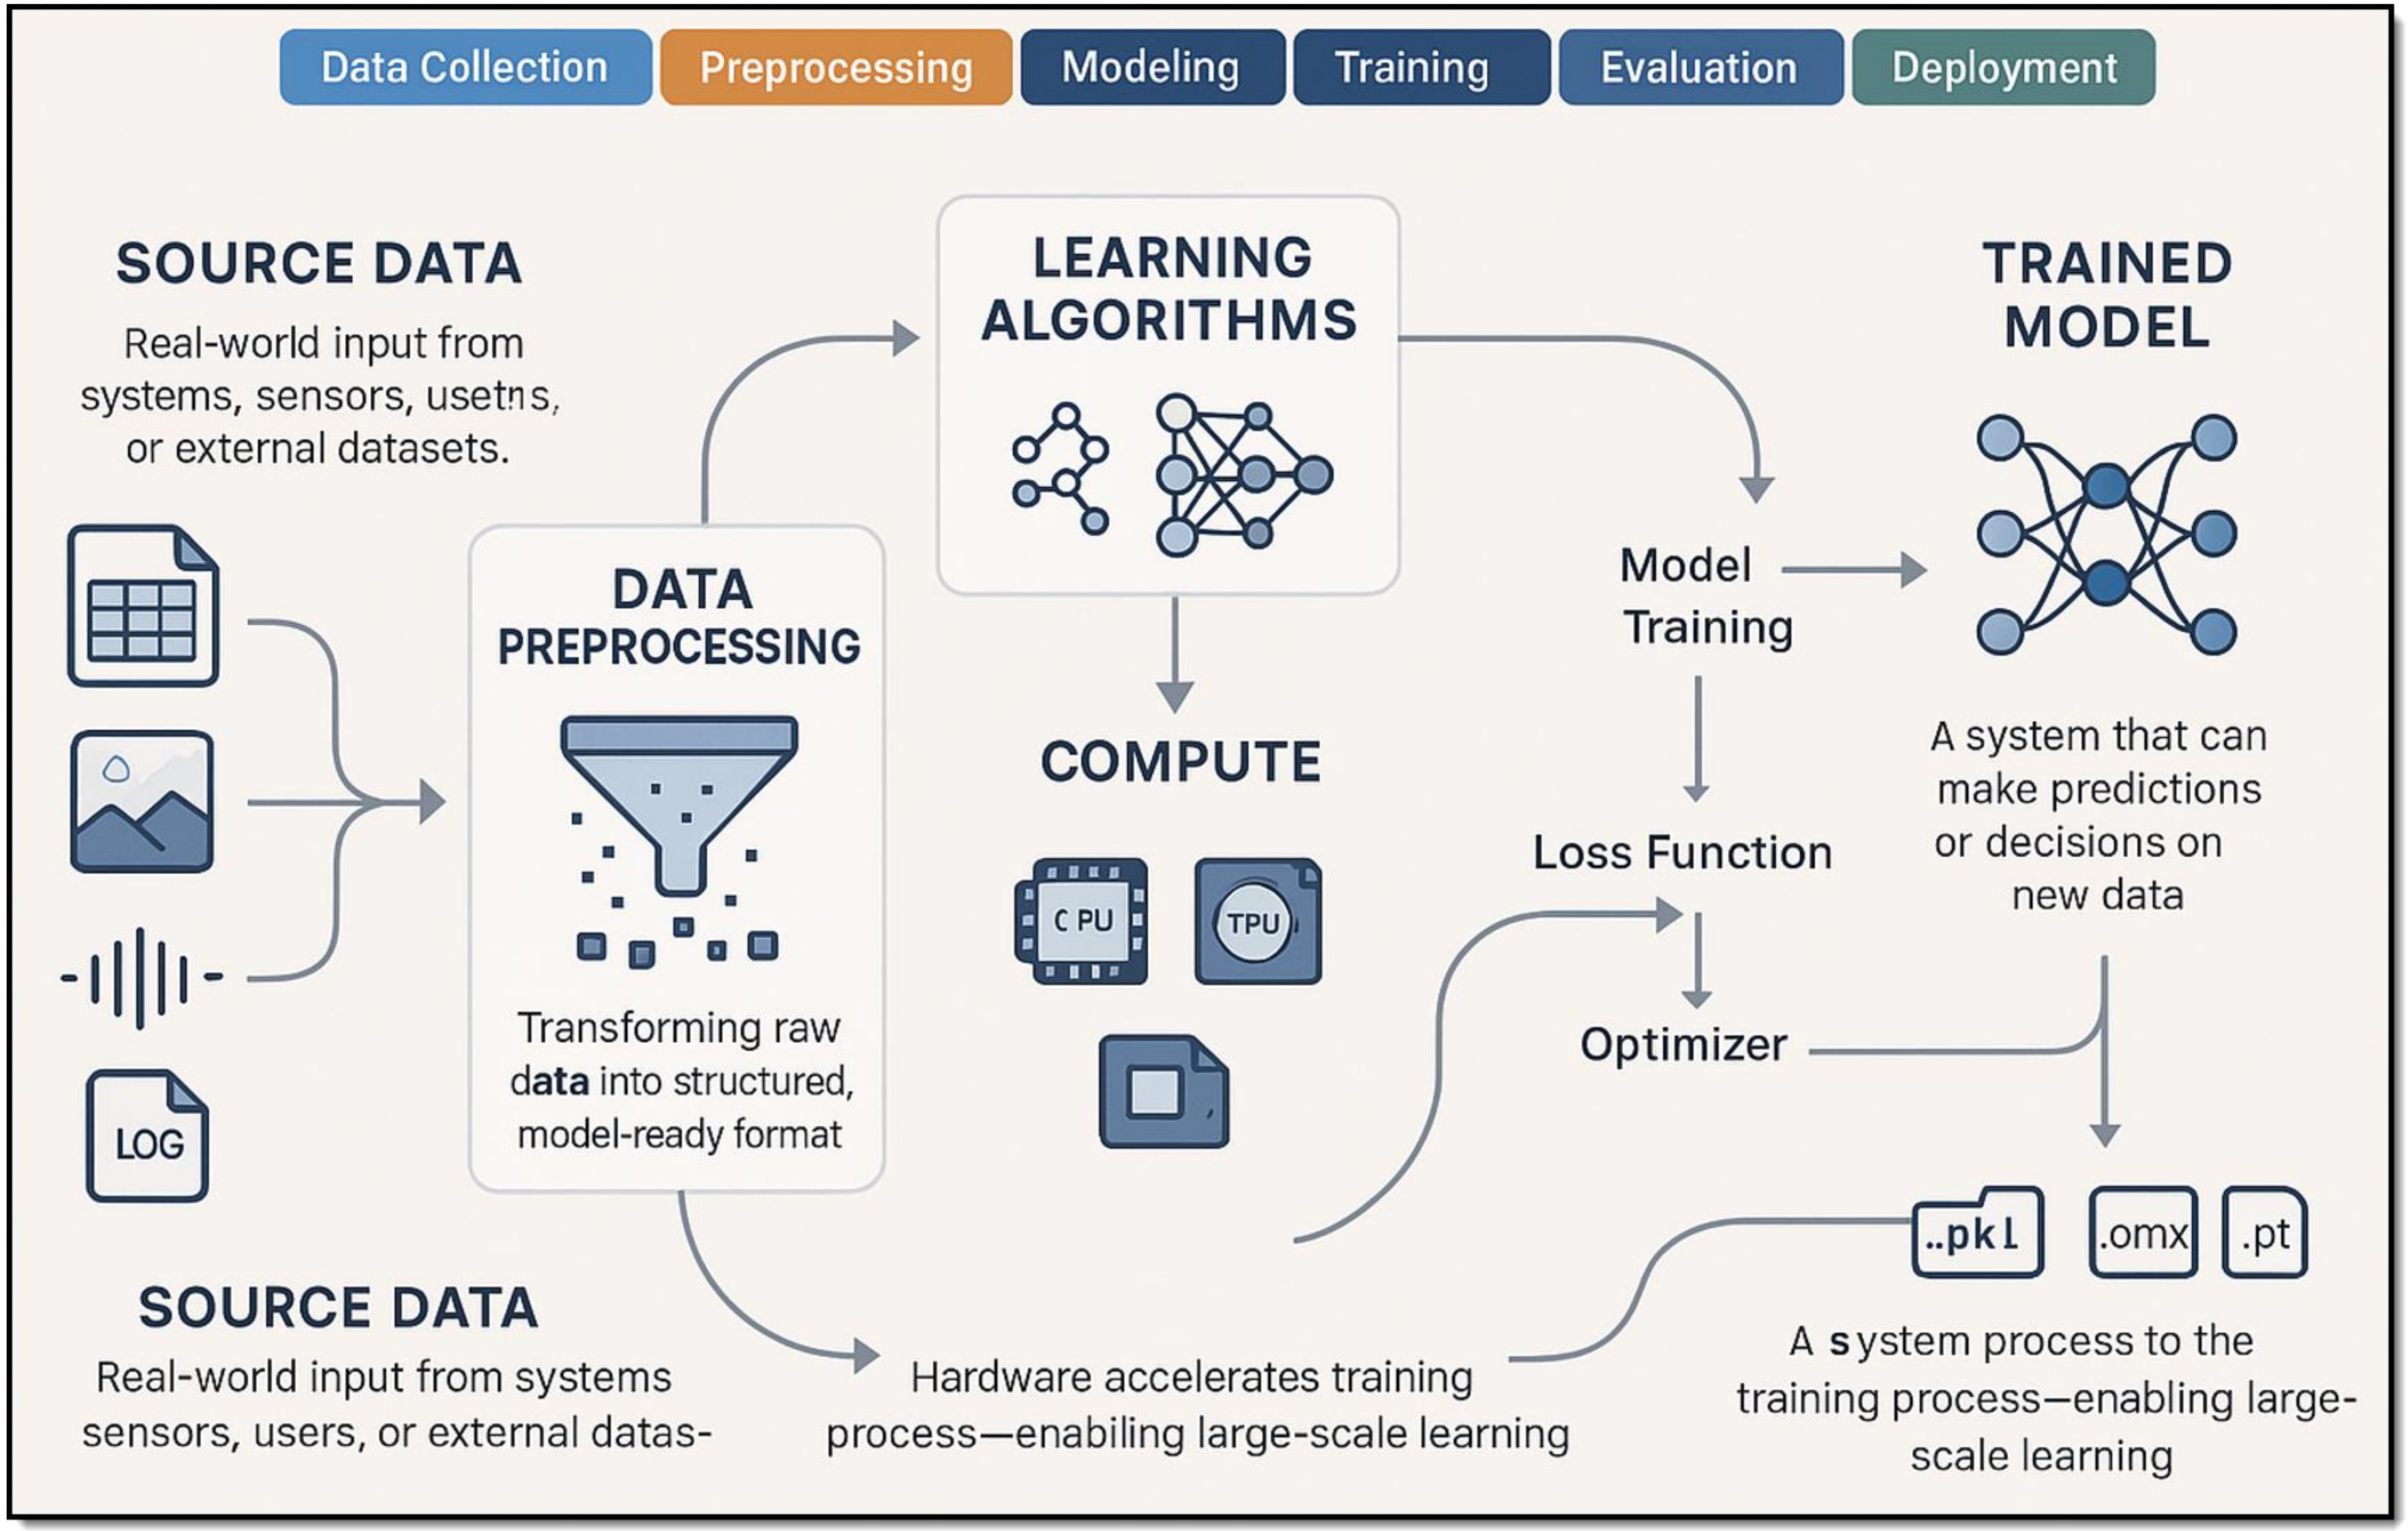

Microsoft

Neural Network

How Characters Become Meaning

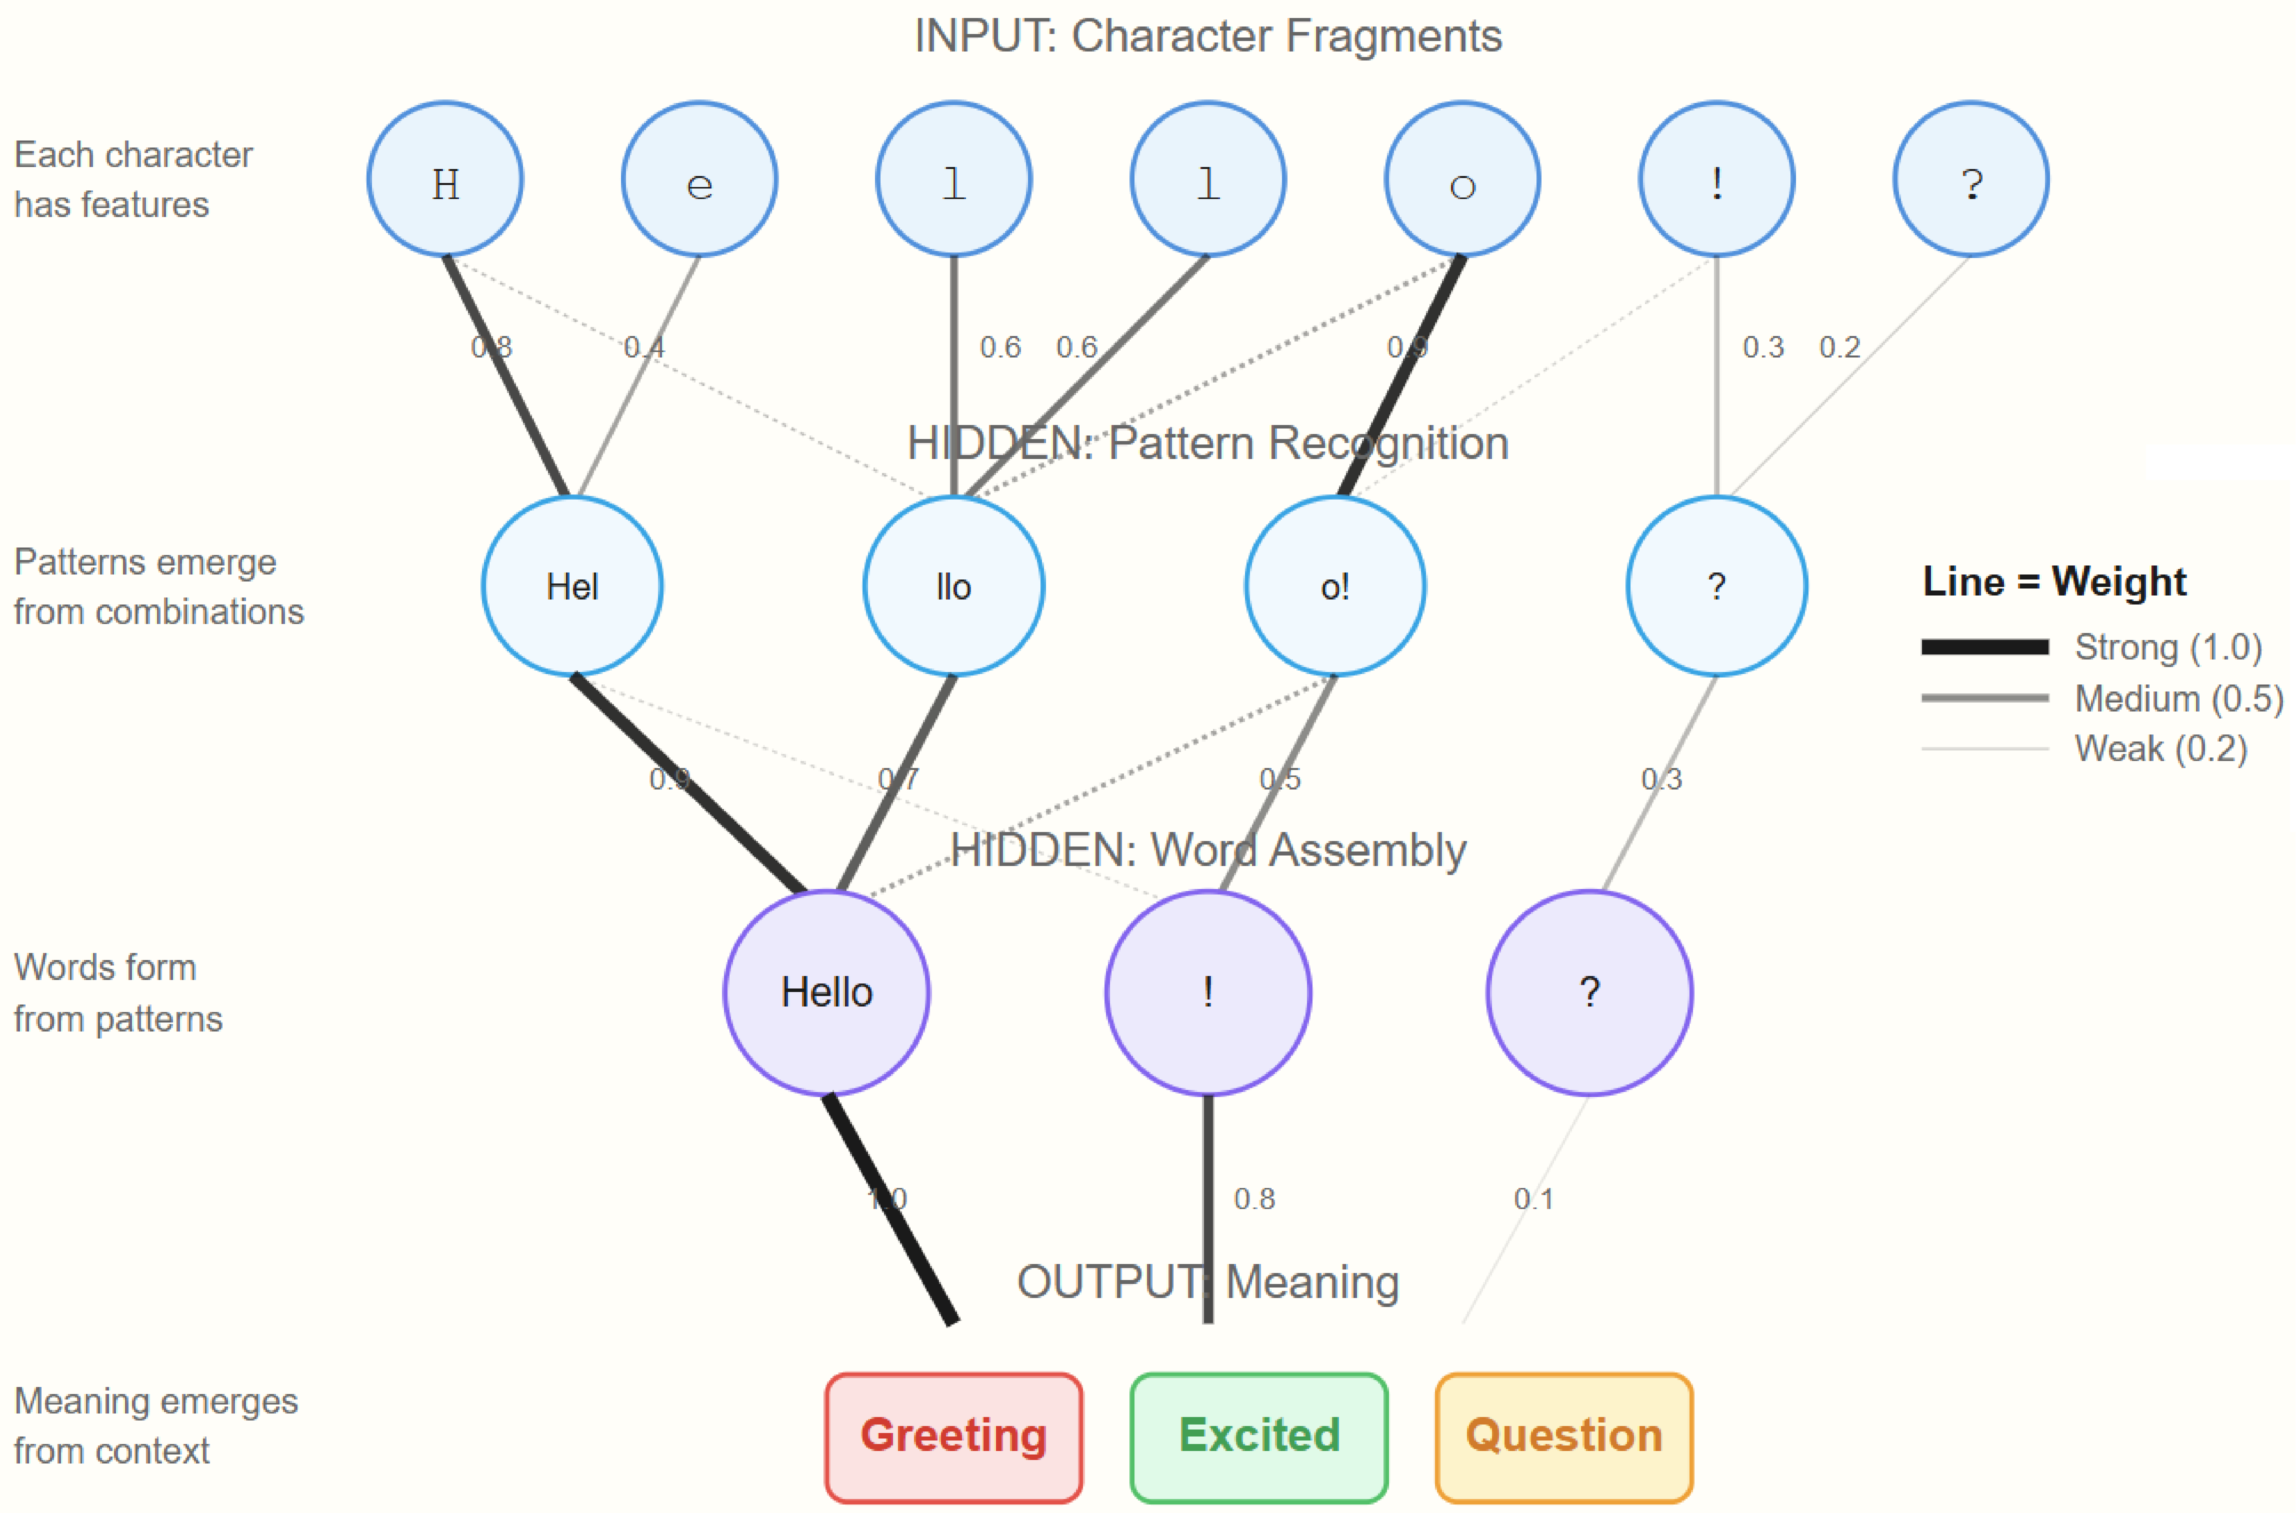

"What I cannot create, I do not understand" - R. Feynman

Neural Network

go.techtrainertim.com/cloud

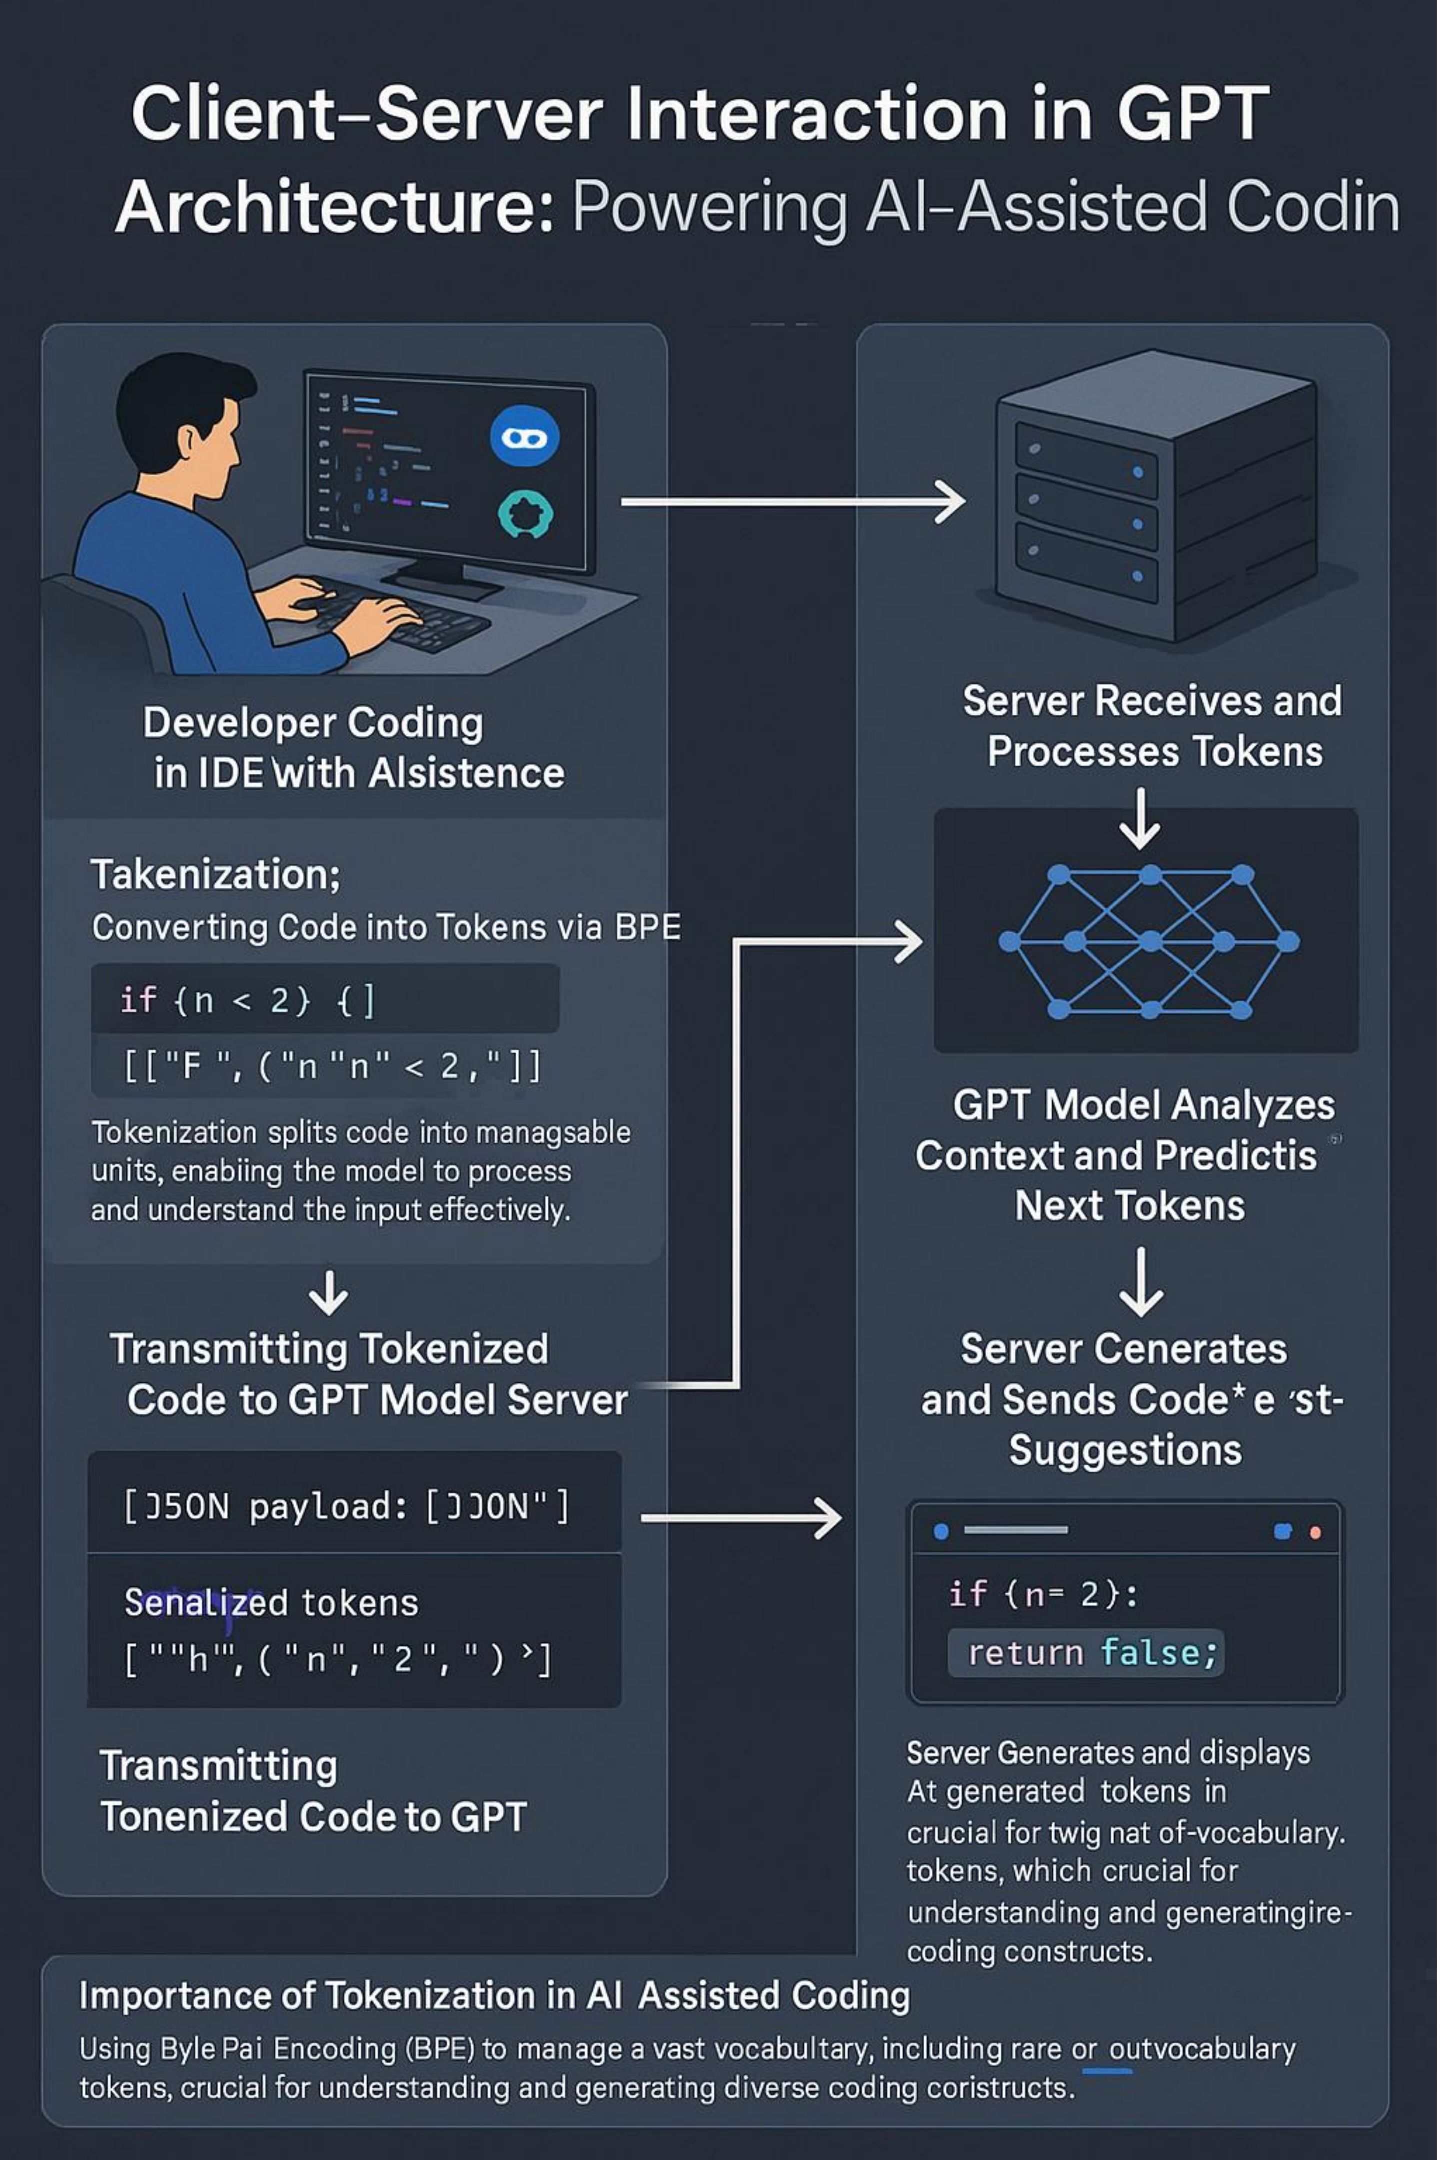

GPT Basic Architecture

- OpenAI GPT 4o Image Generation
- Multimodal generative AI
- Two ways to improve model:
  - Fine tuning
  - Retrieval Augmented Generation (RAG)

Multimodal LLM Processing Flow

From Multiple Input Types to Unified Understanding

Multimodal Inputs

Text

Natural language
Questions, commands

"He
→
→

BPE/
~4 ch

Multimodal Response Generation

Response Components

- Natural Language Text
- Image Generation
- Data Visualization
- Code Generation
- Audio Synthesis

Example Multimodal Response

```
Text: "Based on the image you shared, I can see..."
Code: def analyze_image(img_path): ...
Visual: [Generated diagram or chart]
Audio: [Synthesized explanation]
```

Anthropic Model Context Protocol (MCP)

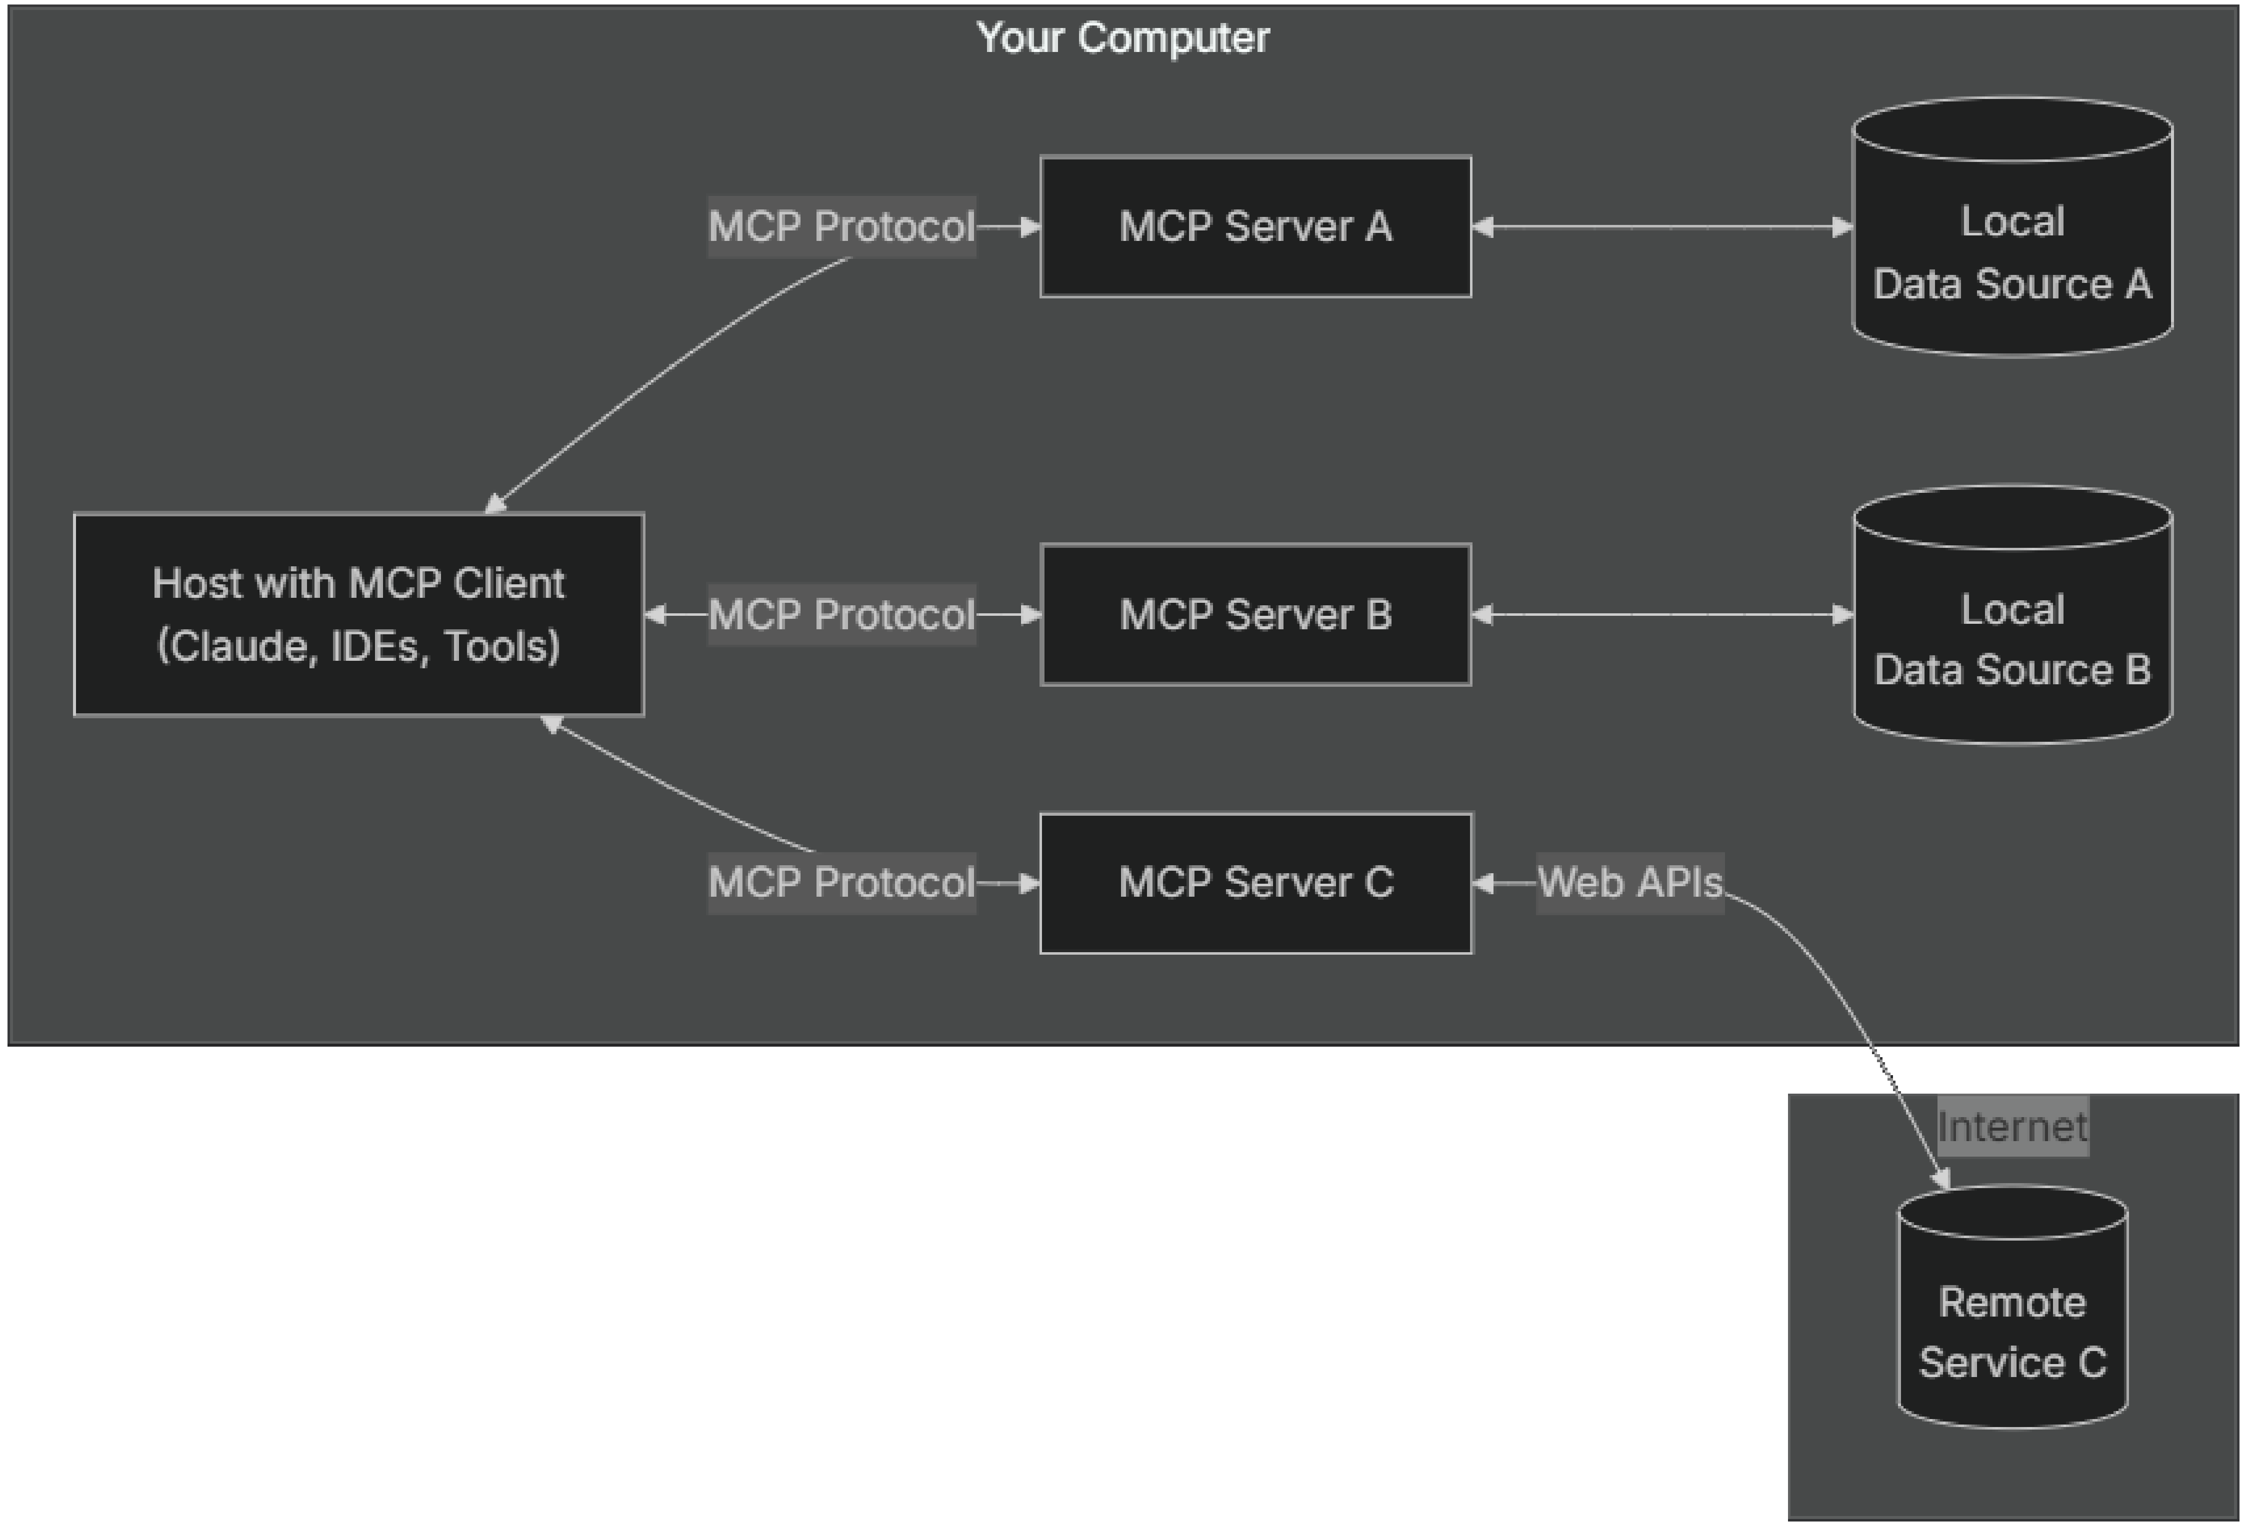

modelcontextprotocol.io

Thank You for Attending!

Please don't forget to complete the O'Reilly course evaluation

Session materials:
github.com/timothywarner-org/agents-pro

Web: techtrainertim.com

COURSE REVAMP • JULY 2026

The Contoso Pinball Gallery Concierge

One Copilot Studio agent, built progressively across four segments.

AGENTS PRO · TIM WARNER

Why One Agent, Built Progressively

We follow one real business scenario, the **Contoso Pinball Gallery Concierge**, and mature it across four segments. Same agent. Rising rigor.

S1 · Foundations

Instructions, knowledge sources, first agent, simulator testing.

S2 · Topics + Flows

Repair Triage topic, Book-a-Service flow, approval step.

S3 · Autonomous

SharePoint event trigger, autonomous intake processing.

S4 · Deploy + ROI

Publish, Analytics, Savings calculator, governance.

Grounded in Microsoft Copilot Studio guidance and the Power Platform Well-Architected Framework.

Power Platform Well-Architected - Mapped

Every segment advances one or more WAF pillars. We name the pillar on the slide it lives on.

SEGMENT 1

Experience Optimization

- Concierge persona & tone
- Instructions block design
- Conversation starters
- Grounded, no-hallucination style

SEGMENT 2

Reliability

- Repair Triage topic
- Grounded inventory answers
- Book-a-Service flow
- Approval gate before booking

SEGMENT 3

Performance Efficiency

- Generative orchestration
- SharePoint event trigger
- Autonomous intake handling
- Copilot Credits budgeting

SEGMENT 4A

Security

- DLP policies
- Least-privilege scopes
- Prompt-injection defense
- Sensitivity labels

SEGMENT 4B

Operational Excellence

- Test panel & Kit suites
- Analytics & Savings calculator
- ALM & solutions
- Governance & CoE

Rule of thumb: if you can’t name the pillar your change advances, the change probably isn’t ready to ship.

learn.microsoft.com/power-platform/well-architected · learn.microsoft.com/microsoft-copilot-studio

Three Sold Patterns, One Agent

The Concierge carries all three

The sell page promises three agent patterns. Instead of three throwaway bots, we fold all three into one Concierge: a **customer-service** front door, an approval-gated booking flow, and an autonomous intake mode.

Inventory Q&A is the customer-service pattern. Book-a-Service with approval is the onboarding pattern. SharePoint-triggered intake is the document-processor pattern.

> **Design heuristic:** one scenario, three modes. Low context-switching for learners, maximum reuse of knowledge and flows.

CUSTOMER SERVICE

Inventory Q&A

- “Do you have Medieval Madness?” The Concierge answers from the
- inventory catalog with maker, year, condition, and price.
- Grounded generative answers, no invented stock or pricing.

ONBOARDING + APPROVAL

Book-a-Service Flow

- Repair Triage runs safe checks, then hands off to the Book-a-Service
- Power Automate flow. An approval step gates the commitment before
- the flow returns a real SVC- reference.

DOCUMENT PROCESSOR

Autonomous Intake

- A repair-intake form lands in SharePoint and fires a “When an item is
- created” event trigger. The agent classifies and routes it autonomously,
- with generative orchestration on.

Source: O’Reilly course description · Contoso Pinball Gallery Concierge asset pack (README.md)

WAF · EXPERIENCE OPTIMIZATION

Build the Concierge by natural language, then ground it in trusted sources.

Inputs we bring

- Concierge scenario & scope
- Four knowledge docs (catalog, history)
- Success metric: grounded answers
- Out-of-scope list (no off-topic)

Artifacts we leave with

- Agent instructions (the persona block)
- Inventory, history, warranty knowledge
- 4 conversation starters
- Simulator smoke test passing

Anti-patterns

- Hundreds of rigid topics
- “Helpful generic chatbot” prompt
- Ungrounded prices and stock
- No refusal or safety surface

```
# Agent instructions - Concierge persona
You are the Contoso Pinball Gallery Concierge, the expert virtual host for a boutique that sells, restores, and services pinball machines. Help with inventory, repairs, and pinball history. Ground every factual answer in the registered knowledge sources; never invent a price or stock. For bookings, quotes, or orders, hand off to the matching deterministic flow.
```

learn.microsoft.com/microsoft-copilot-studio/fundamentals-what-is-copilot-studio · knowledge-copilot-studio

WAF · RELIABILITY

Deterministic where it matters. Generative where it pays.

WHEN IT FIRES

Triggers

- 5–15 phrases per topic, natural variation
- No periods in topic names (breaks exports)
- Run Topic Checker to resolve overlaps
- Prefer intent over keywords

WHAT IT DOES

Topics & variables

- T01 Inventory Lookup + T02 Repair Triage
- *Topic.MachineTitle* scope for clarity
- Question nodes collect symptom & window
- Book-a-Service flow with approval step

WHAT IT KNOWS

Knowledge

- Inventory catalog (the price and stock source)
- Pinball history & research doc
- Warranty and services doc
- Moderation: High for public delivery

Reliability move: disable “AI General Knowledge” so the agent grounds only in configured sources - no invented machine names or prices.

learn.microsoft.com/microsoft-copilot-studio/knowledge-copilot-studio · /guidance/topic-authoring-best-practices

Segment 3 - Autonomous Agents & Event Triggers

WAF · PERFORMANCE EFFICIENCY

Same agent, now it acts on its own when work arrives.

Contoso Pinball Concierge
Generative orchestration

- **SharePoint trigger** — Item created event
- **Repair-intake form** — Lands in a list
- **Autonomous triage** — Reads, classifies
- **Book-a-Service** — Flow + approval

PUBLISH TARGETS

- Teams
- M365 Copilot
- SharePoint
- Web chat
- Mobile

learn.microsoft.com/microsoft-copilot-studio/advanced-generative-actions · autonomous agents & triggers billing

WAF · SECURITY + OPERATIONAL EXCELLENCE

Shipping an agent is a beginning, not an ending.

BEFORE PUBLISH

Testing

- Test panel: golden paths & refusals
- Copilot Studio Kit: batch suites
- Point-test after every YAML edit
- Eval set re-run on each release

AFTER PUBLISH

Analytics

- Engagement / resolution / escalation
- CSAT & topic heatmap
- Savings calculator ROI tile
- Transcript review cadence

ALWAYS

Security

- DLP policies across environments
- Prompt-injection test corpus
- Scope: Sites.Read.All / Files.Read.All
- Sensitivity labels end-to-end

LIFECYCLE

Governance

- Solutions + managed environments
- Dev → Test → Prod ALM
- Approval gates & Responsible AI review
- CoE Starter Kit adoption

≥ 80%	< 5%	100%	0
RESOLUTION RATE	ESCALATION RATE	GROUNDED ANSWERS	DLP VIOLATIONS

learn.microsoft.com/microsoft-copilot-studio/analytics-overview · admin-data-loss-prevention · power-platform/alm

Topic Map - Contoso Pinball Concierge

Two custom topics plus system topics and two flows. Each earns its place through the conversation lifecycle.

T01

Inventory Lookup

- Custom topic
- Grounded in inventory catalog
- Maker, year, condition, price

T02

Repair Triage

- Custom topic
- Grounded in repair playbook
- Hands off to Book-a-Service

SYSTEM

Greeting / Fallback / Escalate

- System topics
- Cover the conversation edges
- Escalate to a human host

FLOW

Book a Service

- Power Automate flow
- Collects machine, symptom, window
- Returns a real SVC-reference

FLOW

Request a Quote

- Power Automate flow
- Maps symptom to a price band
- Returns a QUO-reference

Topic names never contain periods - required for clean solution exports.

Flex the Muscles - Live Demo Prompts

Run these in order during the final build. Each one lights up a different capability you taught.

KNOWLEDGE

1 · Grounded answer

```
Do you have Medieval
Madness, and how much?
```

Grounded answer: Williams 1997, $11,500, in stock.

TOPICS

2 · Repair triage

```
My flipper is weak on my
Godzilla.
```

T02 Repair Triage runs safe checks, then offers a booking.

FLOWS

3 · Book a service

```
Book a service visit for my
Godzilla.
```

Book-a-Service flow returns a real SVC-reference.

GROUNDING

4 · Compare titles

```
Which of your machines are
Bally?
```

Lists Attack from Mars, Twilight Zone, Fireball from catalog.

Instructor tip: run the *refusal prompt* (“give me a competitor’s prices”) second-to-last, not first. Showing refusal after success lands the governance message.

Demo runs in the Test panel; each prompt maps to a row in evals/eval-set.md.

Resources & Next Steps

Everything you need to rebuild this agent Monday morning.

Copilot Studio

- learn.microsoft.com/microsoft-copilot-studio
- /fundamentals-what-is-copilot-studio
- /knowledge-copilot-studio
- /guidance/topic-authoring-best-practices
- /advanced-generative-actions
- /publication-connect-bot-to-microsoft-teams

Well-Architected & ALM

- learn.microsoft.com/power-platform/well-architected
- /power-platform/alm
- /microsoft-copilot-studio/admin-data-loss-prevention
- /microsoft-copilot-studio/analytics-overview
- Copilot Studio Kit (batch testing)

Contoso Pinball Asset Pack

- Contoso Pinball Gallery Concierge/README.md
- instructions.md (the persona block)
- knowledge/ (catalog, history, warranty)
- flows/deterministic-flow-ideas.md
- evals/eval-set.md

Your turn. Load instructions.md plus the four knowledge docs into a blank agent and publish it to your own Teams tenant this week.

`github.com/timothywarner-org/agents-pro · techtrainertim.com`

VERIFIED AGAINST MICROSOFT LEARN · JULY 2026

What’s New Since Last Delivery

The Copilot Studio deltas that moved between writing the deck and teaching it.

Model Lineup Changed the Most

Generative orchestration now lets you pick the reasoning model on the Overview tab.

GA (use in production)

- **GPT-5 Chat**
- **Claude Opus 4.6** (recommended orchestrator)
- Claude Sonnet 4.6, Opus 4.7
- GPT-4.1 (default, not best)

Preview (early, may change)

- GPT-5 Reasoning
- GPT-5 Auto

Experimental (mention only)

- GPT-5.5 Reasoning
- Sonnet 4.5 retired - migrate off it

Claude models are external; tenant admin must approve.

On stage: always say “GPT-5 Chat,” never bare “GPT-5.” The Reasoning and Auto variants are still preview.

learn.microsoft.com/microsoft-copilot-studio/authoring-select-agent-model

Multi-Agent & MCP Moved On

The Segment 3 story is bigger than “connected agents” now.

“Add other agents” taxonomy

- **Child agents** - lightweight, embedded (GA)
- **Connected Copilot Studio agents** - reusable (GA)
- **A2A protocol** - open standard, GA April 2026
- Foundry, Fabric, M365 Agents SDK - preview

Doc page renamed: authoring-add-other-agents.

MCP wiring

- **Streamable HTTP only** - SSE retired after Aug 2025
- **Onboarding wizard** - Tools > Add a tool > MCP
- Custom connector is now the fallback
- MCP requires generative orchestration

FastMCP: serve Streamable HTTP, not stdio/SSE.

Build once, connect everywhere: the same MCP tool serves Claude, ChatGPT, and Copilot Studio.

learn.microsoft.com/microsoft-copilot-studio/authoring-add-other-agents · /mcp-add-existing-server-to-agent

Evaluations Are Built In. Authoring Goes Pro-Code.

The Segment 4 testing story and a developer on-ramp both leveled up.

Native agent evaluations (GA)

- Test sets up to 100 cases
- Generate cases from knowledge & topics
- Activity maps + version compare
- REST API + connector automation
- Multi-turn evaluation now GA

Copilot Studio Kit is now complementary, not the only path.

VS Code extension (GA)

- Clone an agent to your machine
- Edit agent-definition YAML locally
- Git, pull requests, deploy back
- **Claude Code & GitHub Copilot** named as authoring agents

Agent-driven authoring of an agent. Yes, really.

Move the batch story: native evaluations are the new primary, and the deck's YAML can live in Git.

learn.microsoft.com/microsoft-copilot-studio/analytics-agent-evaluation-intro · /visual-studio-code-extension-overview